

```

"dimension": dimension.value,
"contradictions": [self._serialize_contradiction(c) for c in contradictions],
"total_contradictions": len(contradictions),
"high_severity_count": sum(1 for c in contradictions if c.severity > get_param
    eter_loader().get("saaaaaa.analysis.contradiction_deteccion.PolicyContradictionDetector._i
    nitialize_pdm_patterns").get("auto_param_L469_81", 0.7)),
"coherence_metrics": coherence_metrics,
"recommendations": recommendations,
"knowledge_graph_stats": self._get_graph_statistics()
}

def _extract_policy_statements(
    self,
    text: str,
    dimension: PolicyDimension
) -> list[PolicyStatement]:
    """Extrae declaraciones de política estructuradas del texto"""
    doc = self.nlp(text)
    statements = []

    for sent in doc.sents:
        # Analizar entidades nombradas
        entities = [ent.text for ent in sent.ents]

        # Extraer marcadores temporales
        temporal_markers = self._extract_temporal_markers(sent.text)

        # Extraer afirmaciones cuantitativas
        quantitative_claims = self._extract_quantitative_claims(sent.text)

        # Determinar rol semántico
        semantic_role = self._determine_semantic_role(sent)

        # Identificar dependencias
        dependencies = self._identify_dependencies(sent, doc)

        statement = PolicyStatement(
            text=sent.text,
            dimension=dimension,
            position=(sent.start_char, sent.end_char),
            entities=entities,
            temporal_markers=temporal_markers,
            quantitative_claims=quantitative_claims,
            context_window=self._get_context_window(text, sent.start_char,
            sent.end_char),
            semantic_role=semantic_role,
            dependencies=dependencies
        )

        statements.append(statement)

    return statements

def _generate_embeddings(
    self,
    statements: list[PolicyStatement]
) -> list[PolicyStatement]:
    """Genera embeddings semánticos para las declaraciones"""
    texts = [stmt.text for stmt in statements]
    embeddings = self.semantic_model.encode(texts, convert_to_numpy=True)

    # Crear nuevas instancias con embeddings
    enhanced_statements = []
    for stmt, embedding in zip(statements, embeddings, strict=False):
        enhanced = PolicyStatement(
            text=stmt.text,
            dimension=stmt.dimension,
            position=stmt.position,
            embedding=embedding
        )
        enhanced_statements.append(enhanced)

```

```

entities=stmt.entities,
temporal_markers=stmt.temporal_markers,
quantitative_claims=stmt.quantitative_claims,
embedding=embedding,
context_window=stmt.context_window,
semantic_role=stmt.semantic_role,
dependencies=stmt.dependencies
)
enhanced_statements.append(enhanced)

return enhanced_statements

@calibrated_method("saaaaaa.analysis.contradiction_deteccion.PolicyContradictionDetector._build_knowledge_graph")
def _build_knowledge_graph(self, statements: list[PolicyStatement]) -> None:
    """Construye grafo de conocimiento para razonamiento"""
    self.knowledge_graph.clear()

    for i, stmt in enumerate(statements):
        node_id = f"stmt_{i}"
        self.knowledge_graph.add_node(
            node_id,
            text=stmt.text[:100],
            dimension=stmt.dimension.value,
            entities=stmt.entities,
            semantic_role=stmt.semantic_role
        )

    # Conectar con declaraciones relacionadas
    for j, other in enumerate(statements):
        if i != j:
            similarity = self._calculate_similarity(stmt, other)
            if similarity > get_parameter_loader().get("saaaaaa.analysis.contradiction_deteccion.PolicyContradictionDetector._build_knowledge_graph").get("auto_param_L562_36", 0.7): # Umbral de relación
                self.knowledge_graph.add_edge(
                    f"stmt_{i}",
                    f"stmt_{j}",
                    weight=similarity,
                    relation_type=self._determine_relation_type(stmt, other)
                )

def _detect_semantic_contradictions(
    self,
    statements: list[PolicyStatement]
) -> list[ContradictionEvidence]:
    """Detecta contradicciones semánticas usando transformers"""
    contradictions = []

    for i, stmt_a in enumerate(statements):
        for stmt_b in statements[i + 1:]:
            if stmt_a.embedding is not None and stmt_b.embedding is not None:
                # Calcular similaridad coseno
                similarity = 1 - cosine(stmt_a.embedding, stmt_b.embedding)

                # Verificar contradicción usando clasificador
                combined_text = f"{stmt_a.text} [SEP] {stmt_b.text}"
                contradiction_score = self._classify_contradiction(combined_text)

                if contradiction_score > get_parameter_loader().get("saaaaaa.analysis.contradiction_deteccion.PolicyContradictionDetector._build_knowledge_graph").get("auto_param_L587_45", 0.7) and similarity > get_parameter_loader().get("saaaaaa.analysis.contradiction_deteccion.PolicyContradictionDetector._build_knowledge_graph").get("auto_param_L587_66", 0.5):
                    # Calcular confianza Bayesiana
                    confidence = self.bayesian_calculator.calculate_posterior(
                        evidence_strength=contradiction_score,
                        observations=len(stmt_a.entities) + len(stmt_b.entities),

```

```

        domain_weight=self._get_domain_weight(stmt_a.dimension)
    )

evidence = ContradictionEvidence(
    statement_a=stmt_a,
    statement_b=stmt_b,
    contradiction_type=ContradictionType.SEMANTIC_OPPOSITION,
    confidence=confidence,
    severity=self._calculate_severity(stmt_a, stmt_b),
    semantic_similarity=similarity,
    logical_conflict_score=contradiction_score,
    temporal_consistency=True,
    numerical_divergence=None,
    affected_dimensions=[stmt_a.dimension, stmt_b.dimension],
    resolutionSuggestions=self._suggest_resolutions(
        ContradictionType.SEMANTIC_OPPOSITION
    )
)
contradictions.append(evidence)

return contradictions

def _detect_numerical_inconsistencies(
    self,
    statements: list[PolicyStatement]
) -> list[ContradictionEvidence]:
    """Detecta inconsistencias numéricas con análisis estadístico"""
    contradictions = []

    for i, stmt_a in enumerate(statements):
        for stmt_b in statements[i + 1:]:
            if stmt_a.quantitative_claims and stmt_b.quantitative_claims:
                for claim_a in stmt_a.quantitative_claims:
                    for claim_b in stmt_b.quantitative_claims:
                        if self._are_comparable_claims(claim_a, claim_b):
                            divergence = self._calculate_numerical_divergence(
                                claim_a,
                                claim_b
                            )

                            if divergence is not None and divergence > get_parameter_loader().get("saaaaaa.analysis.contradiction_deteccion.PolicyContradictionDetector._build_knowledege_graph").get("auto_param_L632_75", 0.2):
                                # Test estadístico de significancia
                                p_value = self._statistical_significance_test(
                                    claim_a,
                                    claim_b
                                )

                                if p_value < get_parameter_loader().get("saaaaaa.analy sis.contradiction_deteccion.PolicyContradictionDetector._build_knowledge_graph").get("auto_param_L639_49", 0.05): # Significancia estadística
                                    confidence =
self.bayesian_calculator.calculate_posterior(
                                        evidence_strength=1 - p_value,
                                        observations=2,
                                        domain_weight=1.5 # Mayor peso para evidencia
numérica
)
evidence = ContradictionEvidence(
    statement_a=stmt_a,
    statement_b=stmt_b,
contradiction_type=ContradictionType.NUMERICAL_INCONSISTENCY,
    confidence=confidence,
    severity=min(get_parameter_loader().get("saaaa
aa.analysis.contradiction_deteccion.PolicyContradictionDetector._build_knowledge_graph").g

```

```

et("auto_param_L651_57", 1.0), divergence),
    semantic_similarity=get_parameter_loader().get
("saaaaaaa.analysis.contradiction_deteccion.PolicyContradictionDetector._build_knowledge_gr
aph").get("auto_param_L652_64", 0.0),
        logical_conflict_score=divergence,
        temporal_consistency=True,
        numerical_divergence=divergence,
        affected_dimensions=[stmt_a.dimension],

resolutionSuggestions=self._suggest_resolutions(
    ContradictionType.NUMERICAL_INCONSISTENCY
),
    statistical_significance=p_value
)
contradictions.append(evidence)

return contradictions

def _detect_temporal_conflicts(
    self,
    statements: list[PolicyStatement]
) -> list[ContradictionEvidence]:
    """Detecta conflictos temporales usando verificación lógica"""
    contradictions = []

    # Filtrar declaraciones con marcadores temporales
    temporal_statements = [s for s in statements if s.temporal_markers]

    if len(temporal_statements) >= 2:
        is_consistent, conflicts = self.temporal_verifier.verify_temporal_consistency(
            temporal_statements
        )

        for conflict in conflicts:
            stmt_a = conflict['event_a']['statement']
            stmt_b = conflict['event_b']['statement']

            confidence = self.bayesian_calculator.calculate_posterior(
                evidence_strength=get_parameter_loader().get("saaaaaaa.analysis.contrad
iction_deteccion.PolicyContradictionDetector._build_knowledge_graph").get("auto_param_L686
_38", 0.9), # Alta confianza en lógica temporal
                observations=len(conflicts),
                domain_weight=1.2
            )

            evidence = ContradictionEvidence(
                statement_a=stmt_a,
                statement_b=stmt_b,
                contradiction_type=ContradictionType.TEMPORAL_CONFLICT,
                confidence=confidence,
                severity=get_parameter_loader().get("saaaaaaa.analysis.contradiction_de
teccion.PolicyContradictionDetector._build_knowledge_graph").get("auto_param_L696_29",
0.8), # Los conflictos temporales son severos
                semantic_similarity=self._calculate_similarity(stmt_a, stmt_b),
                logical_conflict_score=get_parameter_loader().get("saaaaaaa.analysis.co
ntradiction_deteccion.PolicyContradictionDetector._build_knowledge_graph").get("auto_param
_L698_43", 1.0),
                temporal_consistency=False,
                numerical_divergence=None,
                affected_dimensions=[PolicyDimension.PROGRAMATICO],
                resolutionSuggestions=self._suggest_resolutions(
                    ContradictionType.TEMPORAL_CONFLICT
                )
            )
            contradictions.append(evidence)

    return contradictions

```

```

def _detect_logical_incompatibilities(
    self,
    statements: list[PolicyStatement]
) -> list[ContradictionEvidence]:
    """Detecta incompatibilidades lógicas usando razonamiento en grafo"""
    contradictions = []

    # Buscar ciclos negativos en el grafo (indicativos de contradicción)
    try:
        negative_cycles = nx.negative_edge_cycle(
            self.knowledge_graph,
            weight='weight'
        )

        for cycle in negative_cycles:
            # Extraer declaraciones del ciclo
            stmt_indices = [int(node.split('_')[1]) for node in cycle]
            cycle_statements = [statements[i] for i in stmt_indices]

            # Analizar incompatibilidad lógica
            for i in range(len(cycle_statements)):
                stmt_a = cycle_statements[i]
                stmt_b = cycle_statements[(i + 1) % len(cycle_statements)]

                if self._has_logical_conflict(stmt_a, stmt_b):
                    confidence = self.bayesian_calculator.calculate_posterior(
                        evidence_strength=get_parameter_loader().get("saaaaaa.analysis.contradiction_deteccion.PolicyContradictionDetector._build_knowledge_graph").get("auto_param_L736_46", 0.85),
                        observations=len(cycle),
                        domain_weight = get_parameter_loader().get("saaaaaa.analysis.contradiction_deteccion.PolicyContradictionDetector._build_knowledge_graph").get("domain_weight", 1.0) # Refactored
                    )

                    evidence = ContradictionEvidence(
                        statement_a=stmt_a,
                        statement_b=stmt_b,
                        contradiction_type=ContradictionType.LOGICAL_INCOMPATIBILITY,
                        confidence=confidence,
                        severity=get_parameter_loader().get("saaaaaa.analysis.contradiction_deteccion.PolicyContradictionDetector._build_knowledge_graph").get("auto_param_L746_37", 0.7),
                        semantic_similarity=self._calculate_similarity(stmt_a,
                        stmt_b),
                        logical_conflict_score=get_parameter_loader().get("saaaaaa.analysis.contradiction_deteccion.PolicyContradictionDetector._build_knowledge_graph").get("auto_param_L748_51", 0.9),
                        temporal_consistency=True,
                        numerical_divergence=None,
                        affected_dimensions=[stmt_a.dimension, stmt_b.dimension],
                        resolution_suggestions=self._suggest_resolutions(
                            ContradictionType.LOGICAL_INCOMPATIBILITY
                        ),
                        graph_path=cycle
                    )
                    contradictions.append(evidence)
    except nx.NetworkXError:
        pass # No negative cycles found

    return contradictions

def _detect_resource_conflicts(
    self,
    statements: list[PolicyStatement]
) -> list[ContradictionEvidence]:
    """Detecta conflictos en asignación de recursos"""
    contradictions = []

```

```

resource_allocations = {}

for stmt in statements:
    # Extraer menciones de recursos
    resources = self._extract_resource_mentions(stmt.text)
    for resource_type, amount in resources:
        if resource_type not in resource_allocations:
            resource_allocations[resource_type] = []
        resource_allocations[resource_type].append((stmt, amount))

# Verificar conflictos de asignación
for resource_type, allocations in resource_allocations.items():
    if len(allocations) > 1:
        total_claimed = sum(amount for _, amount in allocations)

        # Verificar si las asignaciones son mutuamente excluyentes
        for i, (stmt_a, amount_a) in enumerate(allocations):
            for stmt_b, amount_b in allocations[i + 1:]:
                if amount_a and amount_b:
                    if self._are_conflicting_allocations(
                        amount_a,
                        amount_b,
                        total_claimed
                    ):
                        confidence = self.bayesian_calculator.calculate_posterior(
                            evidence_strength=get_parameter_loader().get("saaaaaaa.
analysis.contradiction_deteccion.PolicyContradictionDetector._build_knowledge_graph").get(
                            "auto_param_L794_54", 0.8),
                            observations=len(allocations),
                            domain_weight=1.3
                        )

                        evidence = ContradictionEvidence(
                            statement_a=stmt_a,
                            statement_b=stmt_b,

contradiction_type=ContradictionType.RESOURCE_ALLOCATION_MISMATCH,
                            confidence=confidence,
                            severity=get_parameter_loader().get("saaaaaaa.analysis.
contradiction_deteccion.PolicyContradictionDetector._build_knowledge_graph").get("auto_par
am_L804_45", 0.9), # Conflictos de recursos son críticos
                            semantic_similarity=self._calculate_similarity(stmt_a,
stmt_b),
                            logical_conflict_score=get_parameter_loader().get("saa
aaaa.analysis.contradiction_deteccion.PolicyContradictionDetector._build_knowledge_graph")
.get("auto_param_L806_59", 0.8),
                            temporal_consistency=True,
                            numerical_divergence=abs(amount_a - amount_b) /
max(amount_a, amount_b),
                            affected_dimensions=[PolicyDimension.FINANCIERO],
                            resolution_suggestions=self._suggest_resolutions(
                                ContradictionType.RESOURCE_ALLOCATION_MISMATCH
                            )
                        )
                    contradictions.append(evidence)

return contradictions

def _calculate_coherence_metrics(
    self,
    contradictions: list[ContradictionEvidence],
    statements: list[PolicyStatement],
    text: str
) -> dict[str, float]:
    """Calcula métricas avanzadas de coherencia del documento"""

    # Densidad de contradicciones normalizada
    contradiction_density = len(contradictions) / max(1, len(statements))

```

```

# Índice de coherencia semántica global
semantic_coherence = self._calculate_global_semantic_coherence(statements)

# Consistencia temporal
temporal_consistency = sum(
    1 for c in contradictions
    if c.contradiction_type != ContradictionType.TEMPORAL_CONFLICT
) / max(1, len(contradictions))

# Alineación de objetivos
objective_alignment = self._calculate_objective_alignment(statements)

# Índice de fragmentación del grafo
graph_fragmentation = self._calculate_graph_fragmentation()

# Score de coherencia compuesto (weighted harmonic mean)
weights = np.array([get_parameter_loader().get("saaaaaaa.analysis.contradiction_detec-
cion.PolicyContradictionDetector._build_knowledge_graph").get("auto_param_L845_28",
0.3), get_parameter_loader().get("saaaaaaa.analysis.contradiction_deteccion.PolicyContradic-
tionDetector._build_knowledge_graph").get("auto_param_L845_33", 0.25), get_parameter_loade-
r().get("saaaaaaa.analysis.contradiction_deteccion.PolicyContradictionDetector._build_knowl-
edge_graph").get("auto_param_L845_39", 0.2), get_parameter_loader().get("saaaaaaa.analysis.
contradiction_deteccion.PolicyContradictionDetector._build_knowledge_graph").get("auto_par-
am_L845_44", 0.15), get_parameter_loader().get("saaaaaaa.analysis.contradiction_deteccion.P-
olicyContradictionDetector._build_knowledge_graph").get("auto_param_L845_50", 0.1)])]

scores = np.array([
    1 - contradiction_density,
    semantic_coherence,
    temporal_consistency,
    objective_alignment,
    1 - graph_fragmentation
])

# Harmonic mean ponderada para penalizar valores bajos
coherence_score = np.sum(weights) / np.sum(weights) / np.maximum(scores, get_parameter_
loader().get("saaaaaaa.analysis.contradiction_deteccion.PolicyContradictionDetector._bu-
ild_knowledge_graph").get("auto_param_L855_80", 0.01)))

# Entropía de contradicciones
contradiction_entropy = self._calculate_contradiction_entropy(contradictions)

# Complejidad sintáctica del documento
syntactic_complexity = self._calculate_syntactic_complexity(text)

return {
    "coherence_score": float(coherence_score),
    "contradiction_density": float(contradiction_density),
    "semantic_coherence": float(semantic_coherence),
    "temporal_consistency": float(temporal_consistency),
    "objective_alignment": float(objective_alignment),
    "graph_fragmentation": float(graph_fragmentation),
    "contradiction_entropy": float(contradiction_entropy),
    "syntactic_complexity": float(syntactic_complexity),
    "confidence_interval": self._calculate_confidence_interval(coherence_score,
len(statements))
}

def _calculate_global_semantic_coherence(
    self,
    statements: list[PolicyStatement]
) -> float:
    """Calcula coherencia semántica global usando embeddings"""
    if len(statements) < 2:
        return get_parameter_loader().get("saaaaaaa.analysis.contradiction_deteccion.Po-
licyContradictionDetector._build_knowledge_graph").get("auto_param_L881_19", 1.0)

    # Calcular matriz de similitud

```

```

embeddings = [s.embedding for s in statements if s.embedding is not None]
if len(embeddings) < 2:
    return get_parameter_loader().get("saaaaaaa.analysis.contradiction_deteccion.Po
licyContradictionDetector._build_knowledge_graph").get("auto_param_L886_19", 0.5)

similarity_matrix = cosine_similarity(embeddings)

# Calcular coherencia como promedio de similitudes consecutivas
consecutive_similarities = []
for i in range(len(similarity_matrix) - 1):
    consecutive_similarities.append(similarity_matrix[i, i + 1])

# Penalizar alta varianza en similitudes
mean_similarity = np.mean(consecutive_similarities)
std_similarity = np.std(consecutive_similarities)

coherence = mean_similarity * (1 - min(get_parameter_loader().get("saaaaaaa.analysis
s.contradiction_deteccion.PolicyContradictionDetector._build_knowledge_graph").get("auto_p
aram_L899_47", 0.5), std_similarity))

return float(coherence)

def _calculate_objective_alignment(
    self,
    statements: list[PolicyStatement]
) -> float:
    """Calcula alineación entre objetivos declarados"""
    objective_statements = [
        s for s in statements
        if s.semantic_role in ['objective', 'goal', 'target']
    ]

    if len(objective_statements) < 2:
        return get_parameter_loader().get("saaaaaaa.analysis.contradiction_deteccion.Po
licyContradictionDetector._build_knowledge_graph").get("auto_param_L914_19", 1.0)

    # Analizar consistencia direccional de objetivos
    alignment_scores = []
    for i, obj_a in enumerate(objective_statements):
        for obj_b in objective_statements[i + 1:]:
            if obj_a.embedding is not None and obj_b.embedding is not None:
                # Calcular alineación como similitud coseno
                alignment = 1 - cosine(obj_a.embedding, obj_b.embedding)
                alignment_scores.append(alignment)

    if alignment_scores:
        return float(np.mean(alignment_scores))
    return get_parameter_loader().get("saaaaaaa.analysis.contradiction_deteccion.Policy
ContradictionDetector._build_knowledge_graph").get("auto_param_L927_15", 0.5)

@calibrated_method("saaaaaaa.analysis.contradiction_deteccion.PolicyContradictionDetect
or._calculate_graph_fragmentation")
def _calculate_graph_fragmentation(self) -> float:
    """Calcula fragmentación del grafo de conocimiento"""
    if self.knowledge_graph.number_of_nodes() == 0:
        return get_parameter_loader().get("saaaaaaa.analysis.contradiction_deteccion.Po
licyContradictionDetector._calculate_graph_fragmentation").get("auto_param_L933_19", 0.0)

    # Calcular número de componentes conectados
    num_components = nx.number_weakly_connected_components(self.knowledge_graph)
    num_nodes = self.knowledge_graph.number_of_nodes()

    # Fragmentación normalizada
    fragmentation = (num_components - 1) / max(1, num_nodes - 1)

    return float(fragmentation)

def _calculate_contradiction_entropy(

```

```

    self,
    contradictions: list[ContradictionEvidence]
) -> float:
    """Calcula entropía de distribución de tipos de contradicción"""
    if not contradictions:
        return get_parameter_loader().get("saaaaaa.analysis.contradiction_deteccion.Po
licyContradictionDetector._calculate_graph_fragmentation").get("auto_param_L950_19", 0.0)

    # Contar frecuencia de cada tipo
    type_counts = {}
    for c in contradictions:
        type_counts[c.contradiction_type] = type_counts.get(c.contradiction_type, 0) +
1

    # Calcular probabilidades
    total = len(contradictions)
    probabilities = [count / total for count in type_counts.values()]

    # Calcular entropía de Shannon
    entropy = -sum(p * np.log2(p) if p > 0 else 0 for p in probabilities)

    # Normalizar por entropía máxima
    max_entropy = np.log2(len(ContradictionType))
    normalized_entropy = entropy / max_entropy if max_entropy > 0 else 0

    return float(normalized_entropy)

@calibrated_method("saaaaaa.analysis.contradiction_deteccion.PolicyContradictionDetect
or._calculate_syntactic_complexity")
def _calculate_syntactic_complexity(self, text: str) -> float:
    """Calcula complejidad sintáctica del documento"""
    doc = self.nlp(text[:5000]) # Limitar para eficiencia

    # Métricas de complejidad
    avg_sentence_length = np.mean([len(sent.text.split()) for sent in doc.sents])

    # Profundidad promedio del árbol de dependencias
    dependency_depths = []
    for sent in doc.sents:
        depths = [self._get_dependency_depth(token) for token in sent]
        if depths:
            dependency_depths.append(np.mean(depths))

    avg_dependency_depth = np.mean(dependency_depths) if dependency_depths else 0

    # Diversidad léxica (Type-Token Ratio)
    tokens = [token.text.lower() for token in doc if token.is_alpha]
    ttr = len(set(tokens)) / len(tokens) if tokens else 0

    # Combinar métricas
    complexity = (
        min(get_parameter_loader().get("saaaaaa.analysis.contradiction_deteccion.P
olicyContradictionDetector._calculate_syntactic_complexity").get("auto_param_L993_20",
1.0), avg_sentence_length / 50) * get_parameter_loader().get("saaaaaa.analysis.contradicti
on_deteccion.PolicyContradictionDetector._calculate_syntactic_complexity").get("auto_param
_L993_53", 0.3) +
        min(get_parameter_loader().get("saaaaaa.analysis.contradiction_deteccion.P
olicyContradictionDetector._calculate_syntactic_complexity").get("auto_param_L994_20",
1.0), avg_dependency_depth / 10) * get_parameter_loader().get("saaaaaa.analysis.contradicti
on_deteccion.PolicyContradictionDetector._calculate_syntactic_complexity").get("auto_para
m_L994_54", 0.3) +
        ttr * get_parameter_loader().get("saaaaaa.analysis.contradiction_deteccion
.PolicyContradictionDetector._calculate_syntactic_complexity").get("auto_param_L995_22",
0.4)
    )

    return float(complexity)

```

```

@calibrated_method("saaaaaa.analysis.contradiction_deteccion.PolicyContradictionDetect
or._get_dependency_depth")
def _get_dependency_depth(self, token) -> int:
    """Calcula profundidad de un token en el árbol de dependencias"""
    depth = 0
    current = token
    while current.head != current and depth < 20: # Evitar loops infinitos
        current = current.head
        depth += 1
    return depth

def _calculate_confidence_interval(
    self,
    score: float,
    n_observations: int
) -> tuple[float, float]:
    """Calcula intervalo de confianza del 95% para el score"""
    # Usar distribución t de Student para muestras pequeñas
    if n_observations < 30:
        # Error estándar estimado
        se = np.sqrt(score * (1 - score) / n_observations)
        # Valor crítico t para 95% de confianza
        t_critical = stats.t.ppf(get_parameter_loader().get("saaaaaa.analysis.contradi
ction_deteccion.PolicyContradictionDetector._get_dependency_depth").get("auto_param_L1021_
37", 0.975), n_observations - 1)
        margin = t_critical * se
    else:
        # Usar distribución normal para muestras grandes
        se = np.sqrt(score * (1 - score) / n_observations)
        margin = 1.96 * se

    return (
        max(get_parameter_loader().get("saaaaaa.analysis.contradiction_deteccion.Polic
yContradictionDetector._get_dependency_depth").get("auto_param_L1029_16", 0.0), score -
margin),
        min(get_parameter_loader().get("saaaaaa.analysis.contradiction_deteccion.Polic
yContradictionDetector._get_dependency_depth").get("auto_param_L1030_16", 1.0), score +
margin)
    )

def _generate_resolution_recommendations(
    self,
    contradictions: list[ContradictionEvidence]
) -> list[dict[str, Any]]:
    """Genera recomendaciones específicas para resolver contradicciones"""
    recommendations = []

    # Agrupar contradicciones por tipo
    by_type = {}
    for c in contradictions:
        if c.contradiction_type not in by_type:
            by_type[c.contradiction_type] = []
        by_type[c.contradiction_type].append(c)

    # Generar recomendaciones por tipo
    for cont_type, conflicts in by_type.items():
        if cont_type == ContradictionType.NUMERICAL_INCONSISTENCY:
            recommendations.append({
                "type": "numerical_reconciliation",
                "priority": "high",
                "description": "Revisar y reconciliar cifras inconsistentes",
                "specific_actions": [
                    "Verificar fuentes de datos originales",
                    "Establecer línea base única",
                    "Documentar metodología de cálculo"
                ],
                "affected_sections": self._identify_affected_sections(conflicts)
            })

```

```

        elif cont_type == ContradictionType.TEMPORAL_CONFLICT:
            recommendations.append({
                "type": "timeline_adjustment",
                "priority": "high",
                "description": "Ajustar cronograma para resolver conflictos temporales",
                "specific_actions": [
                    "Revisar secuencia de actividades",
                    "Validar plazos con áreas responsables",
                    "Establecer hitos intermedios claros"
                ],
                "affected_sections": self._identify_affected_sections(conflicts)
            })

        elif cont_type == ContradictionType.RESOURCE_ALLOCATION_MISMATCH:
            recommendations.append({
                "type": "budget_reallocation",
                "priority": "critical",
                "description": "Revisar asignación presupuestal",
                "specific_actions": [
                    "Realizar análisis de suficiencia presupuestal",
                    "Priorizar programas según impacto",
                    "Identificar fuentes alternativas de financiación"
                ],
                "affected_sections": self._identify_affected_sections(conflicts)
            })

        elif cont_type == ContradictionType.SEMANTIC_OPPOSITION:
            recommendations.append({
                "type": "conceptual_clarification",
                "priority": "medium",
                "description": "Clarificar conceptos y objetivos opuestos",
                "specific_actions": [
                    "Realizar sesiones de alineación estratégica",
                    "Definir glosario de términos unificado",
                    "Establecer jerarquía clara de objetivos"
                ],
                "affected_sections": self._identify_affected_sections(conflicts)
            })

    # Ordenar por prioridad
    priority_order = {"critical": 0, "high": 1, "medium": 2, "low": 3}
    recommendations.sort(key=lambda x: priority_order.get(x["priority"], 4))

    return recommendations

def _identify_affected_sections(
    self,
    conflicts: list[ContradictionEvidence]
) -> list[str]:
    """Identifica secciones del plan afectadas por contradicciones"""
    affected = set()
    for c in conflicts:
        # Extraer información de sección desde el contexto
        for pattern_name, pattern in self.pdm_patterns.items():
            if pattern.search(c.statement_a.context_window):
                affected.add(pattern_name)
            if pattern.search(c.statement_b.context_window):
                affected.add(pattern_name)

    return list(affected)

def _serialize_contradiction(
    self,
    contradiction: ContradictionEvidence
) -> dict[str, Any]:
    """Serializa evidencia de contradicción para output"""

```

```

return {
    "statement_1": contradiction.statement_a.text,
    "statement_2": contradiction.statement_b.text,
    "position_1": contradiction.statement_a.position,
    "position_2": contradiction.statement_b.position,
    "contradiction_type": contradiction.contradiction_type.name,
    "confidence": float(contradiction.confidence),
    "severity": float(contradiction.severity),
    "semantic_similarity": float(contradiction.semantic_similarity),
    "logical_conflict_score": float(contradiction.logical_conflict_score),
    "temporal_consistency": contradiction.temporal_consistency,
    "numerical_divergence": float(
        contradiction.numerical_divergence) if contradiction.numerical_divergence
    else None,
    "statistical_significance": float(
        contradiction.statistical_significance) if
    contradiction.statistical_significance else None,
}
@calibrated_method("saaaaaaa.analysis.contradiction_deteccion.PolicyContradictionDetect
or._get_graph_statistics")
def _get_graph_statistics(self) -> dict[str, Any]:
    """Obtiene estadísticas del grafo de conocimiento"""
    if self.knowledge_graph.number_of_nodes() == 0:
        return {"nodes": 0, "edges": 0, "components": 0}

    return {
        "nodes": self.knowledge_graph.number_of_nodes(),
        "edges": self.knowledge_graph.number_of_edges(),
        "components": nx.number_weakly_connected_components(self.knowledge_graph),
        "density": nx.density(self.knowledge_graph),
        "average_clustering":
            nx.average_clustering(self.knowledge_graph.to_undirected()),
        "diameter": nx.diameter(self.knowledge_graph.to_undirected()) if
        nx.is_connected(
            self.knowledge_graph.to_undirected()) else -1
    }

# Métodos auxiliares

@calibrated_method("saaaaaaa.analysis.contradiction_deteccion.PolicyContradictionDetect
or._extract_temporal_markers")
def _extract_temporal_markers(self, text: str) -> list[str]:
    """Extrae marcadores temporales del texto"""
    markers = []

    # Patrones de fechas
    date_patterns = [
        r'\d{1,2}\s+de\s+\w+\s+\d{4}', r'\d{4}-\d{2}-\d{2}', r'(enero|febrero|marzo|abril|mayo|junio|julio|agosto|septiembre|octubre|noviem
bre|diciembre)\s+\d{4}', r'(Q[1-4]|trimestre)\s+[1-4])\s+\d{4}', r'20\d{2}', r'(corto|mediano|largo)\s+plazo', r'(primer|segundo|tercer|cuarto)\s+(año|semestre|trimestre)'
    ]

    for pattern in date_patterns:
        matches = re.findall(pattern, text, re.IGNORECASE)
        markers.extend(matches)

    return markers
@calibrated_method("saaaaaaa.analysis.contradiction_deteccion.PolicyContradictionDetect
or._get_temporal_markers")
def _get_temporal_markers(self) -> list[str]:
    """Extrae marcadores temporales del grafo"""
    markers = []

```

```

or._extract_quantitative_claims")
def _extract_quantitative_claims(self, text: str) -> list[dict[str, Any]]:
    """Extrae afirmaciones cuantitativas estructuradas"""
    claims = []

    # Patrones numéricos con contexto
    patterns = [
        (r'(\d+(:[.]|\d+)?)(\s*(%|por\s*ciento)', 'percentage'),
        (r'(\d+(:[.]|\d+)?)(\s*(millones?|mil\s+millones?|)', 'amount'),
        (r'(\$|COP)\s*(\d+(:[.]|\d+)?|', 'currency'),
        (r'(\d+(:[.]|\d+)?)(\s*(personas?|beneficiarios?|familias?|',
        'beneficiaries'),
        (r'(\d+(:[.]|\d+)?)(\s*(hectáreas?|km2?|metros?|', 'area'),
        (r'meta\s+de\s+(\d+(:[.]|\d+)?|', 'target')
    ]

    for pattern, claim_type in patterns:
        matches = re.finditer(pattern, text, re.IGNORECASE)
        for match in matches:
            value_str = match.group(1) if claim_type != 'currency' else match.group(2)
            value = self._parse_number(value_str)

            claims.append({
                'type': claim_type,
                'value': value,
                'raw_text': match.group(0),
                'position': match.span(),
                'context': text[max(0, match.start() - 20):min(len(text), match.end()
+ 20)]
            })
    }

    return claims

@calibrated_method("saaaaaaa.analysis.contradiction_deteccion.PolicyContradictionDetect
or._parse_number")
def _parse_number(self, text: str) -> float:
    """Parsea número desde texto"""
    try:
        # Reemplazar coma decimal
        normalized = text.replace(',', '.')
        return float(normalized)
    except ValueError:
        return get_parameter_loader().get("saaaaaaa.analysis.contradiction_deteccion.Po
licyContradictionDetector._parse_number").get("auto_param_L1227_19", 0.0)

@calibrated_method("saaaaaaa.analysis.contradiction_deteccion.PolicyContradictionDetect
or._extract_resource_mentions")
def _extract_resource_mentions(self, text: str) -> list[tuple[str, float | None]]:
    """Extrae menciones de recursos con montos"""
    resources = []

    # Patrones de recursos específicos de PDM colombiano
    resource_patterns = [
        (r'SGP(\s*[:\s]*\$?\s*(\d+(:[.]|\d+)?))\s*(millones?)?', 'SGP'),
        (r'regalías(\s*[:\s]*\$?\s*(\d+(:[.]|\d+)?))\s*(millones?)?', 'regalías'),
        (r'recursos\s+propios(\s*[:\s]*\$?\s*(\d+(:[.]|\d+)?))\s*(millones?)?',
        'recursos_propios'),
        (r'cofinanciación(\s*[:\s]*\$?\s*(\d+(:[.]|\d+)?))\s*(millones?)?',
        'cofinanciación'),
        (r'presupuesto\s+total(\s*[:\s]*\$?\s*(\d+(:[.]|\d+)?))\s*(millones?)?',
        'presupuesto_total')
    ]

    for pattern, resource_type in resource_patterns:
        matches = re.finditer(pattern, text, re.IGNORECASE)
        for match in matches:
            amount = self._parse_number(match.group(1)) if match.group(1) else None
            if match.group(2) and 'millon' in match.group(2).lower():

```

```

amount = amount * 1000000 if amount else None
resources.append((resource_type, amount))

return resources

@calibrated_method("saaaaaa.analysis.contradiction_deteccion.PolicyContradictionDetect
or._determine_semantic_role")
def _determine_semantic_role(self, sent) -> str | None:
    """Determina el rol semántico de una oración"""
    # Safely extract text (handles both strings and spacy objects)
    text_lower = safe_text_extract(sent).lower()

role_patterns = {
    'objective': ['objetivo', 'meta', 'propósito', 'finalidad'],
    'strategy': ['estrategia', 'línea', 'eje', 'pilar'],
    'action': ['implementar', 'ejecutar', 'desarrollar', 'realizar'],
    'indicator': ['índicador', 'medir', 'evaluar', 'monitorear'],
    'resource': ['presupuesto', 'recurso', 'financiación', 'inversión'],
    'constraint': ['limitación', 'restricción', 'condición', 'requisito']
}

for role, keywords in role_patterns.items():
    if any(keyword in text_lower for keyword in keywords):
        return role

return None

@calibrated_method("saaaaaa.analysis.contradiction_deteccion.PolicyContradictionDetect
or._identify_dependencies")
def _identify_dependencies(self, sent, doc) -> set[str]:
    """Identifica dependencias entre declaraciones"""
    dependencies = set()

    # Buscar referencias a otras secciones
    reference_patterns = [
        r'como\s+se\s+menciona\s+en',
        r'según\s+lo\s+establecido\s+en',
        r'de\s+acuerdo\s+con',
        r'en\s+línea\s+con',
        r'siguiendo\s+lo\s+dispuesto'
    ]

    for pattern in reference_patterns:
        if re.search(pattern, sent.text, re.IGNORECASE):
            # Buscar la sección referenciada
            for other_sent in doc.sents:
                if other_sent != sent:
                    # Usar hash de los primeros 50 caracteres como ID
                    dependencies.add(other_sent.text[:50])

    return dependencies

@calibrated_method("saaaaaa.analysis.contradiction_deteccion.PolicyContradictionDetect
or._get_context_window")
def _get_context_window(self, text: str, start: int, end: int, window_size: int = 200)
-> str:
    """Obtiene ventana de contexto alrededor de una posición"""
    context_start = max(0, start - window_size)
    context_end = min(len(text), end + window_size)
    return text[context_start:context_end]

@calibrated_method("saaaaaa.analysis.contradiction_deteccion.PolicyContradictionDetect
or._calculate_similarity")
def _calculate_similarity(self, stmt_a: PolicyStatement, stmt_b: PolicyStatement) ->
float:
    """Calcula similaridad entre dos declaraciones"""
    if stmt_a.embedding is not None and stmt_b.embedding is not None:
        return float(1 - cosine(stmt_a.embedding, stmt_b.embedding))

```

```

        return get_parameter_loader().get("saaaaaaa.analysis.contradiction_deteccion.Policy
ContradictionDetector._calculate_similarity").get("auto_param_L1310_15", 0.0)

    @calibrated_method("saaaaaaa.analysis.contradiction_deteccion.PolicyContradictionDetect
or._classify_contradiction")
def _classify_contradiction(self, text: str) -> float:
    """Clasifica probabilidad de contradicción en texto"""
    try:
        result = self.contradiction_classifier(text)
        # Buscar score de contradicción
        for item in result:
            if 'contradiction' in item['label'].lower():
                return item['score']
    return get_parameter_loader().get("saaaaaaa.analysis.contradiction_deteccion.Po
licyContradictionDetector._classify_contradiction").get("auto_param_L1321_19", 0.0)
    except Exception as e:
        logger.warning(f"Error en clasificación de contradicción: {e}")
        return get_parameter_loader().get("saaaaaaa.analysis.contradiction_deteccion.Po
licyContradictionDetector._classify_contradiction").get("auto_param_L1324_19", 0.0)

    @calibrated_method("saaaaaaa.analysis.contradiction_deteccion.PolicyContradictionDetect
or._get_domain_weight")
def _get_domain_weight(self, dimension: PolicyDimension) -> float:
    """Obtiene peso específico del dominio"""
    weights = {
        PolicyDimension.DIAGNOSTICO: get_parameter_loader().get("saaaaaaa.analysis.con
tradiction_deteccion.PolicyContradictionDetector._get_domain_weight").get("auto_param_L1330
_41", 0.8),
        PolicyDimension.ESTRATEGICO: 1.2,
        PolicyDimension.PROGRAMATICO: get_parameter_loader().get("saaaaaaa.analysis.con
tradiction_deteccion.PolicyContradictionDetector._get_domain_weight").get("auto_param_L133
2_42", 1.0),
        PolicyDimension.FINANCIERO: 1.5,
        PolicyDimension.SEGUIMIENTO: get_parameter_loader().get("saaaaaaa.analysis.con
tradiction_deteccion.PolicyContradictionDetector._get_domain_weight").get("auto_param_L1334
_41", 0.9),
        PolicyDimension.TERRITORIAL: 1.1
    }
    return weights.get(dimension, get_parameter_loader().get("saaaaaaa.analysis.contradic
tion_deteccion.PolicyContradictionDetector._get_domain_weight").get("auto_param_L1337_38
", 1.0))

    @calibrated_method("saaaaaaa.analysis.contradiction_deteccion.PolicyContradictionDetect
or._suggest_resolutions")
def _suggest_resolutions(self, contradiction_type: ContradictionType) -> list[str]:
    """Sugiere resoluciones específicas por tipo de contradicción"""
    suggestions = {
        ContradictionType.NUMERICAL_INCONSISTENCY: [
            "Verificar fuentes de datos y metodologías de cálculo",
            "Establecer línea base única con validación técnica",
            "Documentar supuestos y proyecciones utilizadas"
        ],
        ContradictionType.TEMPORAL_CONFLICT: [
            "Revisar cronograma maestro del plan",
            "Validar secuencia lógica de actividades",
            "Ajustar plazos según capacidad institucional"
        ],
        ContradictionType.SEMANTIC_OPPOSITION: [
            "Realizar taller de alineación conceptual",
            "Clarificar definiciones en glosario técnico",
            "Priorizar objetivos según Plan Nacional de Desarrollo"
        ],
        ContradictionType.RESOURCE_ALLOCATION_MISMATCH: [
            "Realizar análisis de brechas financieras",
            "Priorizar inversiones según impacto social",
            "Explorar fuentes alternativas de financiación"
        ],
        ContradictionType.LOGICAL_INCOMPATIBILITY: [

```

```

        "Revisar cadena de valor de programas",
        "Validar teoría de cambio del plan",
        "Eliminar duplicidades y solapamientos"
    ]
}
return suggestions.get(contradiction_type, ["Revisar y ajustar según contexto"])

@calibrated_method("saaaaaaa.analysis.contradiction_deteccion.PolicyContradictionDetect
or._are_comparable_claims")
def _are_comparable_claims(self, claim_a: dict, claim_b: dict) -> bool:
    """Determina si dos afirmaciones cuantitativas son comparables"""
    # Mismo tipo y contexto similar
    if claim_a['type'] != claim_b['type']:
        return False

    # Verificar si hablan del mismo concepto
    context_similarity = self._text_similarity(
        claim_a.get('context', ''),
        claim_b.get('context', '')
    )

    return context_similarity > get_parameter_loader().get("saaaaaaa.analysis.contradic
tion_deteccion.PolicyContradictionDetector._are_comparable_claims").get("auto_param_L1384_
36", 0.6)

@calibrated_method("saaaaaaa.analysis.contradiction_deteccion.PolicyContradictionDetect
or._text_similarity")
def _text_similarity(self, text_a: str, text_b: str) -> float:
    """Calcula similaridad simple entre textos"""
    if not text_a or not text_b:
        return get_parameter_loader().get("saaaaaaa.analysis.contradiction_deteccion.Po
licyContradictionDetector._text_similarity").get("auto_param_L1390_19", 0.0)

    # Tokenización simple
    tokens_a = set(text_a.lower().split())
    tokens_b = set(text_b.lower().split())

    if not tokens_a or not tokens_b:
        return get_parameter_loader().get("saaaaaaa.analysis.contradiction_deteccion.Po
licyContradictionDetector._text_similarity").get("auto_param_L1397_19", 0.0)

    # Coeficiente de Jaccard
    intersection = tokens_a & tokens_b
    union = tokens_a | tokens_b

    return len(intersection) / len(union) if union else get_parameter_loader().get("sa
aaaaaa.analysis.contradiction_deteccion.PolicyContradictionDetector._text_similarity").get(
"auto_param_L1403_60", 0.0)

def _calculate_numerical_divergence(
    self,
    claim_a: dict,
    claim_b: dict
) -> float | None:
    """Calcula divergencia entre valores numéricos"""
    value_a = claim_a.get('value', 0)
    value_b = claim_b.get('value', 0)

    if value_a == 0 and value_b == 0:
        return None

    # Divergencia relativa
    max_value = max(abs(value_a), abs(value_b))
    if max_value == 0:
        return None

    divergence = abs(value_a - value_b) / max_value
    return divergence

```

```

def _statistical_significance_test(
    self,
    claim_a: dict,
    claim_b: dict
) -> float:
    """Realiza test de significancia estadística"""
    value_a = claim_a.get('value', 0)
    value_b = claim_b.get('value', 0)

    # Test t de una muestra para diferencia significativa
    # Asumiendo distribución normal con varianza estimada
    diff = abs(value_a - value_b)
    pooled_value = (value_a + value_b) / 2

    if pooled_value == 0:
        return get_parameter_loader().get("saaaaaaa.analysis.contradiction_deteccion.Po
licyContradictionDetector._text_similarity").get("auto_param_L1440_19", 1.0) # No
significativo

    # Estimación conservadora de error estándar
    se = pooled_value * get_parameter_loader().get("saaaaaaa.analysis.contradiction_det
eccion.PolicyContradictionDetector._text_similarity").get("auto_param_L1443_28", 0.1) #
10% de error estimado

    if se == 0:
        return get_parameter_loader().get("saaaaaaa.analysis.contradiction_deteccion.Po
licyContradictionDetector._text_similarity").get("auto_param_L1446_19", 0.0) # Altamente
significativo

    # Estadístico t
    t_stat = diff / se

    # Valor p aproximado (two-tailed)
    p_value = 2 * (1 - stats.norm.cdf(abs(t_stat)))

    return p_value

@calibrated_method("saaaaaaa.analysis.contradiction_deteccion.PolicyContradictionDetect
or._has_logical_conflict")
def _has_logical_conflict(self, stmt_a: PolicyStatement, stmt_b: PolicyStatement) ->
bool:
    """Determina si hay conflicto lógico entre declaraciones"""
    # Verificar si las declaraciones tienen roles incompatibles
    if stmt_a.semantic_role and stmt_b.semantic_role:
        incompatible_roles = [
            ('objective', 'constraint'),
            ('strategy', 'constraint'),
            ('action', 'constraint')
        ]

        for role_pair in incompatible_roles:
            if (stmt_a.semantic_role in role_pair and
                stmt_b.semantic_role in role_pair and
                stmt_a.semantic_role != stmt_b.semantic_role):
                return True

    # Verificar negación explícita
    negation_patterns = ['no', 'nunca', 'ningún', 'sin', 'tampoco']
    has_negation_a = any(pattern in stmt_a.text.lower() for pattern in
negation_patterns)
    has_negation_b = any(pattern in stmt_b.text.lower() for pattern in
negation_patterns)

    # Si una tiene negación y otra no, y son similares, hay conflicto
    if has_negation_a != has_negation_b:
        similarity = self._calculate_similarity(stmt_a, stmt_b)
        if similarity > get_parameter_loader().get("saaaaaaa.analysis.contradiction_det

```

```

eccion.PolicyContradictionDetector._has_logical_conflict").get("auto_param_L1481_28",
0.7):
    return True

    return False

def _are_conflicting_allocations(
    self,
    amount_a: float,
    amount_b: float,
    total: float
) -> bool:
    """Determina si las asignaciones de recursos están en conflicto"""
    # Si la suma excede el total disponible
    if amount_a + amount_b > total * 1.1: # 10% de margen
        return True

    # Si hay una diferencia muy grande entre asignaciones similares
    return abs(amount_a - amount_b) / max(amount_a, amount_b) > get_parameter_loader()
.get("saaaaaaa.analysis.contradiction_deteccion.PolicyContradictionDetector._has_logical_co
nflict").get("auto_param_L1498_68", 0.5)

def _determine_relation_type(
    self,
    stmt_a: PolicyStatement,
    stmt_b: PolicyStatement
) -> str:
    """Determina el tipo de relación entre dos declaraciones"""
    # Analizar roles semánticos
    if stmt_a.semantic_role and stmt_b.semantic_role:
        if stmt_a.semantic_role == stmt_b.semantic_role:
            return "parallel"
        elif stmt_a.semantic_role in ["strategy", "objective"] and
stmt_b.semantic_role == "action":
            return "enables"
        elif stmt_a.semantic_role == "action" and stmt_b.semantic_role in
["indicator", "resource"]:
            return "requires"

    # Analizar dependencias
    if stmt_a.dependencies & {stmt_b.text[:50]}:
        return "depends_on"

    # Por defecto, relación de similaridad
    return "related"

def _calculate_severity(
    self,
    stmt_a: PolicyStatement,
    stmt_b: PolicyStatement
) -> float:
    """Calcula la severidad de una contradicción entre declaraciones"""
    severity = get_parameter_loader().get("saaaaaaa.analysis.contradiction_deteccion.Po
licyContradictionDetector._has_logical_conflict").get("severity", 0.5) # Refactored

    # Incrementar si las declaraciones están en la misma dimensión
    if stmt_a.dimension == stmt_b.dimension:
        severity += get_parameter_loader().get("saaaaaaa.analysis.contradiction_detecci
on.PolicyContradictionDetector._has_logical_conflict").get("auto_param_L1532_24", 0.2)

    # Incrementar si tienen muchas entidades en común
    common_entities = set(stmt_a.entities) & set(stmt_b.entities)
    if len(common_entities) > 0:
        severity += min(get_parameter_loader().get("saaaaaaa.analysis.contradiction_det
eccion.PolicyContradictionDetector._has_logical_conflict").get("auto_param_L1537_28",
0.2), len(common_entities) * get_parameter_loader().get("saaaaaaa.analysis.contradiction_de
teccion.PolicyContradictionDetector._has_logical_conflict").get("auto_param_L1537_56",
0.05))

```

```

# Incrementar si tienen marcadores temporales en conflicto
if stmt_a.temporal_markers and stmt_b.temporal_markers:
    severity += get_parameter_loader().get("saaaaaaa.analysis.contradiction_detection.PolicyContradictionDetector._has_logical_conflict").get("auto_param_L1541_24", 0.1)

    return min(get_parameter_loader().get("saaaaaaa.analysis.contradiction_deteccion.PolicyContradictionDetector._has_logical_conflict").get("auto_param_L1543_19", 1.0),
               severity)

```

===== FILE: src/saaaaaaa/analysis/derek\_beach.py =====

#!/usr/bin/env python3

"""

Causal Deconstruction and Audit Framework (CDAF) v2.0

Framework de Producción para Análisis Causal de Planes de Desarrollo Territorial

THEORETICAL FOUNDATION (Derek Beach):

"A causal mechanism is a system of interlocking parts (entities engaging in activities) that transmits causal forces from X to Y" (Beach 2016: 465)

This framework implements Theory-Testing Process Tracing with mechanistic evidence evaluation using Beach's evidential tests taxonomy (Beach & Pedersen 2019).

Author: AI Systems Architect

Version: 2.0.0 (Beach-Grounded Production Grade)

"""

```

import argparse
import hashlib
import json
import logging
import re
import sys
import warnings
from collections import defaultdict
from dataclasses import asdict, dataclass, field
from pathlib import Path
from typing import (
    from_saaaaaa import get_parameter_loader
    from saaaaaaa.core.calibration.decorators import calibrated_method
    TYPE_CHECKING,
    Any,
    Literal,
    NamedTuple,
    TypedDict,
    cast,
)
if TYPE_CHECKING:
    import fitz
# Core dependencies
try:
    import networkx as nx
    import numpy as np
    import pandas as pd
    import spacy
    import yaml
    from fuzzywuzzy import fuzz, process
    from pydantic import BaseModel, Field, ValidationError, validator
    from pydot import Dot, Edge, Node
    from scipy import stats
    from scipy.spatial.distance import cosine
    from scipy.special import rel_entr
except ImportError as e:
    print(f"ERROR: Dependencia faltante. Ejecute: pip install {e.name}")
    sys.exit(1)

```

```

# DNP Standards Integration
try:
    from dnp_integration import ValidadorDNP, validar_plan_desarrollo_completo

    DNP_AVAILABLE = True
except ImportError:
    DNP_AVAILABLE = False
    warnings.warn("Módulos DNP no disponibles. Validación DNP deshabilitada.", stacklevel=2)

# Refactored Bayesian Engine (F1.2: Architectural Refactoring)
try:
    from inference.bayesian_adapter import BayesianEngineAdapter

    REFACTORED_BAYESIAN_AVAILABLE = True
except ImportError:
    REFACTORED_BAYESIAN_AVAILABLE = False
    warnings.warn("Motor Bayesiano refactorizado no disponible. Usando implementación legacy.", stacklevel=2)

# Configure logging
logging.basicConfig(
    format='%(asctime)s - %(name)s - %(levelname)s - %(message)s',
    datefmt='%Y-%m-%d %H:%M:%S'
)

# =====
# CONSTANTS
# =====
DEFAULT_CONFIG_FILE = "config.yaml"
EXTRACTION_REPORT_SUFFIX = "_extraction_confidence_report.json"
CAUSAL_MODEL_SUFFIX = "_causal_model.json"
DNP_REPORT_SUFFIX = "_dnp_compliance_report.txt"

# Type definitions
NodeType = Literal["programa", "producto", "resultado", "impacto"]
RigorStatus = Literal["fuerte", "débil", "sin_evaluar"]
TestType = Literal["hoop_test", "smoking_gun", "doubly_decisional", "straw_in_wind"]
DynamicsType = Literal["suma", "decreciente", "constante", "indefinido"]

# =====
# BEACH THEORETICAL PRIMITIVES - Added to existing code
# =====

```

class BeachEvidentialTest:

"""

Derek Beach evidential tests implementation (Beach & Pedersen 2019: Ch 5).

FOUR-FOLD TYPOLOGY calibrated by necessity (N) and sufficiency (S):

HOOP TEST [N: High, S: Low]:

- Fail → ELIMINATES hypothesis (definitive knock-out)
- Pass → Hypothesis survives but not proven
- Example: "Responsible entity must be documented"

SMOKING GUN [N: Low, S: High]:

- Pass → Strongly confirms hypothesis
- Fail → Doesn't eliminate (could be false negative)
- Example: "Unique policy instrument only used for this mechanism"

DOUBLY DECISIVE [N: High, S: High]:

- Pass → Conclusively confirms
- Fail → Conclusively eliminates
- Extremely rare in social science

STRAW-IN-WIND [N: Low, S: Low]:

- Pass/Fail → Marginal confidence change
- Used for preliminary screening

REFERENCE: Beach & Pedersen (2019), pp 117-126

"""

```
@staticmethod  
def classify_test(necessity: float, sufficiency: float) -> TestType:  
    """  
        Classify evidential test type based on necessity and sufficiency.  
    """
```

Beach calibration:

- Necessity > 0.7 → High necessity
- Sufficiency > 0.7 → High sufficiency

"""

```
high_n = necessity > 0.7  
high_s = sufficiency > 0.7
```

```
if high_n and high_s:  
    return "doubly_decisional"  
elif high_n and not high_s:  
    return "hoop_test"  
elif not high_n and high_s:  
    return "smoking_gun"  
else:  
    return "straw_in_wind"
```

```
@staticmethod  
def apply_test_logic(test_type: TestType, evidence_found: bool,  
                     prior: float, bayes_factor: float) -> tuple[float, str]:  
    """
```

Apply Beach test-specific logic to Bayesian updating.

CRITICAL RULES:

1. Hoop Test FAIL → posterior ≈ 0 (knock-out)
2. Smoking Gun PASS → multiply prior by large BF (>10)
3. Doubly Decisive → extreme updates (BF > 100 or < 0.01)

Returns: (posterior\_confidence, interpretation)

"""

```
if test_type == "hoop_test":  
    if not evidence_found:  
        # KNOCK-OUT per Beach: "hypothesis must jump through hoop"  
        return 0.01, "HOOP_TEST_FAILURE: Hypothesis eliminated"  
    else:  
        # Pass: necessary condition met, use standard Bayesian  
        posterior = min(0.95, prior * bayes_factor)  
        return posterior, "HOOP_TEST_PASSED: Hypothesis survives, not proven"  
  
elif test_type == "smoking_gun":  
    if evidence_found:  
        # Strong confirmation: unique evidence found  
        posterior = min(0.98, prior * max(bayes_factor, 10.0))  
        return posterior, "SMOKING_GUN_FOUND: Strong confirmation"  
    else:  
        # Doesn't eliminate: could be false negative  
        posterior = prior * 0.9 # slight penalty  
        return posterior, "SMOKING_GUN_NOT_FOUND: Doesn't eliminate"  
  
elif test_type == "doubly_decisional":  
    if evidence_found:  
        return 0.99, "DOUBLY_DECISIVE_CONFIRMED: Conclusive"  
    else:  
        return 0.01, "DOUBLY_DECISIVE_ELIMINATED: Conclusive"  
  
# Marginal update only  
elif evidence_found:  
    posterior = min(0.95, prior * min(bayes_factor, 2.0))  
    return posterior, "STRAW_IN_WIND: Weak support"  
else:
```

```

posterior = max(0.05, prior / min(bayes_factor, 2.0))
return posterior, "STRAW_IN_WIND: Weak disconfirmation"

# =====
# Custom Exceptions - Structured Error Semantics
# =====

class CDAFException(Exception):
    """Base exception for CDAF framework with structured payloads"""

    def __init__(self, message: str, details: dict[str, Any] | None = None,
                 stage: str | None = None, recoverable: bool = False) -> None:
        self.message = message
        self.details = details or {}
        self.stage = stage
        self.recoverable = recoverable
        super().__init__(self._format_message())

    @calibrated_method("saaaaaa.analysis.derek_beach.CDAFException._format_message")
    def _format_message(self) -> str:
        """Format error message with structured information"""
        parts = "[CDAF Error]"
        if self.stage:
            parts.append(f"[Stage: {self.stage}]")
        parts.append(self.message)
        if self.details:
            parts.append(f"Details: {json.dumps(self.details, indent=2)}")
        return " ".join(parts)

    @calibrated_method("saaaaaa.analysis.derek_beach.CDAFException.to_dict")
    def to_dict(self) -> dict[str, Any]:
        """Convert exception to structured dictionary"""
        return {
            'error_type': self.__class__.__name__,
            'message': self.message,
            'details': self.details,
            'stage': self.stage,
            'recoverable': self.recoverable
        }

class CDAFValidationError(CDAFException):
    """Configuration or data validation error"""
    pass

class CDAFProcessingError(CDAFException):
    """Error during document processing"""
    pass

class CDAFBayesianError(CDAFException):
    """Error during Bayesian inference"""
    pass

class CDAFConfigError(CDAFException):
    """Configuration loading or validation error"""
    pass

# =====
# Pydantic Configuration Models - Schema Validation at Load Time
# =====

class BayesianThresholdsConfig(BaseModel):
    """Bayesian inference thresholds configuration"""
    kl_divergence: float = Field(
        default=get_parameter_loader().get("saaaaaa.analysis.derek_beach.CDAFException.to_dict").get("auto_param_L255_16", 0.01),
        ge=get_parameter_loader().get("saaaaaa.analysis.derek_beach.CDAFException.to_dict").get("auto_param_L256_11", 0.0),
        le=get_parameter_loader().get("saaaaaa.analysis.derek_beach.CDAFException.to_dict")
    )

```

```

).get("auto_param_L257_11", 1.0),
    description="KL divergence threshold for convergence"
)
convergence_min_evidence: int = Field(
    default=2,
    ge=1,
    description="Minimum evidence count for convergence check"
)
prior_alpha: float = Field(
    default=2.0,
    ge=get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CDAFException.to_dict")
).get("auto_param_L267_11", 0.1),
    description="Default alpha parameter for Beta prior"
)
prior_beta: float = Field(
    default=2.0,
    ge=get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CDAFException.to_dict")
).get("auto_param_L272_11", 0.1),
    description="Default beta parameter for Beta prior"
)
laplace_smoothing: float = Field(
    default=get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CDAFException.to_dict").get("auto_param_L276_16", 1.0),
    ge=get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CDAFException.to_dict")
).get("auto_param_L277_11", 0.0),
    description="Laplace smoothing parameter"
)

class MechanismTypeConfig(BaseModel):
    """Mechanism type prior probabilities"""
    administrativo: float = Field(default=get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CDAFException.to_dict").get("auto_param_L283_42", 0.30), ge=get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CDAFException.to_dict").get("auto_param_L283_51", 0.0), le=get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CDAFException.to_dict").get("auto_param_L283_59", 1.0))
    tecnico: float = Field(default=get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CDAFException.to_dict").get("auto_param_L284_35", 0.25), ge=get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CDAFException.to_dict").get("auto_param_L284_44", 0.0), le=ge_t_parameter_loader().get("saaaaaaa.analysis.derek_beach.CDAFException.to_dict").get("auto_param_L284_52", 1.0))
    financiero: float = Field(default=get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CDAFException.to_dict").get("auto_param_L285_38", 0.20), ge=get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CDAFException.to_dict").get("auto_param_L285_47", 0.0), le=get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CDAFException.to_dict").get("auto_param_L285_55", 1.0))
    politico: float = Field(default=get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CDAFException.to_dict").get("auto_param_L286_36", 0.15), ge=get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CDAFException.to_dict").get("auto_param_L286_45", 0.0), le=g_et_parameter_loader().get("saaaaaaa.analysis.derek_beach.CDAFException.to_dict").get("auto_param_L286_53", 1.0))
    mixto: float = Field(default=get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CDAFException.to_dict").get("auto_param_L287_33", 0.10), ge=get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CDAFException.to_dict").get("auto_param_L287_42", 0.0), le=get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CDAFException.to_dict").get("auto_param_L287_50", 1.0))

    @validator('*', pre=True, always=True)
    def check_sum_to_one(cls, v, values):
        """Validate that probabilities sum to approximately get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CDAFException.to_dict").get("auto_param_L291_60", 1.0)"""
        if len(values) == 4: # All fields loaded
            total = sum(values.values()) + v
            if abs(total - get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CDAFException.to_dict").get("auto_param_L294_27", 1.0)) > get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CDAFException.to_dict").get("auto_param_L294_34", 0.01):
                raise ValueError(f"Mechanism type priors must sum to get_parameter_loader().get('saaaaaaa.analysis.derek_beach.CDAFException.to_dict').get('auto_param_L295_69', 1.0), got {total}")

```

```

    return v

class PerformanceConfig(BaseModel):
    """Performance and optimization settings"""
    enable_vectorized_ops: bool = Field(
        default=True,
        description="Use vectorized numpy operations where possible"
    )
    enable_async_processing: bool = Field(
        default=False,
        description="Enable async processing for large PDFs (experimental)"
    )
    max_context_length: int = Field(
        default=1000,
        ge=100,
        description="Maximum context length for spaCy processing"
    )
    cache_embeddings: bool = Field(
        default=True,
        description="Cache spaCy embeddings for reuse"
    )

class SelfReflectionConfig(BaseModel):
    """Self-reflective learning configuration"""
    enable_prior_learning: bool = Field(
        default=False,
        description="Enable learning from audit feedback to update priors"
    )
    feedback_weight: float = Field(
        default=get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CDAFException.to_dict").get("auto_param_L325_16", 0.1),
        ge=get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CDAFException.to_dict").get("auto_param_L326_11", 0.0),
        le=get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CDAFException.to_dict").get("auto_param_L327_11", 1.0),
        description="Weight for feedback in prior updates (0=ignore, 1=full)"
    )
    prior_history_path: str | None = Field(
        default=None,
        description="Path to save/load historical priors"
    )
    min_documents_for_learning: int = Field(
        default=5,
        ge=1,
        description="Minimum documents before applying learned priors"
    )

class CDAFConfigSchema(BaseModel):
    """Complete CDAF configuration schema with validation"""
    patterns: dict[str, str] = Field(
        description="Regex patterns for document parsing"
    )
    lexicons: dict[str, Any] = Field(
        description="Lexicons for causal logic, classification, etc."
    )
    entity_aliases: dict[str, str] = Field(
        description="Entity name aliases and mappings"
    )
    verb_sequences: dict[str, int] = Field(
        description="Verb sequence ordering for temporal coherence"
    )
    bayesian_thresholds: BayesianThresholdsConfig = Field(
        default_factory=BayesianThresholdsConfig,
        description="Bayesian inference thresholds"
    )
    mechanism_type_priors: MechanismTypeConfig = Field(
        default_factory=MechanismTypeConfig,
        description="Prior probabilities for mechanism types"
    )

```

```

)
performance: PerformanceConfig = Field(
    default_factory=PerformanceConfig,
    description="Performance and optimization settings"
)
self_reflection: SelfReflectionConfig = Field(
    default_factory=SelfReflectionConfig,
    description="Self-reflective learning configuration"
)

```

```

class Config:
    extra = 'allow' # Allow additional fields for extensibility

```

```

class GoalClassification(NamedTuple):
    """
    Classification structure for goals
    """

```

```

    type: NodeType
    dynamics: DynamicsType
    test_type: TestType
    confidence: float

```

```

class EntityActivity(NamedTuple):
    """
    Entity-Activity tuple for mechanism parts (Beach 2016).
    """

```

**BEACH DEFINITION:**

"A mechanism part consists of an entity (organization, actor, structure) engaging in an activity that transmits causal forces" (Beach 2016: 465)

This is the FUNDAMENTAL UNIT of mechanistic evidence in Process Tracing.

"""

```

entity: str
activity: str
verb_lemma: str
confidence: float

```

```

class CausalLink(TypedDict):
    """
    Structure for causal links in the graph
    """

```

```

    source: str
    target: str
    logic: str
    strength: float
    evidence: list[str]
    posterior_mean: float | None
    posterior_std: float | None
    kl_divergence: float | None
    converged: bool | None

```

```

class AuditResult(TypedDict):
    """
    Audit result structure
    """

```

```

    passed: bool
    warnings: list[str]
    errors: list[str]
    recommendations: list[str]

```

@dataclass

```

class MetaNode:
    """
    Comprehensive node structure for goals/metas
    """

```

```

    id: str
    text: str
    type: NodeType
    baseline: float | str | None = None
    target: float | str | None = None
    unit: str | None = None
    responsible_entity: str | None = None
    entity_activity: EntityActivity | None = None
    financial_allocation: float | None = None
    unit_cost: float | None = None
    rigor_status: RigorStatus = "sin_evaluar"

```

```

dynamics: DynamicsType = "indefinido"
test_type: TestType = "straw_in_wind"
contextual_risks: list[str] = field(default_factory=list)
causal_justification: list[str] = field(default_factory=list)
audit_flags: list[str] = field(default_factory=list)
confidence_score: float = get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CDA
FException.to_dict").get("auto_param_L434_30", 0.0)

class ConfigLoader:
    """External configuration management with Pydantic schema validation"""

    def __init__(self, config_path: Path) -> None:
        self.logger = logging.getLogger(self.__class__.__name__)
        self.config_path = config_path
        self.config: dict[str, Any] = {}
        self.validated_config: CDAFConfigSchema | None = None
        # HARMONIC FRONT 4: Track uncertainty over iterations
        self._uncertainty_history: list[float] = []
        self._load_config()
        self._validate_config()
        self._load_uncertainty_history()

    @calibrated_method("saaaaaaa.analysis.derek_beach.ConfigLoader._load_config")
    def _load_config(self) -> None:
        """Load YAML configuration file"""
        # Delegate to factory for I/O operation
        from .factory import load_yaml

        try:
            self.config = load_yaml(self.config_path)
            self.logger.info(f"Configuración cargada desde {self.config_path}")
        except FileNotFoundError:
            self.logger.warning(f"Archivo de configuración no encontrado:
{self.config_path}")
            self._load_default_config()
        except Exception as e:
            raise CDAFConfigError(
                "Error cargando configuración",
                details={'path': str(self.config_path), 'error': str(e)},
                stage="config_load",
                recoverable=True
            )

    @calibrated_method("saaaaaaa.analysis.derek_beach.ConfigLoader._load_default_config")
    def _load_default_config(self) -> None:
        """Load default configuration if custom fails"""
        self.config = {
            'patterns': {
                'section_titles': r'^(?:CAPÍTULO|ARTÍCULO|PARTE)\s+[\dIVX]+',
                'goal_codes': r'[MP][RIP]-\d{3}',
                'numeric_formats': r'[\d.]+(?:\.\d+)?%?',
                'table_headers': r'(?:PROGRAMA|META|INDICADOR|LÍNEA BASE|VALOR ESPERADO)',
                'financial_headers': r'(?:PRESUPUESTO|VALOR|MONTO|INVERSIÓN)'
            },
            'lexicons': {
                'causal_logic': [
                    'gracias a', 'con el fin de', 'para lograr', 'mediante',
                    'a través de', 'como resultado de', 'debido a', 'porque',
                    'por medio de', 'permitirá', 'contribuirá a'
                ],
                'goal_classification': {
                    'tasa': 'decreciente',
                    'índice': 'constante',
                    'número': 'suma',
                    'porcentaje': 'constante',
                    'cantidad': 'suma',
                    'cobertura': 'suma'
                }
            }
        }

```

```

'contextual_factors': [
    'riesgo', 'amenaza', 'obstáculo', 'limitación',
    'restricción', 'desafío', 'brecha', 'déficit',
    'vulnerabilidad', 'hipótesis alternativa'
],
'administrative_keywords': [
    'gestión', 'administración', 'coordinación', 'regulación',
    'normativa', 'institucional', 'gobernanza', 'reglamento',
    'decreto', 'resolución', 'acuerdo'
]
},
'entity_aliases': {
    'SEC GOB': 'Secretaría de Gobierno',
    'SEC PLAN': 'Secretaría de Planeación',
    'SEC HAC': 'Secretaría de Hacienda',
    'SEC SALUD': 'Secretaría de Salud',
    'SEC EDU': 'Secretaría de Educación',
    'SEC INFRA': 'Secretaría de Infraestructura'
},
'verb_sequences': {
    'diagnosticar': 1,
    'identificar': 2,
    'analizar': 3,
    'diseñar': 4,
    'planificar': 5,
    'implementar': 6,
    'ejecutar': 7,
    'monitorear': 8,
    'evaluar': 9
},
# Bayesian thresholds - now externalized
'bayesian_thresholds': {
    'kl_divergence': get_parameter_loader().get("aaaaaaaa.analysis.derek_beach.ConfigLoader._load_default_config").get("auto_param_L527_33", 0.01),
    'convergence_min_evidence': 2,
    'prior_alpha': 2.0,
    'prior_beta': 2.0,
    'laplace_smoothing': get_parameter_loader().get("aaaaaaaa.analysis.derek_beach.ConfigLoader._load_default_config").get("auto_param_L531_37", 1.0)
},
# Mechanism type priors - now externalized
'mechanism_type_priors': {
    'administrativo': get_parameter_loader().get("aaaaaaaa.analysis.derek_beach.ConfigLoader._load_default_config").get("auto_param_L535_34", 0.30),
    'técnico': get_parameter_loader().get("aaaaaaaa.analysis.derek_beach.ConfigLoader._load_default_config").get("auto_param_L536_27", 0.25),
    'financiero': get_parameter_loader().get("aaaaaaaa.analysis.derek_beach.ConfigLoader._load_default_config").get("auto_param_L537_30", 0.20),
    'político': get_parameter_loader().get("aaaaaaaa.analysis.derek_beach.ConfigLoader._load_default_config").get("auto_param_L538_28", 0.15),
    'mixto': get_parameter_loader().get("aaaaaaaa.analysis.derek_beach.ConfigLoader._load_default_config").get("auto_param_L539_25", 0.10)
},
# Performance settings
'performance': {
    'enable_vectorized_ops': True,
    'enable_async_processing': False,
    'max_context_length': 1000,
    'cache_embeddings': True
},
# Self-reflection settings
'self_reflection': {
    'enable_prior_learning': False,
    'feedback_weight': get_parameter_loader().get("aaaaaaaa.analysis.derek_beach.ConfigLoader._load_default_config").get("auto_param_L551_35", 0.1),
    'prior_history_path': None,
    'min_documents_for_learning': 5
}
}

```

```

        }
        self.logger.warning("Usando configuración por defecto")

@calibrated_method("saaaaaaa.analysis.derek_beach.ConfigLoader._validate_config")
def _validate_config(self) -> None:
    """Validate configuration structure using Pydantic schema"""
    try:
        # Validate with Pydantic schema
        self.validated_config = CDAFConfigSchema(**self.config)
        self.logger.info("✓ Configuración validada exitosamente con esquema Pydantic")
    except ValidationError as e:
        error_details = {
            'validation_errors': [
                {
                    'field': '.'.join(str(x) for x in err['loc']),
                    'error': err['msg'],
                    'type': err['type']
                }
                for err in e.errors()
            ]
        }
        raise CDAFValidationError(
            "Configuración inválida - errores de esquema",
            details=error_details,
            stage="config_validation",
            recoverable=False
        )

# Legacy validation for required sections
required_sections = ['patterns', 'lexicons', 'entity_aliases', 'verb_sequences']
for section in required_sections:
    if section not in self.config:
        self.logger.warning(f"Sección faltante en configuración: {section}")
        self.config[section] = {}

@calibrated_method("saaaaaaa.analysis.derek_beach.ConfigLoader.get")
def get(self, key: str, default: Any = None) -> Any:
    """Get configuration value with dot notation support"""
    keys = key.split('.')
    value = self.config
    for k in keys:
        if isinstance(value, dict):
            value = value.get(k, default)
        else:
            return default
    return value

@calibrated_method("saaaaaaa.analysis.derek_beach.ConfigLoader.get_bayesian_threshold")
def get_bayesian_threshold(self, key: str) -> float:
    """Get Bayesian threshold with type safety"""
    if self.validated_config:
        return getattr(self.validated_config.bayesian_thresholds, key)
    return self.get('bayesian_thresholds.{key}', get_parameter_loader().get("saaaaaaa.analysis.derek_beach.ConfigLoader.get_bayesian_threshold").get("auto_param_L607_54", 0.01))

@calibrated_method("saaaaaaa.analysis.derek_beach.ConfigLoader.get_mechanism_prior")
def get_mechanism_prior(self, mechanism_type: str) -> float:
    """Get mechanism type prior probability with type safety"""
    if self.validated_config:
        return getattr(self.validated_config.mechanism_type_priors, mechanism_type, get_parameter_loader().get("saaaaaaa.analysis.derek_beach.ConfigLoader.get_mechanism_prior").get("auto_param_L613_88", 0.0))
    return self.get('mechanism_type_priors.{mechanism_type}', get_parameter_loader().get("saaaaaaa.analysis.derek_beach.ConfigLoader.get_mechanism_prior").get("auto_param_L614_67", 0.0))

```

```

@calibrated_method("saaaaaaa.analysis.derek_beach.ConfigLoader.get_performance_setting")
def get_performance_setting(self, key: str) -> Any:
    """Get performance setting with type safety"""
    if self.validated_config:
        return getattr(self.validated_config.performance, key)
    return self.get(f"performance.{key}")

@calibrated_method("saaaaaaa.analysis.derek_beach.ConfigLoader.update_priors_from_feedback")
def update_priors_from_feedback(self, feedback_data: dict[str, Any]) -> None:
    """
    Self-reflective loop: Update priors based on audit feedback
    Implements frontier paradigm of learning from results
    """

HARMONIC FRONT 4 ENHANCEMENT:
- Applies penalties to mechanism types with implementation_failure flags
- Heavily penalizes "miracle" mechanisms failing necessity/sufficiency tests
- Ensures mean mech_uncertainty decreases by ≥5% over iterations
"""

if not self.validated_config or not
self.validated_config.self_reflection.enable_prior_learning:
    self.logger.debug("Prior learning disabled")
    return

feedback_weight = self.validated_config.self_reflection.feedback_weight

# Track initial priors for uncertainty measurement
initial_priors = {}
for attr in ['administrativo', 'tecnico', 'financiero', 'politico', 'mixto']:
    if hasattr(self.validated_config.mechanism_type_priors, attr):
        initial_priors[attr] =
getattr(self.validated_config.mechanism_type_priors, attr)

# Update mechanism type priors based on observed frequencies
if 'mechanism_frequencies' in feedback_data:
    for mech_type, observed_freq in
feedback_data['mechanism_frequencies'].items():
        if hasattr(self.validated_config.mechanism_type_priors, mech_type):
            current_prior = getattr(self.validated_config.mechanism_type_priors,
mech_type)
            # Weighted update: new_prior = (1-weight)*current + weight*observed
            updated_prior = (1 - feedback_weight) * current_prior +
feedback_weight * observed_freq
            setattr(self.validated_config.mechanism_type_priors, mech_type,
updated_prior)
            self.config['mechanism_type_priors'][mech_type] = updated_prior

# NEW: Apply penalty factors for failing mechanism types
if 'penalty_factors' in feedback_data:
    penalty_weight = feedback_weight * 1.5 # Heavier penalty than positive
feedback
    for mech_type, penalty_factor in feedback_data['penalty_factors'].items():
        if hasattr(self.validated_config.mechanism_type_priors, mech_type):
            current_prior = getattr(self.validated_config.mechanism_type_priors,
mech_type)
            # Apply penalty: reduce prior for frequently failing types
            penalized_prior = current_prior * penalty_factor
            # Blend with current
            updated_prior = (1 - penalty_weight) * current_prior + penalty_weight
            * penalized_prior
            setattr(self.validated_config.mechanism_type_priors, mech_type,
updated_prior)
            self.config['mechanism_type_priors'][mech_type] = updated_prior
            self.logger.info(f"Applied penalty to {mech_type}: {current_prior:.4f}
-> {updated_prior:.4f}")

# NEW: Heavy penalty for "miracle" mechanisms failing necessity/sufficiency
test_failures = feedback_data.get('test_failures', {})
```

```

if test_failures.get('necessity_failures', 0) > 0 or
test_failures.get('sufficiency_failures', 0) > 0:
    # If failures exist, apply additional penalty to 'politico' (often "miracle"
type)
    # and 'mixto' (vague mechanism types)
    miracle_types = ['politico', 'mixto']
    miracle_penalty = get_parameter_loader().get("saaaaaaa.analysis.derek_beach.Con
figLoader.update_priors_from_feedback").get("miracle_penalty", 0.85) # Refactored
    for mech_type in miracle_types:
        if hasattr(self.validated_config.mechanism_type_priors, mech_type):
            current_prior = getattr(self.validated_config.mechanism_type_priors,
mech_type)
            updated_prior = current_prior * miracle_penalty
            setattr(self.validated_config.mechanism_type_priors, mech_type,
updated_prior)
            self.config['mechanism_type_priors'][mech_type] = updated_prior
            self.logger.info(
                f"Miracle mechanism penalty for {mech_type}: {current_prior:.4f}
-> {updated_prior:.4f}")

# Renormalize to ensure priors sum to get_parameter_loader().get("saaaaaaa.analysis
.derek_beach.ConfigLoader.update_priors_from_feedback").get("auto_param_L686_46", 1.0)
total_prior = sum(
    getattr(self.validated_config.mechanism_type_priors, attr)
    for attr in ['administrativo', 'tecnico', 'financiero', 'politico', 'mixto']
    if hasattr(self.validated_config.mechanism_type_priors, attr)
)
if total_prior > 0:
    for attr in ['administrativo', 'tecnico', 'financiero', 'politico', 'mixto']:
        if hasattr(self.validated_config.mechanism_type_priors, attr):
            current = getattr(self.validated_config.mechanism_type_priors, attr)
            normalized = current / total_prior
            setattr(self.validated_config.mechanism_type_priors, attr, normalized)
            self.config['mechanism_type_priors'][attr] = normalized

# Calculate uncertainty reduction for quality criteria
final_priors = {}
for attr in ['administrativo', 'tecnico', 'financiero', 'politico', 'mixto']:
    if hasattr(self.validated_config.mechanism_type_priors, attr):
        final_priors[attr] = getattr(self.validated_config.mechanism_type_priors,
attr)

# Calculate entropy as uncertainty measure
initial_entropy = -sum(p * np.log(p + 1e-10) for p in initial_priors.values() if p
> 0)
final_entropy = -sum(p * np.log(p + 1e-10) for p in final_priors.values() if p >
0)
uncertainty_reduction = ((initial_entropy - final_entropy) / max(initial_entropy,
1e-10)) * 100

self.logger.info(f"Uncertainty reduction: {uncertainty_reduction:.2f}%")

# Save updated priors if history path configured
if self.validated_config.self_reflection.prior_history_path:
    self._save_prior_history(feedback_data, uncertainty_reduction)

self.logger.info(f"Prioras actualizadas con peso de retroalimentación
{feedback_weight}")

@calibrated_method("saaaaaaa.analysis.derek_beach.ConfigLoader._save_prior_history")
def _save_prior_history(self, feedback_data: dict[str, Any] | None = None,
                        uncertainty_reduction: float | None = None) -> None:
    """
    Save prior history for learning across documents
    """

HARMONIC FRONT 4 ENHANCEMENT:
- Tracks uncertainty reduction over iterations

```

```

- Records penalty applications and test failures
"""
if not self.validated_config or not
self.validated_config.self_reflection.prior_history_path:
    return

try:
    history_path = Path(self.validated_config.self_reflection.prior_history_path)
    history_path.parent.mkdir(parents=True, exist_ok=True)

    # Load existing history if available
    # Delegate to factory for I/O operation
    from .factory import load_json, save_json

    history_records = []
    if history_path.exists():
        try:
            existing_data = load_json(history_path)
            if isinstance(existing_data, list):
                history_records = existing_data
            elif isinstance(existing_data, dict) and 'history' in existing_data:
                history_records = existing_data['history']
        except json.JSONDecodeError:
            self.logger.warning("Existing history file corrupted, starting fresh")

    # Create new record
    history_record = {
        'mechanism_type_priors': dict(self.config.get('mechanism_type_priors',
        {})),
        'timestamp': pd.Timestamp.now().isoformat(),
        'version': '2.0'
    }

    # Add feedback metrics if available
    if feedback_data:
        history_record['audit_quality'] = feedback_data.get('audit_quality', {})
        history_record['test_failures'] = feedback_data.get('test_failures', {})
        history_record['penalty_factors'] = feedback_data.get('penalty_factors',
        {})

    if uncertainty_reduction is not None:
        history_record['uncertainty_reduction_percent'] = uncertainty_reduction

    history_records.append(history_record)

    # Save complete history
    history_data = {
        'version': '2.0',
        'harmonic_front': 4,
        'last_updated': pd.Timestamp.now().isoformat(),
        'total_iterations': len(history_records),
        'history': history_records
    }

    save_json(history_data, history_path)

    self.logger.info(f"Historial de priors guardado en {history_path} (iteración
{len(history_records)}))")
except Exception as e:
    self.logger.warning(f"Error guardando historial de priors: {e}")

```

```

@calibrated_method("saaaaaa.analysis.derek_beach.ConfigLoader._load_uncertainty_history")
def _load_uncertainty_history(self) -> None:
"""
Load historical uncertainty measurements

```

HARMONIC FRONT 4: Required for tracking ≥5% reduction over 10 iterations

```

"""
if not self.validated_config or not
self.validated_config.self_reflection.prior_history_path:
    return

# Delegate to factory for I/O operation
from .factory import load_json

try:
    history_path = Path(self.validated_config.self_reflection.prior_history_path)
    if history_path.exists():
        history_data = load_json(history_path)
        if isinstance(history_data, dict) and 'history' in history_data:
            # Extract uncertainty from each record
            for record in history_data['history']:
                if 'uncertainty_reduction_percent' in record:
                    self._uncertainty_history.append(
                        record['uncertainty_reduction_percent']
                    )
            self.logger.info(f"Loaded {len(self._uncertainty_history)} uncertainty
measurements")
    except Exception as e:
        self.logger.warning(f"Could not load uncertainty history: {e}")

@calibrated_method("saaaaaa.analysis.derek_beach.ConfigLoader.check_uncertainty_reduct
ion_criterion")
def check_uncertainty_reduction_criterion(self, current_uncertainty: float) ->
dict[str, Any]:
"""

Check if mean mechanism_type uncertainty has decreased ≥5% over 10 iterations

HARMONIC FRONT 4 QUALITY CRITERIA:
Success verified if mean mech_uncertainty decreases by ≥5% over 10 sequential PDM
analyses
"""

self._uncertainty_history.append(current_uncertainty)

# Keep only last 10 iterations
recent_history = self._uncertainty_history[-10:]

result = {
    'current_uncertainty': current_uncertainty,
    'iterations_tracked': len(recent_history),
    'criterion_met': False,
    'reduction_percent': get_parameter_loader().get("saaaaaa.analysis.derek_beach.
ConfigLoader.check_uncertainty_reduction_criterion").get("auto_param_L830_33", 0.0),
    'status': 'insufficient_data'
}

if len(recent_history) >= 10:
    initial_uncertainty = recent_history[0]
    final_uncertainty = recent_history[-1]

    if initial_uncertainty > 0:
        reduction_percent = ((initial_uncertainty - final_uncertainty) /
initial_uncertainty) * 100
        result['reduction_percent'] = reduction_percent
        result['criterion_met'] = reduction_percent >= 5.0
        result['status'] = 'success' if result['criterion_met'] else
'needs_improvement'

        self.logger.info(
            f"Uncertainty reduction over 10 iterations: {reduction_percent:.2f}% "
            f"(criterion: ≥5%, met: {result['criterion_met']})"
        )
    else:
        self.logger.info(
            f"Uncertainty tracking: {len(recent_history)}/10 iterations "

```

```

        f"(need {10 - len(recent_history)} more for criterion check)"
    )

return result

class PDFProcessor:
    """Advanced PDF processing and extraction"""

    def __init__(self, config: ConfigLoader, retry_handler=None) -> None:
        self.logger = logging.getLogger(self.__class__.__name__)
        self.config = config
        self.document: fitz.Document | None = None
        self.text_content: str = ""
        self.tables: list[pd.DataFrame] = []
        self.metadata: dict[str, Any] = {}
        self.retry_handler = retry_handler

    @calibrated_method("saaaaaa.analysis.derek_beach.PDFProcessor.load_document")
    def load_document(self, pdf_path: Path) -> bool:
        """Load PDF document with retry logic"""
        if self.retry_handler:
            try:
                from retry_handler import DependencyType

                @self.retry_handler.with_retry(
                    DependencyType.PDF_PARSER,
                    operation_name="open_pdf",
                    exceptions=(IOError, OSError, RuntimeError)
                )
            except Exception as e:
                self.logger.error(f"Error cargando PDF: {e}")
                return False
            else:
                # Fallback without retry
                # Delegate to factory for I/O operation
                from .factory import open_pdf_with_fitz

                doc = open_pdf_with_fitz(pdf_path)
                self.logger.info(f"PDF cargado: {pdf_path.name} ({len(doc)} páginas)")
                return doc

                self.document = load_with_retry()
                self.metadata = self.document.metadata
                return True
        else:
            try:
                self.document = open_pdf_with_fitz(pdf_path)
                self.metadata = self.document.metadata
                self.logger.info(f"PDF cargado: {pdf_path.name} ({len(self.document)} páginas)")
                return True
            except Exception as e:
                self.logger.error(f"Error cargando PDF: {e}")
                return False

    @calibrated_method("saaaaaa.analysis.derek_beach.PDFProcessor.extract_text")
    def extract_text(self) -> str:
        """Extract all text from PDF"""
        if not self.document:
            return ""

        text_parts = []
        for page_num, page in enumerate(self.document, 1):
            try:
                text = page.get_text()
                text_parts.append(text)
            except Exception as e:
                self.logger.error(f"Error extracting text from page {page_num}: {e}")

        return "\n".join(text_parts)

```

```

        text_parts.append(text)
        self.logger.debug(f"Texto extraído de página {page_num}")
    except Exception as e:
        self.logger.warning(f"Error extrayendo texto de página {page_num}: {e}")

    self.text_content = "\n".join(text_parts)
    self.logger.info(f"Texto total extraído: {len(self.text_content)} caracteres")
    return self.text_content

@calibrated_method("saaaaaaa.analysis.derek_beach.PDFProcessor.extract_tables")
def extract_tables(self) -> list[pd.DataFrame]:
    """Extract tables from PDF"""
    if not self.document:
        return []

    table_pattern = re.compile(
        self.config.get('patterns.table_headers', r'PROGRAMA|META|INDICADOR'),
        re.IGNORECASE
    )

    for page_num, page in enumerate(self.document, 1):
        try:
            tabs = page.find_tables()
            if tabs:
                for tab in tabs:
                    try:
                        df = pd.DataFrame(tab.extract())
                        if not df.empty and len(df.columns) > 1:
                            # Check if this is a relevant table
                            header_text = ''.join(str(cell) for cell in df.iloc[0] if
cell)
                            if table_pattern.search(header_text):
                                self.tables.append(df)
                                self.logger.info(f"Tabla extraída de página
{page_num}: {df.shape}")
                    except Exception as e:
                        self.logger.warning(f"Error procesando tabla en página
{page_num}: {e}")
                    except Exception as e:
                        self.logger.debug(f"Error extrayendo tablas de página {page_num}: {e}")

            self.logger.info(f"Total de tablas extraídas: {len(self.tables)}")
            return self.tables

        @calibrated_method("saaaaaaa.analysis.derek_beach.PDFProcessor.extract_sections")
        def extract_sections(self) -> dict[str, str]:
            """Extract document sections based on patterns"""
            sections = {}
            section_pattern = re.compile(
                self.config.get('patterns.section_titles',
r'^(?:CAPÍTULO|ARTÍCULO)\s+[\dIVX]+'),
                re.MULTILINE | re.IGNORECASE
            )

            matches = list(section_pattern.finditer(self.text_content))

            for i, match in enumerate(matches):
                section_title = match.group().strip()
                start_pos = match.end()
                end_pos = matches[i + 1].start() if i + 1 < len(matches) else
len(self.text_content)
                sections[section_title] = self.text_content[start_pos:end_pos].strip()

            self.logger.info(f"Secciones identificadas: {len(sections)}")
            return sections

    class CausalExtractor:
        """Extract and structure causal chains from text"""

```

```

def __init__(self, config: ConfigLoader, nlp_model: spacy.Language) -> None:
    self.logger = logging.getLogger(self.__class__.__name__)
    self.config = config
    self.nlp = nlp_model
    self.graph = nx.DiGraph()
    self.nodes: dict[str, MetaNode] = {}
    self.causal_chains: list[CausalLink] = []

    @calibrated_method("saaaaaaa.analysis.derek_beach.CausalExtractor.extract_causal_hierar
chy")
    def extract_causal_hierarchy(self, text: str) -> nx.DiGraph:
        """Extract complete causal hierarchy from text"""
        # Extract goals/metas
        goals = self._extract_goals(text)

        # Build hierarchy
        for goal in goals:
            self._add_node_to_graph(goal)

        # Extract causal connections
        self._extract_causal_links(text)

        # Build hierarchy based on goal types
        self._build_type_hierarchy()

        self.logger.info(f"Grafo causal construido: {self.graph.number_of_nodes()} nodos,
"
                        f"{self.graph.number_of_edges()} aristas")
        return self.graph

    @calibrated_method("saaaaaaa.analysis.derek_beach.CausalExtractor._extract_goals")
    def _extract_goals(self, text: str) -> list[MetaNode]:
        """Extract all goals from text"""
        goals = []
        goal_pattern = re.compile(
            self.config.get('patterns.goal_codes', r'[MP][RIP]-\d{3}'),
            re.IGNORECASE
        )

        for match in goal_pattern.finditer(text):
            goal_id = match.group().upper()
            context_start = max(0, match.start() - 500)
            context_end = min(len(text), match.end() + 500)
            context = text[context_start:context_end]

            goal = self._parse_goal_context(goal_id, context)
            if goal:
                goals.append(goal)
                self.nodes[goal.id] = goal

        self.logger.info(f"Metas extraídas: {len(goals)}")
        return goals

    @calibrated_method("saaaaaaa.analysis.derek_beach.CausalExtractor._parse_goal_context")
    def _parse_goal_context(self, goal_id: str, context: str) -> MetaNode | None:
        """Parse goal context to extract structured information"""
        # Determine goal type
        if goal_id.startswith('MP'):
            node_type = 'producto'
        elif goal_id.startswith('MR'):
            node_type = 'resultado'
        elif goal_id.startswith('MI'):
            node_type = 'impacto'
        else:
            node_type = 'programa'

        # Extract numerical values

```

```

numeric_pattern = re.compile(
    self.config.get('patterns.numeric_formats', r'[\d,]+(?:\.\d+)?%?')
)
numbers = numeric_pattern.findall(context)

# Process with spaCy
doc = self.nlp(context[:1000])

# Extract entities
entities = [ent.text for ent in doc.ents if ent.label_ in ['ORG', 'PER', 'LOC']]

# Create goal node
goal = MetaNode(
    id=goal_id,
    text=context[:200].strip(),
    type=cast("NodeType", node_type),
    baseline=numbers[0] if len(numbers) > 0 else None,
    target=numbers[1] if len(numbers) > 1 else None,
    responsible_entity=entities[0] if entities else None
)
return goal

```

```

@calibrated_method("saaaaaaa.analysis.derek_beach.CausalExtractor._extract_goal_text")
def _extract_goal_text(self, text: str, **kwargs) -> str | None:
    """
    Extract the text content associated with a specific goal ID.
    """

```

This method extracts goal text from the provided document text. It can work in two modes:

1. If a goal\_id is provided in kwargs, it extracts text for that specific goal
2. Otherwise, it returns the first goal text found in the document

Args:

text: The full document text  
 \*\*kwargs: Additional parameters including optional 'goal\_id', 'data',  
          'sentences', 'tables'

Returns:

The extracted text for the goal, or None if not found

```

# Get goal_id from kwargs if provided, otherwise look for data parameter
goal_id = kwargs.get('goal_id')
kwargs.get('data')

```

# If no goal\_id specified, try to extract the first goal from text

```

if not goal_id:
    goal_pattern = re.compile(
        r'\b[MP][RIP]\-\d{3}\b',
        re.IGNORECASE
    )
    match = goal_pattern.search(text)
    if match:
        goal_id = match.group().upper()
    else:
        # No goal found in text
        return None

```

# Now extract the context around the goal\_id

```

goal_pattern = re.compile(
    rf'\b{re.escape(goal_id)}\b',
    re.IGNORECASE
)

```

```

match = goal_pattern.search(text)
if not match:
    return None

```

```

# Extract context around the goal ID
context_start = max(0, match.start() - 500)
context_end = min(len(text), match.end() + 500)
context = text[context_start:context_end]

return context.strip()

@calibrated_method("saaaaaaa.analysis.derek_beach.CausalExtractor._add_node_to_graph")
def _add_node_to_graph(self, node: MetaNode) -> None:
    """Add node to causal graph"""
    node_dict = asdict(node)
    # Convert NamedTuple to dict for JSON serialization
    if node.entity_activity:
        node_dict['entity_activity'] = node.entity_activity._asdict()
    self.graph.add_node(node.id, **node_dict)

@calibrated_method("saaaaaaa.analysis.derek_beach.CausalExtractor._extract_causal_links")
def _extract_causal_links(self, text: str) -> None:
    """
    AGUJA I: El Prior Informado Adaptativo
    Extract causal links using Bayesian inference with adaptive priors
    """
    causal_keywords = self.config.get('lexicons.causal_logic', [])

    # Get externalized thresholds from configuration
    kl_threshold = self.config.get_bayesian_threshold('kl_divergence')
    convergence_min_evidence =
        self.config.get_bayesian_threshold('convergence_min_evidence')

    # Track evidence for each potential link
    link_evidence: dict[tuple[str, str], list[dict[str, Any]]] = defaultdict(list)

    # Phase 1: Collect all evidence
    for keyword in causal_keywords:
        pattern = re.compile(
            rf'({"|".join(re.escape(nid) for nid in self.nodes)})'
            rf'\s+{re.escape(keyword)}\s+'
            rf'({"|".join(re.escape(nid) for nid in self.nodes)})',
            re.IGNORECASE
        )

        for match in pattern.finditer(text):
            source = match.group(1).upper()
            target = match.group(2).upper()
            logic = match.group(0)

            if source in self.nodes and target in self.nodes:
                # Extract context around the match for language specificity analysis
                context_start = max(0, match.start() - 100)
                context_end = min(len(text), match.end() + 100)
                match_context = text[context_start:context_end]

                # Calculate evidence components
                evidence = {
                    'keyword': keyword,
                    'logic': logic,
                    'match_position': match.start(),
                    'semantic_distance': self._calculate_semantic_distance(source,
                target),
                    'type_transition_prior':
                        self._calculate_type_transition_prior(source, target),
                    'language_specificity':
                        self._calculate_language_specificity(keyword, None, match_context),
                    'temporal_coherence': self._assess_temporal_coherence(source,
                target),
                    'financial_consistency':
                        self._assess_financial_consistency(source, target),
                }

```

```

        'textual_proximity': self._calculate_textual_proximity(source,
target, text)
    }

    link_evidence[(source, target)].append(evidence)

# Phase 2: Bayesian inference for each link
for (source, target), evidences in link_evidence.items():
    # Initialize prior distribution
    prior_mean, prior_alpha, prior_beta = self._initialize_prior(source, target)

    # Incremental Bayesian update
    posterior_alpha = prior_alpha
    posterior_beta = prior_beta
    kl_divs = []

    for evidence in evidences:
        # Calculate likelihood components
        likelihood = self._calculate_composite_likelihood(evidence)

        # Update Beta distribution parameters
        # Using Beta-Binomial conjugate prior
        posterior_alpha += likelihood
        posterior_beta += (1 - likelihood)

        # Calculate KL divergence for convergence check
        if len(kl_divs) > 0:
            prior_dist = np.array([posterior_alpha - likelihood, posterior_beta -
(1 - likelihood)])
            prior_dist = prior_dist / prior_dist.sum()
            posterior_dist = np.array([posterior_alpha, posterior_beta])
            posterior_dist = posterior_dist / posterior_dist.sum()
            kl_div = float(np.sum(rel_entr(posterior_dist, prior_dist)))
            kl_divs.append(kl_div)

        # Calculate posterior statistics
        posterior_mean = posterior_alpha / (posterior_alpha + posterior_beta)
        posterior_var = (posterior_alpha * posterior_beta) / (
            (posterior_alpha + posterior_beta) ** 2 * (posterior_alpha +
posterior_beta + 1))
        posterior_std = np.sqrt(posterior_var)

    # AUDIT POINT 2.1: Structural Veto (D6-Q2)
    # TeoriaCambio validation - caps Bayesian posterior ≤get_parameter_loader().ge
    t("saaaaaaa.analysis.derek_beach.CausalExtractor._extract_causal_links").get("auto_param_L1
217_65", 0.6) for impermissible links
    # Implements axiomatic-Bayesian fusion per Goertz & Mahoney 2012
    structuralViolation = self._check_structural_violation(source, target)
    if structuralViolation:
        # Deterministic veto: cap posterior at get_parameter_loader().get("saaaaaaa
.analysis.derek_beach.CausalExtractor._extract_causal_links").get("auto_param_L1221_55",
0.6) despite high semantic evidence
        originalPosterior = posterior_mean
        posterior_mean = min(posterior_mean, get_parameter_loader().get("saaaaaaa.a
nalysis.derek_beach.CausalExtractor._extract_causal_links").get("auto_param_L1223_53",
0.6))
        self.logger.warning(
            f"STRUCTURAL VETO (D6-Q2): Link {source}→{target} violates causal
hierarchy."
            f"Posterior capped from {originalPosterior:.3f} to
{posterior_mean:.3f}."
            f"Violation: {structuralViolation}"
        )

    # Check convergence (require minimum evidence count)
    converged = (len(kl_divs) >= convergence_min_evidence and
        len(kl_divs) > 0 and kl_divs[-1] < kl_threshold)

```

```

final_kl = kl_divs[-1] if len(kl_divs) > 0 else get_parameter_loader().get("sa
aaaaaa.analysis.derek_beach.CausalExtractor._extract_causal_links").get("auto_param_L1233_6
0", 0.0)

# Add edge with posterior distribution
self.graph.add_edge(
    source, target,
    logic=evidences[0]['logic'],
    keyword=evidences[0]['keyword'],
    strength=float(posterior_mean),
    posterior_mean=float(posterior_mean),
    posterior_std=float(posterior_std),
    posterior_alpha=float(posterior_alpha),
    posterior_beta=float(posterior_beta),
    kl_divergence=float(final_kl),
    converged=converged,
    evidence_count=len(evidences),
    structuralViolation=structuralViolation,
    veto_applied=structuralViolation is not None
)
self.causal_chains.append({
    'source': source,
    'target': target,
    'logic': evidences[0]['logic'],
    'strength': float(posterior_mean),
    'evidence': [e['keyword'] for e in evidences],
    'posterior_mean': float(posterior_mean),
    'posterior_std': float(posterior_std),
    'kl_divergence': float(final_kl),
    'converged': converged
})

```

```

self.logger.info(f"Enlaces causales extraídos: {len(self.causal_chains)} "
                f"(con inferencia Bayesiana)")

```

```

@calibrated_method("aaaaaaaa.analysis.derek_beach.CausalExtractor._calculate_semantic_d
istance")

```

```

def _calculate_semantic_distance(self, source: str, target: str) -> float:
    """

```

Calculate semantic distance between nodes using spaCy embeddings

PERFORMANCE NOTE: This method can be optimized with:

1. Vectorized operations using numpy for batch processing
2. Embedding caching to avoid recomputing spaCy vectors
3. Async processing for large documents with many nodes
4. Alternative: BERT/transformer embeddings for higher fidelity (SOTA)

Current implementation prioritizes determinism over speed.

Enable performance.cache\_embeddings in config for production use.

```

try:

```

```

    source_node = self.nodes.get(source)
    target_node = self.nodes.get(target)

```

```

    if not source_node or not target_node:

```

```

        return get_parameter_loader().get("aaaaaaaa.analysis.derek_beach.CausalExtr
actor._calculate_semantic_distance").get("auto_param_L1286_23", 0.5)
    
```

```

# TODO: Implement embedding cache if performance.cache_embeddings is enabled
# This would save ~60% computation time on large documents

```

```

# Use spaCy to get embeddings
max_context = self.config.get_performance_setting('max_context_length') or

```

1000

```

    source_doc = self.nlp(source_node.text[:max_context])
    target_doc = self.nlp(target_node.text[:max_context])

```

```

if source_doc.vector.any() and target_doc.vector.any():
    # Calculate cosine similarity (1 - distance)
    # PERFORMANCE NOTE: Could vectorize this with numpy.dot for batch
operations
    similarity = 1 - cosine(source_doc.vector, target_doc.vector)
    return max(get_parameter_loader().get("saaaaaaa.analysis.derek_beach.Causal
Extractor._calculate_semantic_distance").get("auto_param_L1300_27", 0.0), min(get_parameter
r_loader().get("saaaaaaa.analysis.derek_beach.CausalExtractor._calculate_semantic_distance"
).get("auto_param_L1300_36", 1.0), similarity))

    return get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CausalExtracto
r._calculate_semantic_distance").get("auto_param_L1302_19", 0.5)
except Exception:
    return get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CausalExtracto
r._calculate_semantic_distance").get("auto_param_L1304_19", 0.5)

@calibrated_method("saaaaaaa.analysis.derek_beach.CausalExtractor._calculate_type_trans
ition_prior")
def _calculate_type_transition_prior(self, source: str, target: str) -> float:
    """Calculate prior based on historical transition frequencies between goal
types"""
    source_type = self.nodes[source].type
    target_type = self.nodes[target].type

    # Define transition probabilities based on logical flow
    # programa → producto → resultado → impacto
    transition_priors = {
        ('programa', 'producto'): get_parameter_loader().get("saaaaaaa.analysis.derek_b
each.CausalExtractor._calculate_type_transition_prior").get("auto_param_L1315_38", 0.85),
        ('producto', 'resultado'): get_parameter_loader().get("saaaaaaa.analysis.derek_
beach.CausalExtractor._calculate_type_transition_prior").get("auto_param_L1316_39", 0.80),
        ('resultado', 'impacto'): get_parameter_loader().get("saaaaaaa.analysis.derek_b
each.CausalExtractor._calculate_type_transition_prior").get("auto_param_L1317_38", 0.75),
        ('programa', 'resultado'): get_parameter_loader().get("saaaaaaa.analysis.derek_
beach.CausalExtractor._calculate_type_transition_prior").get("auto_param_L1318_39", 0.60),
        ('producto', 'impacto'): get_parameter_loader().get("saaaaaaa.analysis.derek_be
ach.CausalExtractor._calculate_type_transition_prior").get("auto_param_L1319_37", 0.50),
        ('programa', 'impacto'): get_parameter_loader().get("saaaaaaa.analysis.derek_be
ach.CausalExtractor._calculate_type_transition_prior").get("auto_param_L1320_37", 0.30),
    }

    # Reverse transitions are less likely
    reverse_key = (target_type, source_type)
    if reverse_key in transition_priors:
        return transition_priors[reverse_key] * get_parameter_loader().get("saaaaaaa.an
alysis.derek_beach.CausalExtractor._calculate_type_transition_prior").get("auto_param_L132
6_52", 0.3)

    return transition_priors.get((source_type, target_type), get_parameter_loader().ge
t("saaaaaaa.analysis.derek_beach.CausalExtractor._calculate_type_transition_prior").get("au
to_param_L1328_65", 0.40))

@calibrated_method("saaaaaaa.analysis.derek_beach.CausalExtractor._check_structural_vio
lation")
def _check_structuralViolation(self, source: str, target: str) -> str | None:
    """
AUDIT POINT 2.1: Structural Veto (D6-Q2)

Check if causal link violates structural hierarchy based on TeoriaCambio axioms.
Implements set-theoretic constraints per Goertz & Mahoney 2012.

Returns:
    None if link is valid, otherwise a string describing the violation
    """
    source_type = self.nodes[source].type
    target_type = self.nodes[target].type

    # Define causal hierarchy levels (following TeoriaCambio axioms)

```

```

# Lower levels cannot causally influence higher levels
hierarchy_levels = {
    'programa': 1,
    'producto': 2,
    'resultado': 3,
    'impacto': 4
}

source_level = hierarchy_levels.get(source_type, 0)
target_level = hierarchy_levels.get(target_type, 0)

# Impermissible links: jumping more than 2 levels or reverse causation
if target_level < source_level:
    # Reverse causation (e.g., Impacto → Producto)
    return f"reverse_causation:{source_type}→{target_type}"

if target_level - source_level > 2:
    # Skipping levels (e.g., Programa → Impacto without intermediates)
    return f"level_skip:{source_type}→{target_type} (skips {target_level - source_level - 1} levels)"

# Special case: Producto → Impacto is impermissible (must go through Resultado)
if source_type == 'producto' and target_type == 'impacto':
    return "missing_intermediate:producto→impacto requires resultado"

return None

@calibrated_method("saaaaaa.analysis.derek_beach.CausalExtractor._calculate_language_specificity")
def _calculate_language_specificity(self, keyword: str, policy_area: str | None = None,
                                     context: str | None = None) -> float:
    """Assess specificity of causal language (epistemic certainty)

Harmonic Front 3 - Enhancement 4: Language Specificity Assessment
Enhanced to check policy-specific vocabulary (patrones_verificacion) for current
Policy Area (P1–P10), not just generic causal keywords.

For D6-Q5 (Contextual/Differential Focus): rewards use of specialized terminology
that anchors intervention in social/cultural context (e.g., "catastro
multipropósito",
"reparación integral", "mujeres rurales", "guardia indígena").
"""

# Strong causal indicators
strong_indicators = ['causa', 'produce', 'genera', 'resulta en', 'conduce a']
# Moderate indicators
moderate_indicators = ['permite', 'contribuye', 'facilita', 'mediante', 'a través
de']
# Weak indicators
weak_indicators = ['con el fin de', 'para', 'porque']

keyword_lower = keyword.lower()

# Base score from causal indicators
base_score = get_parameter_loader().get("saaaaaa.analysis.derek_beach.CausalExtractor._calculate_language_specificity").get("base_score", 0.6) # Refactored
if any(ind in keyword_lower for ind in strong_indicators):
    base_score = get_parameter_loader().get("saaaaaa.analysis.derek_beach.CausalExtractor._calculate_language_specificity").get("base_score", 0.9) # Refactored
elif any(ind in keyword_lower for ind in moderate_indicators):
    base_score = get_parameter_loader().get("saaaaaa.analysis.derek_beach.CausalExtractor._calculate_language_specificity").get("base_score", 0.7) # Refactored
elif any(ind in keyword_lower for ind in weak_indicators):
    base_score = get_parameter_loader().get("saaaaaa.analysis.derek_beach.CausalExtractor._calculate_language_specificity").get("base_score", 0.5) # Refactored

# HARMONIC FRONT 3 - Enhancement 4: Policy-specific vocabulary boost
# Check for specialized terminology per policy area

```

```

policy_area_vocabulary = {
    'P1': [ # Ordenamiento Territorial
        'catastro multipropósito', 'pot', 'pbott', 'eot', 'uaf', 'suelo de
protección',
        'zonificación', 'uso del suelo', 'densificación', 'expansión urbana'
    ],
    'P2': [ # Víctimas y Paz
        'reparación integral', 'restitución de tierras', 'víctimas del conflicto',
        'desplazamiento forzado', 'despojo', 'acción integral', 'enfoque
diferencial étnico',
        'construcción de paz', 'reconciliación', 'memoria histórica'
    ],
    'P3': [ # Desarrollo Rural
        'mujeres rurales', 'extensión agropecuaria', 'asistencia técnica rural',
        'adecuación de tierras', 'comercialización campesina', 'economía
campesina',
        'soberanía alimentaria', 'fondo de tierras'
    ],
    'P4': [ # Grupos Étnicos
        'guardia indígena', 'guardia cimarrona', 'territorios colectivos',
        'autoridades ancestrales', 'consulta previa', 'consentimiento libre',
        'medicina tradicional', 'sistema de salud propio indígena', 'jurisdicción
especial indígena'
    ],
    'P5': [ # Infraestructura y Conectividad
        'terciarias', 'vías terciarias', 'transporte intermodal', 'último
kilómetro',
        'conectividad digital', 'internet rural', 'electrificación rural'
    ],
    'P6': [ # Salud Rural
        'red hospitalaria', 'atención primaria', 'promotores de salud',
        'prevención de enfermedades tropicales', 'saneamiento básico', 'agua
segura'
    ],
    'P7': [ # Educación Rural
        'escuela nueva', 'modelos flexibles', 'post-primaria rural',
        'educación propia', 'alfabetización', 'deserción escolar rural'
    ],
    'P8': [ # Vivienda y Habitabilidad
        'mejoramiento de vivienda rural', 'materiales locales', 'construcción
sostenible',
        'vivienda de interés social rural', 'titulación predial'
    ],
    'P9': [ # Medio Ambiente
        'páramos', 'humedales', 'áreas protegidas', 'corredores biológicos',
        'servicios ecosistémicos', 'pago por servicios ambientales', 'restauración
ecológica'
    ],
    'P10': [ # Reactivación Económica
        'encadenamientos productivos', 'economía solidaria', 'cooperativas',
        'microcrédito', 'emprendimiento asociativo', 'fondo rotatorio'
    ]
}
}

# General contextual/differential focus vocabulary (D6-Q5)
contextual_vocabulary = [
    'enfoque diferencial', 'enfoque de género', 'enfoque étnico',
    'acción sin daño', 'pertinencia cultural', 'contexto territorial',
    'restrictiones territoriales', 'barreras culturales', 'inequidad',
    'discriminación', 'exclusión', 'vulnerabilidad', 'marginalidad',
    'ruralidad dispersa', 'aislamiento geográfico', 'baja densidad poblacional',
    'población dispersa', 'difícil acceso'
]

# Check for policy-specific vocabulary boost
specificity_boost = get_parameter_loader().get("saaaaaaa.analysis.derek_beach.Causa
IExtractor._calculate_language_specificity").get("specificity_boost", 0.0) # Refactored
text_to_check = (keyword_lower + ' ' + (context or "")).lower()

```

```

if policy_area and policy_area in policy_area_vocabulary:
    for term in policy_area_vocabulary[policy_area]:
        if term.lower() in text_to_check:
            specificity_boost = max(specificity_boost, get_parameter_loader().get(
                "saaaaaaa.analysis.derek_beach.CausalExtractor._calculate_language_specificity").get("auto_"
                "param_L1467_63", 0.15))
            self.logger.debug(f"Policy-specific term detected: '{term}' for"
                f"{policy_area}")
            break

    # Check for general contextual vocabulary (D6-Q5)
    for term in contextual_vocabulary:
        if term.lower() in text_to_check:
            specificity_boost = max(specificity_boost, get_parameter_loader().get("saa
aaaa.analysis.derek_beach.CausalExtractor._calculate_language_specificity").get("auto_para
m_L1474_59", 0.10))
            self.logger.debug(f"Contextual term detected: '{term}'")
            break

    final_score = min(get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CausalE
xtractor._calculate_language_specificity").get("auto_param_L1478_26", 1.0), base_score +
specificity_boost)

return final_score

@calibrated_method("saaaaaaa.analysis.derek_beach.CausalExtractor._assess_temporal_coh
herence")
def _assess_temporal_coherence(self, source: str, target: str) -> float:
    """Assess temporal coherence based on verb sequences"""
    source_node = self.nodes.get(source)
    target_node = self.nodes.get(target)

    if not source_node or not target_node:
        return get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CausalExtracto
r._assess_temporal_coherence").get("auto_param_L1489_19", 0.5)

    # Extract verbs from entity-activity if available
    if source_node.entity_activity and target_node.entity_activity:
        source_verb = source_node.entity_activity.verb_lemma
        target_verb = target_node.entity_activity.verb_lemma

    # Define logical verb sequences
    verb_sequences = {
        'diagnosticar': 1, 'planificar': 2, 'ejecutar': 3, 'evaluar': 4,
        'diseñar': 2, 'implementar': 3, 'monitorear': 4
    }

    source_seq = verb_sequences.get(source_verb, 5)
    target_seq = verb_sequences.get(target_verb, 5)

    if source_seq < target_seq:
        return get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CausalExtr
actor._assess_temporal_coherence").get("auto_param_L1506_23", 0.85)
    elif source_seq == target_seq:
        return get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CausalExtr
actor._assess_temporal_coherence").get("auto_param_L1508_23", 0.60)
    else:
        return get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CausalExtr
actor._assess_temporal_coherence").get("auto_param_L1510_23", 0.30)

    return get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CausalExtractor._a
ssess_temporal_coherence").get("auto_param_L1512_15", 0.50)

@calibrated_method("saaaaaaa.analysis.derek_beach.CausalExtractor._assess_financial_con
sistency")
def _assess_financial_consistency(self, source: str, target: str) -> float:
    """Assess financial alignment between connected nodes"""

```

```

source_node = self.nodes.get(source)
target_node = self.nodes.get(target)

if not source_node or not target_node:
    return get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CausalExtractor._assess_financial_consistency").get("auto_param_L1521_19", 0.5)

source_budget = source_node.financial_allocation
target_budget = target_node.financial_allocation

if source_budget and target_budget:
    # Check if budgets are aligned (target should be <= source)
    ratio = target_budget / source_budget if source_budget > 0 else 0

    if get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CausalExtractor._assess_financial_consistency").get("auto_param_L1530_15", 0.1) <= ratio <= get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CausalExtractor._assess_financial_consistency").get("auto_param_L1530_31", 1.0):
        return get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CausalExtractor._assess_financial_consistency").get("auto_param_L1531_23", 0.85)
    elif ratio > get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CausalExtractor._assess_financial_consistency").get("auto_param_L1532_25", 1.0) and ratio <= 1.5:
        return get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CausalExtractor._assess_financial_consistency").get("auto_param_L1533_23", 0.60)
    else:
        return get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CausalExtractor._assess_financial_consistency").get("auto_param_L1535_23", 0.30)

    return get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CausalExtractor._assess_financial_consistency").get("auto_param_L1537_15", 0.50)

@calibrated_method("saaaaaaa.analysis.derek_beach.CausalExtractor._calculate_textual_proximity")
def _calculate_textual_proximity(self, source: str, target: str, text: str) -> float:
    """Calculate how often node IDs appear together in text windows"""
    window_size = 200 # characters
    co_occurrences = 0
    total_windows = 0

    source_positions = [m.start() for m in re.finditer(re.escape(source), text, re.IGNORECASE)]
    target_positions = [m.start() for m in re.finditer(re.escape(target), text, re.IGNORECASE)]

    for source_pos in source_positions:
        total_windows += 1
        for target_pos in target_positions:
            if abs(source_pos - target_pos) <= window_size:
                co_occurrences += 1
                break

    if total_windows > 0:
        proximity_score = co_occurrences / total_windows
        return proximity_score

    return get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CausalExtractor._calculate_textual_proximity").get("auto_param_L1560_15", 0.5)

@calibrated_method("saaaaaaa.analysis.derek_beach.CausalExtractor._initialize_prior")
def _initialize_prior(self, source: str, target: str) -> tuple[float, float, float]:
    """Initialize prior distribution for causal link"""
    # Use type transition as base prior
    type_prior = self._calculate_type_transition_prior(source, target)

    # Beta distribution parameters - now externalized
    prior_alpha = self.config.get_bayesian_threshold('prior_alpha')
    prior_beta = self.config.get_bayesian_threshold('prior_beta')

```

```

# Adjust based on type transition
prior_mean = type_prior
prior_strength = prior_alpha + prior_beta

adjusted_alpha = prior_mean * prior_strength
adjusted_beta = (1 - prior_mean) * prior_strength

return prior_mean, adjusted_alpha, adjusted_beta

@calibrated_method("saaaaaaa.analysis.derek_beach.CausalExtractor._calculate_composite_
likelihood")
def _calculate_composite_likelihood(self, evidence: dict[str, Any]) -> float:
    """Calculate composite likelihood from multiple evidence components

Enhanced with:
- Nonlinear transformation rewarding triangulation
- Evidence diversity verification across analytical domains
"""

# Weight different evidence types
weights = {
    'semantic_distance': get_parameter_loader().get("saaaaaaa.analysis.derek_beach.
CausalExtractor._calculate_composite_likelihood").get("auto_param_L1591_33", 0.25),
    'type_transition_prior': get_parameter_loader().get("saaaaaaa.analysis.derek_be
ach.CausalExtractor._calculate_composite_likelihood").get("auto_param_L1592_37", 0.20),
    'language_specificity': get_parameter_loader().get("saaaaaaa.analysis.derek_bea
ch.CausalExtractor._calculate_composite_likelihood").get("auto_param_L1593_36", 0.20),
    'temporal_coherence': get_parameter_loader().get("saaaaaaa.analysis.derek_beach
.CausalExtractor._calculate_composite_likelihood").get("auto_param_L1594_34", 0.15),
    'financial_consistency': get_parameter_loader().get("saaaaaaa.analysis.derek_be
ach.CausalExtractor._calculate_composite_likelihood").get("auto_param_L1595_37", 0.10),
    'textual_proximity': get_parameter_loader().get("saaaaaaa.analysis.derek_beach.
CausalExtractor._calculate_composite_likelihood").get("auto_param_L1596_33", 0.10)
}

# Basic weighted average
likelihood = get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CausalExtrac
tor._calculate_composite_likelihood").get("likelihood", 0.0) # Refactored
evidence_count = 0
domain_diversity = set()

for component, weight in weights.items():
    if component in evidence:
        likelihood += evidence[component] * weight
        evidence_count += 1

    # Track evidence diversity across domains
    if component in ['semantic_distance', 'textual_proximity']:
        domain_diversity.add('semantic')
    elif component in ['temporal_coherence']:
        domain_diversity.add('temporal')
    elif component in ['financial_consistency']:
        domain_diversity.add('financial')
    elif component in ['type_transition_prior', 'language_specificity']:
        domain_diversity.add('structural')

# Triangulation bonus: Exponentially reward multiple independent observations
# D6-Q4/Q5 (Adaptiveness/Context) - evidence across different analytical domains
diversity_count = len(domain_diversity)
if diversity_count >= 3:
    # Strong triangulation across semantic, temporal, and financial domains
    triangulation_bonus = get_parameter_loader().get("saaaaaaa.analysis.derek_beach
.CausalExtractor._calculate_composite_likelihood").get("auto_param_L1624_34", 1.0) + get_p
arameter_loader().get("saaaaaaa.analysis.derek_beach.CausalExtractor._calculate_composite_l
ikelihood").get("auto_param_L1624_40", 0.15) * np.exp(diversity_count - 2)
    elif diversity_count == 2:
        # Moderate triangulation
        triangulation_bonus = get_parameter_loader().get("saaaaaaa.analysis.derek_beach
.CausalExtractor._calculate_composite_likelihood").get("auto_param_L1627_34", 1.0)5

```

```

else:
    # Weak or no triangulation
    triangulation_bonus = get_parameter_loader().get("saaaaaaa.analysis.derek_beach
.CausalExtractor._calculate_composite_likelihood").get("triangulation_bonus", 1.0) #
Refactored

    # Apply nonlinear transformation
    enhanced_likelihood = min(get_parameter_loader().get("saaaaaaa.analysis.derek_beach
.CausalExtractor._calculate_composite_likelihood").get("auto_param_L1633_34", 1.0),
likelihood * triangulation_bonus)

    # Penalty for insufficient evidence diversity
    if evidence_count < 3:
        enhanced_likelihood *= get_parameter_loader().get("saaaaaaa.analysis.derek_beac
h.CausalExtractor._calculate_composite_likelihood").get("auto_param_L1637_35", 0.85)

return enhanced_likelihood

```

```

@calibrated_method("saaaaaaa.analysis.derek_beach.CausalExtractor._build_type_hierarchy")
def _build_type_hierarchy(self) -> None:
    """Build hierarchy based on goal types"""

    nodes_by_type: dict[str, list[str]] = defaultdict(list)
    for node_id in self.graph.nodes():
        node_type = self.graph.nodes[node_id].get('type', 'programa')
        nodes_by_type[node_type].append(node_id)

    # Connect productos to programas
    for prod in nodes_by_type.get('producto', []):
        for prog in nodes_by_type.get('programa', []):
            if not self.graph.has_edge(prog, prod):
                self.graph.add_edge(prog, prod, logic='inferido', strength=get_paramet
er_loader().get("saaaaaaa.analysis.derek_beach.CausalExtractor._build_type_hierarchy").get(
"auto_param_L1654_79", 0.5))

    # Connect resultados to productos
    for res in nodes_by_type.get('resultado', []):
        for prod in nodes_by_type.get('producto', []):
            if not self.graph.has_edge(prod, res):
                self.graph.add_edge(prod, res, logic='inferido', strength=get_paramete
r_loader().get("saaaaaaa.analysis.derek_beach.CausalExtractor._build_type_hierarchy").get(
"auto_param_L1660_78", 0.5))

```

```

@calibrated_method("saaaaaaa.analysis.derek_beach.CausalExtractor._calculate_confidence")
def _calculate_confidence(self, node: MetaNode, link_text: str = "") -> float:
    """

```

Calculate confidence score for a causal link.

Args:

node: The node to calculate confidence for  
link\_text: Optional text describing the causal link

Returns:

Confidence score between 0 and 1

"""

```

confidence = get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CausalExtrac
tor._calculate_confidence").get("confidence", 0.5) # Refactored

```

# Increase confidence if node has quantitative targets

if node.target and node.baseline:

try:

float(str(node.target).replace(',', '').replace('%', ''))

confidence += get\_parameter\_loader().get("saaaaaaa.analysis.derek\_beach.Cau
salExtractor.\_calculate\_confidence").get("auto\_param\_L1680\_30", 0.2)

except (ValueError, TypeError):

pass

```

# Increase confidence if text has causal indicators
if link_text:
    causal_words = ['porque', 'debido', 'mediante', 'a través', 'permite',
'genera', 'produce']
    if any(word in link_text.lower() for word in causal_words):
        confidence += get_parameter_loader().get("saaaaaaa.analysis.derek_beach.Cau
salExtractor._calculate_confidence").get("auto_param_L1688_30", 0.15)

# Increase confidence based on rigor status
if hasattr(node, 'rigor_status'):
    if node.rigor_status == 'fuerte':
        confidence += get_parameter_loader().get("saaaaaaa.analysis.derek_beach.Cau
salExtractor._calculate_confidence").get("auto_param_L1693_30", 0.15)
    elif node.rigor_status == 'débil':
        confidence -= get_parameter_loader().get("saaaaaaa.analysis.derek_beach.Cau
salExtractor._calculate_confidence").get("auto_param_L1695_30", 0.1)

    return min(get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CausalExtracto
r._calculate_confidence").get("auto_param_L1697_19", 1.0), max(get_parameter_loader().get(
"saaaaaaa.analysis.derek_beach.CausalExtractor._calculate_confidence").get("auto_param_L169
7_28", 0.0), confidence))

```

`@calibrated_method("saaaaaaa.analysis.derek_beach.CausalExtractor._classify_goal_type")`

`def _classify_goal_type(self, text: str) -> str:`

`"""`

Classify the type of a goal based on its text.

Args:

`text: Goal text to classify`

Returns:

`Goal type (programa, producto, resultado, impacto)`

`"""`

`text_lower = text.lower()`

`# Keywords for each type`

```

if any(word in text_lower for word in ['programa', 'línea estratégica',
'componente', 'eje']):

```

`return 'programa'`

```

elif any(word in text_lower for word in ['producto', 'servicio', 'bien',
'actividad']):

```

`return 'producto'`

```

elif any(word in text_lower for word in ['resultado', 'efecto', 'cambio',
'mejora']):

```

`return 'resultado'`

```

elif any(word in text_lower for word in ['impacto', 'transformación',
'desarrollo', 'bienestar']):

```

`return 'impacto'`

`# Default classification based on position and complexity`

`elif len(text) < 100:`

`return 'producto'`

`else:`

`return 'resultado'`

`@calibrated_method("saaaaaaa.analysis.derek_beach.CausalExtractor._extract_causal_justi
fications")`

`def _extract_causal_justifications(self, text: str) -> list[dict[str, Any]]:`

`"""`

Extract causal justifications from text.

Args:

`text: Text to extract justifications from`

Returns:

`List of justifications with text and confidence`

`"""`

`justifications = []`

```

# Patterns that indicate causal justifications
patterns = [
    r'porque\s+([^.]+)',
    r'debido\s+a\s+([^.]+)',
    r'mediante\s+([^.]+)',
    r'a\s+través\s+de\s+([^.]+)',
    r'se\s+logra\s+mediante\s+([^.]+)',
    r'permite\s+([^.]+)',
    r'genera\s+([^.]+)',
]
for pattern in patterns:
    matches = re.finditer(pattern, text, re.IGNORECASE)
    for match in matches:
        justification_text = match.group(1).strip()
        justifications.append({
            'text': justification_text,
            'confidence': get_parameter_loader().get("saaaaaaa.analysis.derek_beach"
.CausalExtractor._extract_causal_justifications").get("auto_param_L1757_34", 0.7),
            'type': 'causal_explanation'
        })
return justifications

class MechanismPartExtractor:
    """Extract Entity-Activity pairs for mechanism parts"""

    def __init__(self, config: ConfigLoader, nlp_model: spacy.Language) -> None:
        self.logger = logging.getLogger(self.__class__.__name__)
        self.config = config
        self.nlp = nlp_model
        self.entity_aliases = config.get('entity_aliases', {})

    @calibrated_method("saaaaaaa.analysis.derek_beach.MechanismPartExtractor.extract_entity"
_activity")
    def extract_entity_activity(self, text: str) -> EntityActivity | None:
        """Extract Entity-Activity tuple from text"""
        doc = self.nlp(text)

        # Find main verb (activity)
        main_verb = None
        for token in doc:
            if token.pos_ == 'VERB' and token.dep_ in ['ROOT', 'ccomp']:
                main_verb = token
                break

        if not main_verb:
            return None

        # Find subject entity
        entity = None
        for child in main_verb.children:
            if child.dep_ in ['nsubj', 'nsubjpass']:
                entity = self._normalize_entity(child.text)
                break

        if not entity:
            # Try to find entity from NER
            for ent in doc.ents:
                if ent.label_ in ['ORG', 'PER']:
                    entity = self._normalize_entity(ent.text)
                    break

        if entity and main_verb:
            return EntityActivity(
                entity=entity,
                activity=main_verb.text,
            )

```

```

verb_lemma=main_verb.lemma_
confidence = get_parameter_loader().get("saaaaaaa.analysis.derek_beach.MechanismPartExtractor.extract_entity_activity").get("confidence", 0.85) # Refactored
)

return None

@calibrated_method("saaaaaaa.analysis.derek_beach.MechanismPartExtractor._normalize_entity")
def _normalize_entity(self, entity: str) -> str:
    """Normalize entity name using aliases"""
    entity_upper = entity.upper().strip()
    return self.entity_aliases.get(entity_upper, entity)

@calibrated_method("saaaaaaa.analysis.derek_beach.MechanismPartExtractor._calculate_ea_confidence")
def _calculate_ea_confidence(self, entity: str, activity: str, context: str = "") ->
float:
    """
    Calculate confidence for an entity-activity pair.

    Args:
        entity: Entity text
        activity: Activity text
        context: Surrounding context

    Returns:
        Confidence score between 0 and 1
    """
    confidence = get_parameter_loader().get("saaaaaaa.analysis.derek_beach.MechanismPartExtractor._calculate_ea_confidence").get("confidence", 0.5) # Refactored

    # Higher confidence if entity is in known aliases
    if entity.upper() in self.entity_aliases:
        confidence += get_parameter_loader().get("saaaaaaa.analysis.derek_beach.MechanismPartExtractor._calculate_ea_confidence").get("auto_param_L1834_26", 0.2)

    # Higher confidence if activity is a strong verb
    strong_verbs = ['ejecutar', 'implementar', 'desarrollar', 'gestionar',
'coordinar']
    if any(verb in activity.lower() for verb in strong_verbs):
        confidence += get_parameter_loader().get("saaaaaaa.analysis.derek_beach.MechanismPartExtractor._calculate_ea_confidence").get("auto_param_L1839_26", 0.15)

    # Higher confidence if there's clear grammatical connection in context
    if entity in context and activity in context:
        confidence += get_parameter_loader().get("saaaaaaa.analysis.derek_beach.MechanismPartExtractor._calculate_ea_confidence").get("auto_param_L1843_26", 0.15)

    return min(get_parameter_loader().get("saaaaaaa.analysis.derek_beach.MechanismPartExtractor._calculate_ea_confidence").get("auto_param_L1845_19", 1.0), confidence)

@calibrated_method("saaaaaaa.analysis.derek_beach.MechanismPartExtractor._find_action_verb")
def _find_action_verb(self, text: str) -> str | None:
    """
    Find the main action verb in text.

    Args:
        text: Text to analyze

    Returns:
        Main action verb or None
    """
    doc = self.nlp(text)

    # Find main verb
    for token in doc:

```

```

if token.pos_ == 'VERB' and token.dep_ in ['ROOT', 'ccomp', 'xcomp']:
    return token.text

# Fallback: any verb
for token in doc:
    if token.pos_ == 'VERB':
        return token.text

return None

@calibrated_method("saaaaaa.analysis.derek_beach.MechanismPartExtractor._find_subject_
entity")
def _find_subject_entity(self, text: str) -> str | None:
    """
    Find the subject entity in text.

    Args:
        text: Text to analyze

    Returns:
        Subject entity or None
    """
    doc = self.nlp(text)

    # Find subject
    for token in doc:
        if token.dep_ in ['nsubj', 'nsubjpass']:
            return self._normalize_entity(token.text)

    # Try NER
    for ent in doc.ents:
        if ent.label_ in ['ORG', 'PER', 'GPE']:
            return self._normalize_entity(ent.text)

    return None

@calibrated_method("saaaaaa.analysis.derek_beach.MechanismPartExtractor._validate_enti
ty_activity")
def _validate_entity_activity(self, entity: str, activity: str) -> bool:
    """
    Validate that an entity-activity pair makes sense.

    Args:
        entity: Entity text
        activity: Activity text

    Returns:
        True if valid pair
    """
    # Basic validation
    if not entity or not activity:
        return False

    # Entity should not be too short or generic
    if len(entity) < 3 or entity.lower() in ['el', 'la', 'los', 'las', 'un', 'una']:
        return False

    # Activity should be a reasonable verb
    return not len(activity) < 3

class FinancialAuditor:
    """
    Financial traceability and auditing
    """

    def __init__(self, config: ConfigLoader) -> None:
        self.logger = logging.getLogger(self.__class__.__name__)
        self.config = config
        self.financial_data: dict[str, dict[str, float]] = {}
        self.unit_costs: dict[str, float] = {}

```

```

self.successful_parses = 0
self.failed_parses = 0
self.d3_q3_analysis: dict[str, Any] = {} # Harmonic Front 3 - D3-Q3 metrics

@calibrated_method("saaaaaaa.analysis.derek_beach.FinancialAuditor.trace_financial_allocation")
def trace_financial_allocation(self, tables: list[pd.DataFrame],
                               nodes: dict[str, MetaNode],
                               graph: nx.DiGraph | None = None) -> dict[str, float]:
    """Trace financial allocations to programs/goals

Harmonic Front 3 - Enhancement 5: Single-Case Counterfactual Budget Check
Incorporates logic from single-case counterfactuals to test minimal sufficiency.
For D3-Q3 (Traceability/Resources): checks if resource X (BPIN code) were removed,
would the mechanism (Product) still execute? Only boosts budget traceability score
if allocation is tied to a specific project.
"""

for i, table in enumerate(tables):
    try:
        self.logger.info(f"Procesando tabla financiera {i + 1}/{len(tables)}")
        self._process_financial_table(table, nodes)
        self.successful_parses += 1
    except Exception as e:
        self.logger.error(f"Error procesando tabla financiera {i + 1}: {e}")
        self.failed_parses += 1
    continue

# HARMONIC FRONT 3 - Enhancement 5: Counterfactual sufficiency check
if graph is not None:
    self._perform_counterfactual_budget_check(nodes, graph)

self.logger.info(f"Asignaciones financieras trazadas: {len(self.financial_data)}")
self.logger.info(f"Tablas parseadas exitosamente: {self.successful_parses}, "
                f"Fallidas: {self.failed_parses}")
return self.unit_costs

@calibrated_method("saaaaaaa.analysis.derek_beach.FinancialAuditor._process_financial_table")
def _process_financial_table(self, table: pd.DataFrame,
                            nodes: dict[str, MetaNode]) -> None:
    """Process a single financial table"""
    # Try to identify relevant columns
    amount_pattern = re.compile(
        self.config.get('patterns.financial_headers', r'PRESUPUESTO|VALOR|MONTO'),
        re.IGNORECASE
    )
    program_pattern = re.compile(r'PROGRAMA|META|CÓDIGO', re.IGNORECASE)

    amount_col = None
    program_col = None

    # Search in column names
    for col in table.columns:
        col_str = str(col)
        if amount_pattern.search(col_str) and not amount_col:
            amount_col = col
        if program_pattern.search(col_str) and not program_col:
            program_col = col

    # If not found in column names, search in first row
    if not amount_col or not program_col:
        first_row = table.iloc[0]
        for i, val in enumerate(first_row):
            val_str = str(val)
            if amount_pattern.search(val_str) and not amount_col:
                amount_col = i
            table.columns = table.iloc[0]
            table = table[1:]

```

```

if program_pattern.search(val_str) and not program_col:
    program_col = i
    table.columns = table.iloc[0]
    table = table[1:]

if amount_col is None or program_col is None:
    self.logger.warning("No se encontraron columnas financieras relevantes")
    return

for _, row in table.iterrows():
    try:
        program_id = str(row[program_col]).strip().upper()
        amount = self._parse_amount(row[amount_col])

        if amount and program_id:
            matched_node = self._match_program_to_node(program_id, nodes)
            if matched_node:
                self.financial_data[matched_node] = {
                    'allocation': amount,
                    'source': 'budget_table'
                }

            # Update node
            nodes[matched_node].financial_allocation = amount

            # Calculate unit cost if possible
            node = nodes.get(matched_node)
            if node and node.target:
                try:
                    target_val = float(str(node.target).replace(',', '').replace('%', ""))
                    if target_val > 0:
                        unit_cost = amount / target_val
                        self.unit_costs[matched_node] = unit_cost
                        nodes[matched_node].unit_cost = unit_cost
                except (ValueError, TypeError):
                    pass
    except Exception as e:
        self.logger.debug(f"Error procesando fila financiera: {e}")
        continue

@calibrated_method("saaaaaa.analysis.derek_beach.FinancialAuditor._parse_amount")
def _parse_amount(self, value: Any) -> float | None:
    """Parse monetary amount from various formats"""
    if pd.isna(value):
        return None

    try:
        clean_value = str(value).replace('$', '').replace(',', '').replace(' ', '').replace('.', '')
        # Handle millions/thousands notation
        if 'M' in clean_value.upper() or 'MILLONES' in clean_value.upper():
            clean_value = clean_value.upper().replace('M', '').replace('MILLONES', '')
            return float(clean_value) * 1_000_000
        return float(clean_value)
    except (ValueError, TypeError):
        return None

@calibrated_method("saaaaaa.analysis.derek_beach.FinancialAuditor._match_program_to_node")
def _match_program_to_node(self, program_id: str,
                           nodes: dict[str, MetaNode]) -> str | None:
    """Match program ID to existing node using fuzzy matching

Enhanced for D1-Q3 / D3-Q3 Financial Traceability:
- Implements confidence penalty if fuzzy match ratio < 100
- Reduces node.financial_allocation confidence by 15% for imperfect matches

```

- Tracks match quality for overall financial traceability scoring

"""

```

if program_id in nodes:
    # Perfect match - no penalty
    return program_id

# Try fuzzy matching
best_match = process.extractOne(
    program_id,
    nodes.keys(),
    scorer=fuzz.ratio,
    score_cutoff=80
)

if best_match:
    matched_node_id = best_match[0]
    match_ratio = best_match[1]

    # D1-Q3 / D3-Q3: Apply confidence penalty for non-perfect matches
    if match_ratio < 100:
        penalty_factor = get_parameter_loader().get("saaaaaa.analysis.derek_beach.
FinancialAuditor._match_program_to_node").get("penalty_factor", 0.85) # Refactored
        node = nodes[matched_node_id]

        # Track original allocation before penalty
        if not hasattr(node, '_original_financial_allocation'):
            node._original_financial_allocation = node.financial_allocation

        # Apply penalty to financial allocation confidence
        if node.financial_allocation:
            penalized_allocation = node.financial_allocation * penalty_factor
            self.logger.debug(
                f"Fuzzy match penalty applied to {matched_node_id}: "
                f"ratio={match_ratio}, penalty={penalty_factor:.2f}, "
                f"allocation {node.financial_allocation:.0f} ->
{penalized_allocation:.0f}"
            )
            node.financial_allocation = penalized_allocation

        # Store match confidence for D1-Q3 / D3-Q3 scoring
        if not hasattr(node, 'financial_match_confidence'):
            node.financial_match_confidence = match_ratio / 10get_parameter_loader
            () .get("saaaaaa.analysis.derek_beach.FinancialAuditor._match_program_to_node").get("auto_p
aram_L2098_70", 0.0)
        else:
            # Average if multiple matches
            node.financial_match_confidence = (node.financial_match_confidence +
            match_ratio / 10get_parameter_loader().get("saaaaaa.analysis.derek_beach.FinancialAuditor.
            _match_program_to_node").get("auto_param_L2101_105", 0.0)) / 2

    return matched_node_id

return None

```

@calibrated\_method("saaaaaa.analysis.derek\_beach.FinancialAuditor.\_perform\_counterfact
ual\_budget\_check")

```

def _perform_counterfactual_budget_check(self, nodes: dict[str, MetaNode],
                                         graph: nx.DiGraph) -> None:
"""

```

Harmonic Front 3 - Enhancement 5: Counterfactual Sufficiency Test for D3-Q3

Tests minimal sufficiency: if resource X (BPIN code) were removed, would the mechanism (Product) still execute? Only boosts budget traceability score if allocation is tied to a specific project.

For D3-Q3 (Traceability/Resources): ensures funding is necessary for the mechanism and prevents false positives from generic or disconnected budget entries.

"""

```

d3_q3_scores = {}

for node_id, node in nodes.items():
    if node.type != 'producto':
        continue

    # Check if node has financial allocation
    has_budget = node.financial_allocation is not None and
node.financial_allocation > 0

    # Check if node has entity-activity (mechanism)
    has_mechanism = node.entity_activity is not None

    # Check if node has dependencies (successors in graph)
    successors = list(graph.successors(node_id)) if graph.has_node(node_id) else
[]

    has_dependencies = len(successors) > 0

    # Counterfactual test: Would mechanism still execute without this budget?
    # Check if there are alternative funding sources or generic allocations
    financial_source = self.financial_data.get(node_id, {}).get('source',
'unknown')
    is_specific_allocation = financial_source == 'budget_table' # From specific
table entry

    # Calculate counterfactual necessity score
    # High score = budget is necessary for execution
    # Low score = budget may be generic/disconnected
    necessity_score = get_parameter_loader().get("saaaaaaaa.analysis.derek_beach.Fin
ancialAuditor._perform_counterfactual_budget_check").get("necessity_score", 0.0) #
Refactored

    if has_budget and has_mechanism:
        necessity_score += get_parameter_loader().get("saaaaaaaa.analysis.derek_beac
h.FinancialAuditor._perform_counterfactual_budget_check").get("auto_param_L2147_35", 0.40)
        # Budget + mechanism present

    if has_budget and has_dependencies:
        necessity_score += get_parameter_loader().get("saaaaaaaa.analysis.derek_beac
h.FinancialAuditor._perform_counterfactual_budget_check").get("auto_param_L2150_35", 0.30)
        # Budget supports downstream goals

    if is_specific_allocation:
        necessity_score += get_parameter_loader().get("saaaaaaaa.analysis.derek_beac
h.FinancialAuditor._perform_counterfactual_budget_check").get("auto_param_L2153_35", 0.30)
        # Specific allocation (not generic)

    # D3-Q3 quality criteria
    d3_q3_quality = 'insuficiente'
    if necessity_score >= get_parameter_loader().get("saaaaaaaa.analysis.derek_beach
.FinancialAuditor._perform_counterfactual_budget_check").get("auto_param_L2157_34", 0.85):
        d3_q3_quality = 'excelente'
    elif necessity_score >= get_parameter_loader().get("saaaaaaaa.analysis.derek_bea
ch.FinancialAuditor._perform_counterfactual_budget_check").get("auto_param_L2159_36",
0.70):
        d3_q3_quality = 'bueno'
    elif necessity_score >= get_parameter_loader().get("saaaaaaaa.analysis.derek_bea
ch.FinancialAuditor._perform_counterfactual_budget_check").get("auto_param_L2161_36",
0.50):
        d3_q3_quality = 'aceptable'

    d3_q3_scores[node_id] = {
        'necessity_score': necessity_score,
        'd3_q3_quality': d3_q3_quality,
        'has_budget': has_budget,
        'has_mechanism': has_mechanism,
        'has_dependencies': has_dependencies,
        'is_specific_allocation': is_specific_allocation,
    }

```

```

'counterfactual_sufficient': necessity_score < get_parameter_loader().get(
    "saaaaaaa.analysis.derek_beach.FinancialAuditor._perform_counterfactual_budget_check").get(
    "auto_param_L2171_63", 0.50), # Would still execute without budget
    'budget_necessary': necessity_score >= get_parameter_loader().get("saaaaaaa
.analysis.derek_beach.FinancialAuditor._perform_counterfactual_budget_check").get("auto_pa
ram_L2172_55", 0.70) # Budget is necessary
}

# Store in node for later retrieval
node.audit_flags = node.audit_flags or []
if necessity_score < get_parameter_loader().get("saaaaaaa.analysis.derek_beach.
FinancialAuditor._perform_counterfactual_budget_check").get("auto_param_L2177_33", 0.50):
    node.audit_flags.append('budget_not_necessary')
    self.logger.warning(
        f"D3-Q3: {node_id} may execute without allocated budget
(score={necessity_score:.2f})")
    elif necessity_score >= get_parameter_loader().get("saaaaaaa.analysis.derek_bea
ch.FinancialAuditor._perform_counterfactual_budget_check").get("auto_param_L2181_36",
0.85):
        node.audit_flags.append('budget_well_traced')
        self.logger.info(f"D3-Q3: {node_id} has well-traced, necessary budget
(score={necessity_score:.2f})")

# Store aggregate D3-Q3 metrics
self.d3_q3_analysis = {
    'node_scores': d3_q3_scores,
    'total_products_analyzed': len(d3_q3_scores),
    'well_traced_count': sum(1 for s in d3_q3_scores.values() if
s['d3_q3_quality'] == 'excelente'),
    'average_necessity_score': sum(s['necessity_score'] for s in
d3_q3_scores.values()) / max(len(d3_q3_scores),
1)
}

self.logger.info(f"D3-Q3 Counterfactual Budget Check completed: "
    f"{self.d3_q3_analysis['well_traced_count']}/{len(d3_q3_scores)}"
)
"""

    f"products with excellent traceability")

```

@calibrated\_method("saaaaaaa.analysis.derek\_beach.FinancialAuditor.\_calculate\_sufficiency")

def \_calculate\_sufficiency(self, allocation: float, target: float) -> float:

"""

Calculate if financial allocation is sufficient for target.

Args:

allocation: Financial allocation amount

target: Target value

Returns:

Sufficiency ratio (get\_parameter\_loader().get("saaaaaaa.analysis.derek\_beach.Fi
nancialAuditor.\_calculate\_sufficiency").get("auto\_param\_L2208\_31", 1.0) = exactly
sufficient, >get\_parameter\_loader().get("saaaaaaa.analysis.derek\_beach.FinancialAuditor.\_ca
lculate\_sufficiency").get("auto\_param\_L2208\_58", 1.0) = oversufficient)

"""

if not target or target == 0:

return get\_parameter\_loader().get("saaaaaaa.analysis.derek\_beach.FinancialAudit
or.\_calculate\_sufficiency").get("auto\_param\_L2211\_19", 0.0)

# Calculate unit cost implied by allocation and target
allocation / target

# Compare with historical/expected unit costs if available

# For now, return simple ratio

return allocation / target if target > 0 else get\_parameter\_loader().get("saaaaaaa.
analysis.derek\_beach.FinancialAuditor.\_calculate\_sufficiency").get("auto\_param\_L2218\_54",
0.0)

```

@calibrated_method("saaaaaaa.analysis.derek_beach.FinancialAuditor._detect_allocation_gaps")
def _detect_allocation_gaps(self, nodes: dict[str, MetaNode]) -> list[dict[str, Any]]:
    """
    Detect gaps in financial allocations.

    Args:
        nodes: Dictionary of nodes

    Returns:
        List of detected gaps
    """
    gaps = []

    for node_id, node in nodes.items():
        # Check for missing allocation
        if node.type in ['producto', 'programa'] and not node.financial_allocation:
            gaps.append({
                'node_id': node_id,
                'type': 'missing_allocation',
                'severity': 'high',
                'message': f"No financial allocation for {node.type} {node_id}"
            })

        # Check for insufficient allocation
        if node.financial_allocation and node.target:
            try:
                target_val = float(str(node.target).replace(',', '').replace('%', ''))
                if target_val > 0:
                    sufficiency = self._calculate_sufficiency(node.financial_allocation, target_val)
                    if sufficiency < get_parameter_loader().get("saaaaaaa.analysis.dere
k_beach.FinancialAuditor._detect_allocation_gaps").get("auto_param_L2249_41", 0.5):
                        gaps.append({
                            'node_id': node_id,
                            'type': 'insufficient_allocation',
                            'severity': 'medium',
                            'message': f"Low sufficiency ratio {sufficiency:.2f} for
{node_id}",
                            'sufficiency': sufficiency
                        })
            except (ValueError, TypeError):
                pass

    return gaps

```

```

@calibrated_method("saaaaaaa.analysis.derek_beach.FinancialAuditor._match_goal_to_budget")
def _match_goal_to_budget(self, goal_text: str, budget_entries: list[dict[str, Any]]) -> dict[str, Any] | None:
    """
    Match a goal to budget entries.
    """

```

```

    Args:
        goal_text: Goal text to match
        budget_entries: List of budget entries

    Returns:
        Best matching budget entry or None
    """
    if not budget_entries:
        return None

    # Extract potential identifiers from goal text
    goal_words = set(goal_text.lower().split())

    best_match = None

```



```

detector = PolicyContradictionDetectorV2(device='cpu')
quantitative_claims =
detector._extract_structured_quantitative_claims(node.text)
except Exception as e:
    self.logger.debug(f"Could not extract quantitative claims: {e}")

# Check baseline
baseline_valid = False
if not node.baseline or str(node.baseline).upper() in ['ND', 'POR DEFINIR',
'N/A', 'NONE']:
    result['errors'].append(f'Línea base no definida para {node_id}')
    result['passed'] = False
    node.rigor_status = 'débil'
    node.audit_flags.append('sin_linea_base')
else:
    baseline_valid = True
    # Cross-check baseline against quantitative claims (D3-Q1)
    if quantitative_claims:
        baseline_in_claims = any(
            claim.get('type') in ['indicator', 'target', 'percentage',
            'beneficiaries']
            for claim in quantitative_claims
        )
        if not baseline_in_claims:
            result['warnings'].append(f'Línea base no verificada en claims
cuantitativos para {node_id}')

# Check target
target_valid = False
if not node.target or str(node.target).upper() in ['ND', 'POR DEFINIR', 'N/A',
'NONE']:
    result['errors'].append(f'Meta no definida para {node_id}')
    result['passed'] = False
    node.rigor_status = 'débil'
    node.audit_flags.append('sin_meta')
else:
    target_valid = True
    # Cross-check target against quantitative claims (D3-Q1)
    if quantitative_claims:
        meta_in_claims = any(
            claim.get('type') == 'target' or 'meta' in claim.get('context',
            '').lower()
            for claim in quantitative_claims
        )
        if not meta_in_claims:
            result['warnings'].append(f'Meta no verificada en claims
cuantitativos para {node_id}')

# D3-Q1 Ficha Técnica compliance check for producto nodes
if node.type == 'producto':
    # Check if has all minimum DNP INDICATOR_STRUCTURE elements
    has_complete_ficha = (
        baseline_valid and
        target_valid and
        'sin_linea_base' not in node.audit_flags and
        'sin_meta' not in node.audit_flags
    )

    if has_complete_ficha and quantitative_claims:
        # Node passes D3-Q1 compliance
        producto_nodes_passed += 1
        result['recommendations'].append(f'D3-Q1 Ficha Técnica completa para
{node_id}')
    elif has_complete_ficha:
        # Has baseline/target but no quantitative claims verification
        producto_nodes_passed += get_parameter_loader().get("saaaaaaa.analysis.
derek_beach.OperationalizationAuditor.audit_evidence_traceability").get("auto_param_L2401_
45", 0.5) # Partial credit

```

```

        result['warnings'].append(f"D3-Q1 parcial: Ficha básica sin
verificación cuantitativa en {node_id}")

    # Check responsible entity
    if not node.responsible_entity:
        result['warnings'].append(f"Entidad responsable no identificada para
{node_id}")
        node.audit_flags.append('sin_responsable')

    # Check financial traceability
    if not node.financial_allocation:
        result['warnings'].append(f"Sin trazabilidad financiera para {node_id}")
        node.audit_flags.append('sin_presupuesto')

    # Set rigor status if passed all checks
    if result['passed'] and len(result['warnings']) == 0:
        node.rigor_status = 'fuerte'

    self.audit_results[node_id] = result

# Calculate D3-Q1 compliance score
if producto_nodes_count > 0:
    d3_q1_compliance_pct = (producto_nodes_passed / producto_nodes_count) * 100
    self.logger.info(f"D3-Q1 Ficha Técnica Compliance: {d3_q1_compliance_pct:.1f}%
"
                    f"\n({producto_nodes_passed}/{producto_nodes_count}
productos)")

    if d3_q1_compliance_pct >= 80:
        self.logger.info("D3-Q1 Score: EXCELENTE (≥80% productos con Ficha Técnica
completa)")
    elif d3_q1_compliance_pct >= 60:
        self.logger.info("D3-Q1 Score: BUENO (60-80% compliance)")
    else:
        self.logger.warning("D3-Q1 Score: INSUFICIENTE (<60% compliance)")

    passed_count = sum(1 for r in self.audit_results.values() if r['passed'])
    self.logger.info(f"Auditoría de trazabilidad: {passed_count}/{len(nodes)} nodos
aprobados")

return self.audit_results

@calibrated_method("saaaaaa.analysis.derek_beach.OperationalizationAuditor.audit_seque
nce_logic")
def audit_sequence_logic(self, graph: nx.DiGraph) -> list[str]:
    """Audit logical sequence of activities"""
    warnings = []

    # Group nodes by program
    programs: dict[str, list[str]] = defaultdict(list)
    for node_id in graph.nodes():
        node_data = graph.nodes[node_id]
        if node_data.get('type') == 'programa':
            for successor in graph.successors(node_id):
                if graph.nodes[successor].get('type') == 'producto':
                    programs[node_id].append(successor)

    # Check sequence within each program
    for program_id, product_goals in programs.items():
        if len(product_goals) < 2:
            continue

        activities = []
        for goal_id in product_goals:
            node = graph.nodes[goal_id]
            ea = node.get('entity_activity')
            if ea and isinstance(ea, dict):
                verb = ea.get('verb_lemma', "")

```

```

sequence_num = self.verb_sequences.get(verb, 999)
activities.append((goal_id, verb, sequence_num))

# Check for sequence violations
activities.sort(key=lambda x: x[2])
for i in range(len(activities) - 1):
    if activities[i][2] > activities[i + 1][2]:
        warning = (f"Violación de secuencia en {program_id}: "
                   f"{activities[i][1]} ({activities[i][0]}) "
                   f"antes de {activities[i + 1][1]} ({activities[i +
1][0]})")
        warnings.append(warning)
        self.logger.warning(warning)

self.sequence_warnings = warnings
return warnings

@calibrated_method("saaaaaaa.analysis.derek_beach.OperationalizationAuditor.bayesian_co
unterfactual_audit")
def bayesian_counterfactual_audit(self, nodes: dict[str, MetaNode],
                                  graph: nx.DiGraph,
                                  historical_data: dict[str, Any] | None = None,
                                  pdet_alignment: float | None = None) -> dict[str,
Any]:
    """
    AGUJA III: El Auditor Contrafactual Bayesiano
    Perform counterfactual audit using Bayesian causal reasoning

    Harmonic Front 3: Enhanced to consume pdet_alignment scores for D4-Q5 and D5-Q4
    integration
    """
    self.logger.info("Iniciando auditoría contrafactual Bayesiana...")

    # Build implicit Structural Causal Model (SCM)
    scm_dag = self._build_normative_dag()

    # Initialize historical priors
    if historical_data is None:
        historical_data = self._get_default_historical_priors()

    # Audit results by layers
    layer1_results = self._audit_direct_evidence(nodes, scm_dag, historical_data)
    layer2_results = self._audit_causal_implications(nodes, graph, layer1_results)
    layer3_results = self._audit_systemic_risk(nodes, graph, layer1_results,
                                             layer2_results, pdet_alignment)

    # Generate optimal remediation recommendations
    recommendations = self._generate_optimal_remediations(
        layer1_results, layer2_results, layer3_results
    )

    audit_report = {
        'direct_evidence': layer1_results,
        'causal_implications': layer2_results,
        'systemic_risk': layer3_results,
        'recommendations': recommendations,
        'summary': {
            'total_nodes': len(nodes),
            'critical_omissions': sum(1 for r in layer1_results.values()
                                      if r.get('omission_severity') == 'critical'),
            'expected_success_probability': layer3_results.get('success_probability',
get_parameter_loader().get("saaaaaaa.analysis.derek_beach.OperationalizationAuditor.bayesi
n_counterfactual_audit").get("auto_param_L2518_90", 0.0)),
            'risk_score': layer3_results.get('risk_score', get_parameter_loader().get(
"saaaaaaa.analysis.derek_beach.OperationalizationAuditor.bayesian_counterfactual_audit").ge
t("auto_param_L2519_63", 0.0))
        }
    }
}

```

```

        self.logger.info(f"Auditoría contrafactual completada: "
                       f"\'{audit_report['summary']['critical_omissions']}\' omisiones
                       críticas detectadas")

    return audit_report

@calibrated_method("saaaaaaa.analysis.derek_beach.OperationalizationAuditor._build_normative_dag")
def _build_normative_dag(self) -> nx.DiGraph:
    """Build normative DAG of expected relationships in well-formed plans"""
    dag = nx.DiGraph()

    # Define normative structure
    # Each goal type should have these attributes
    dag.add_node('baseline', type='required_attribute')
    dag.add_node('target', type='required_attribute')
    dag.add_node('entity', type='required_attribute')
    dag.add_node('budget', type='recommended_attribute')
    dag.add_node('mechanism', type='recommended_attribute')
    dag.add_node('timeline', type='optional_attribute')
    dag.add_node('risk_factors', type='optional_attribute')

    # Causal relationships
    dag.add_edge('baseline', 'target', relation='defines_gap')
    dag.add_edge('entity', 'mechanism', relation='executes')
    dag.add_edge('budget', 'mechanism', relation='enables')
    dag.add_edge('mechanism', 'target', relation='achieves')
    dag.add_edge('risk_factors', 'target', relation='threatens')

    return dag

@calibrated_method("saaaaaaa.analysis.derek_beach.OperationalizationAuditor._get_default_historical_priors")
def _get_default_historical_priors(self) -> dict[str, Any]:
    """Get default historical priors if no data is available"""
    return {
        'entity_presence_success_rate': get_parameter_loader().get("saaaaaaa.analysis.derek_beach.OperationalizationAuditor._get_default_historical_priors").get("auto_param_L255_6_44", 0.94),
        'baseline_presence_success_rate': get_parameter_loader().get("saaaaaaa.analysis.derek_beach.OperationalizationAuditor._get_default_historical_priors").get("auto_param_L255_557_46", 0.89),
        'target_presence_success_rate': get_parameter_loader().get("saaaaaaa.analysis.derek_beach.OperationalizationAuditor._get_default_historical_priors").get("auto_param_L255_8_44", 0.92),
        'budget_presence_success_rate': get_parameter_loader().get("saaaaaaa.analysis.derek_beach.OperationalizationAuditor._get_default_historical_priors").get("auto_param_L255_9_44", 0.78),
        'mechanism_presence_success_rate': get_parameter_loader().get("saaaaaaa.analysis.derek_beach.OperationalizationAuditor._get_default_historical_priors").get("auto_param_L2560_47", 0.65),
        'complete_documentation_success_rate': get_parameter_loader().get("saaaaaaa.analysis.derek_beach.OperationalizationAuditor._get_default_historical_priors").get("auto_param_L2561_51", 0.82),
        'node_type_success_rates': {
            'producto': get_parameter_loader().get("saaaaaaa.analysis.derek_beach.OperationalizationAuditor._get_default_historical_priors").get("auto_param_L2563_28", 0.85),
            'resultado': get_parameter_loader().get("saaaaaaa.analysis.derek_beach.OperationalizationAuditor._get_default_historical_priors").get("auto_param_L2564_29", 0.72),
            'impacto': get_parameter_loader().get("saaaaaaa.analysis.derek_beach.OperationalizationAuditor._get_default_historical_priors").get("auto_param_L2565_27", 0.58)
        }
    }

@calibrated_method("saaaaaaa.analysis.derek_beach.OperationalizationAuditor._audit_direct_evidence")
def _audit_direct_evidence(self, nodes: dict[str, MetaNode],

```

```

        scm_dag: nx.DiGraph,
        historical_data: dict[str, Any]) -> dict[str, dict[str,
Any]]:
    """Layer 1: Audit direct evidence of required components

    Enhanced with highly specific Bayesian priors for rare evidence items.
    Example: D2-Q4 risk matrix, D5-Q5 unwanted effects are rare in poor PDMs.
    """
    results = {}

    # Load highly specific priors for rare evidence types
    # D2-Q4: Risk matrices are rare in poor PDMs (high probative value as Smoking Gun)
    rare_evidence_priors = {
        'risk_matrix': {
            'prior_alpha': 1.5, # Low alpha = rare occurrence
            'prior_beta': 12.0, # High beta = high failure rate when absent
            'keywords': ['matriz de riesgo', 'análisis de riesgo', 'gestión de
riesgo', 'riesgos identificados']
        },
        'unwanted_effects': {
            'prior_alpha': 1.8, # D5-Q5: Effects analysis is also rare
            'prior_beta': 1get_parameter_loader().get("aaaaaaa.analysis.derek_beach.Op
erationalizationAuditor._audit_direct_evidence").get("auto_param_L2590_31", 0.5),
            'keywords': ['efectos no deseados', 'efectos adversos', 'impactos
negativos',
            'consecuencias no previstas']
        },
        'theory_of_change': {
            'prior_alpha': 1.2,
            'prior_beta': 15.0,
            'keywords': ['teoría de cambio', 'teoría del cambio', 'cadena causal',
'modelo lógico']
        }
    }

    for node_id, node in nodes.items():
        omissions = []
        omission_probs = {}
        rare_evidence_found = {}

        # Check for rare, high-value evidence in node text
        node_text_lower = node.text.lower()
        for evidence_type, prior_config in rare_evidence_priors.items():
            if any(kw in node_text_lower for kw in prior_config['keywords']):
                # Rare evidence found! Strong Smoking Gun
                rare_evidence_found[evidence_type] = {
                    'prior_alpha': prior_config['prior_alpha'],
                    'prior_beta': prior_config['prior_beta'],
                    'posterior_strength': prior_config['prior_alpha'] / (
                        prior_config['prior_alpha'] +
                        prior_config['prior_beta'])
                }
                self.logger.info(f"Rare evidence '{evidence_type}' found in {node_id}
- Strong Smoking Gun!")

        # Check baseline
        if not node.baseline or str(node.baseline).upper() in ['ND', 'POR DEFINIR',
'N/A', 'NONE']:
            p_failure_given_omission = get_parameter_loader().get("aaaaaaa.analysis.de
rek_beach.OperationalizationAuditor._audit_direct_evidence").get("auto_param_L2621_43",
1.0) - historical_data.get('baseline_presence_success_rate', get_parameter_loader().get("s
aaaaaaa.analysis.derek_beach.OperationalizationAuditor._audit_direct_evidence").get("auto_p
aram_L2621_103", 0.89))
            omissions.append('baseline')
            omission_probs['baseline'] = p_failure_given_omission

        # Check target
        if not node.target or str(node.target).upper() in ['ND', 'POR DEFINIR', 'N/A',

```

'NONE']:

```
p_failure_given_omission = get_parameter_loader().get("saaaaaaa.analysis.de
rek_beach.OperationalizationAuditor._audit_direct_evidence").get("auto_param_L2627_43",
1.0) - historical_data.get('target_presence_success_rate', get_parameter_loader().get("sa
aaaa.analysis.derek_beach.OperationalizationAuditor._audit_direct_evidence").get("auto_par
am_L2627_101", 0.92))
omissions.append('target')
omission_probs['target'] = p_failure_given_omission

# Check entity
if not node.responsible_entity:
    p_failure_given_omission = get_parameter_loader().get("saaaaaaa.analysis.de
rek_beach.OperationalizationAuditor._audit_direct_evidence").get("auto_param_L2633_43",
1.0) - historical_data.get('entity_presence_success_rate', get_parameter_loader().get("sa
aaaa.analysis.derek_beach.OperationalizationAuditor._audit_direct_evidence").get("auto_par
am_L2633_101", 0.94))
    omisions.append('entity')
    omission_probs['entity'] = p_failure_given_omission

# Check budget
if not node.financial_allocation:
    p_failure_given_omission = get_parameter_loader().get("saaaaaaa.analysis.de
rek_beach.OperationalizationAuditor._audit_direct_evidence").get("auto_param_L2639_43",
1.0) - historical_data.get('budget_presence_success_rate', get_parameter_loader().get("sa
aaaa.analysis.derek_beach.OperationalizationAuditor._audit_direct_evidence").get("auto_par
am_L2639_101", 0.78))
    omisions.append('budget')
    omission_probs['budget'] = p_failure_given_omission

# Check mechanism
if not node.entity_activity:
    p_failure_given_omission = get_parameter_loader().get("saaaaaaa.analysis.de
rek_beach.OperationalizationAuditor._audit_direct_evidence").get("auto_param_L2645_43",
1.0) - historical_data.get('mechanism_presence_success_rate', get_parameter_loader().get("sa
aaaaaaa.analysis.derek_beach.OperationalizationAuditor._audit_direct_evidence").get("auto_
param_L2645_104", 0.65))
    omisions.append('mechanism')
    omission_probs['mechanism'] = p_failure_given_omission

# Determine severity
severity = 'none'
if omission_probs:
    max_failure_prob = max(omission_probs.values())
    if max_failure_prob > get_parameter_loader().get("saaaaaaa.analysis.derek_b
each.OperationalizationAuditor._audit_direct_evidence").get("auto_param_L2653_38", 0.15):
        severity = 'critical'
    elif max_failure_prob > get_parameter_loader().get("saaaaaaa.analysis.derek
_beach.OperationalizationAuditor._audit_direct_evidence").get("auto_param_L2655_40",
0.10):
        severity = 'high'
    elif max_failure_prob > get_parameter_loader().get("saaaaaaa.analysis.derek
_beach.OperationalizationAuditor._audit_direct_evidence").get("auto_param_L2657_40",
0.05):
        severity = 'medium'
    else:
        severity = 'low'

results[node_id] = {
    'omissions': omissions,
    'omission_probabilities': omission_probs,
    'omission_severity': severity,
    'node_type': node.type,
    'rare_evidence_found': rare_evidence_found # Add rare evidence to results
}

return results
```

@calibrated\_method("saaaaaaa.analysis.derek\_beach.OperationalizationAuditor.\_audit\_caus

```

al_implications")
def _audit_causal_implications(self, nodes: dict[str, MetaNode],
                                graph: nx.DiGraph,
                                direct_evidence: dict[str, dict[str, Any]]) ->
dict[str, dict[str, Any]]:
    """Layer 2: Audit causal implications of omissions"""
    implications = {}

    for node_id, node in nodes.items():
        node_omissions = direct_evidence[node_id]['omissions']
        causal_effects = {}

        # If baseline is missing
        if 'baseline' in node_omissions:
            # P(target_miscalibrated | missing_baseline)
            causal_effects['target_miscalibration'] = {
                'probability': get_parameter_loader().get("saaaaaaa.analysis.derek_beac",
h.OperationalizationAuditor._audit_causal_implications").get("auto_param_L2687_35", 0.73),
                'description': 'Sin línea base, la meta probablemente está mal
calibrada'
            }

        # If entity and high budget are missing
        if 'entity' in node_omissions and node.financial_allocation and
node.financial_allocation > 1000000:
            causal_effects['implementation_failure'] = {
                'probability': get_parameter_loader().get("saaaaaaa.analysis.derek_beac",
h.OperationalizationAuditor._audit_causal_implications").get("auto_param_L2694_35", 0.89),
                'description': 'Alto presupuesto sin entidad responsable indica alto
riesgo de falla'
            }

        elif 'entity' in node_omissions:
            causal_effects['implementation_failure'] = {
                'probability': get_parameter_loader().get("saaaaaaa.analysis.derek_beac",
h.OperationalizationAuditor._audit_causal_implications").get("auto_param_L2699_35", 0.65),
                'description': 'Sin entidad responsable, la implementación es
incierta'
            }

        # If mechanism is missing
        if 'mechanism' in node_omissions:
            causal_effects['unclear_pathway'] = {
                'probability': get_parameter_loader().get("saaaaaaa.analysis.derek_beac",
h.OperationalizationAuditor._audit_causal_implications").get("auto_param_L2706_35", 0.70),
                'description': 'Sin mecanismo definido, la vía causal es opaca'
            }

        # Check downstream effects
        successors = list(graph.successors(node_id)) if graph.has_node(node_id) else
[]

        if node_omissions and successors:
            causal_effects['cascade_risk'] = {
                'probability': min(get_parameter_loader().get("saaaaaaa.analysis.derek_
beach.OperationalizationAuditor._audit_causal_implications").get("auto_param_L2714_39",
0.95), get_parameter_loader().get("saaaaaaa.analysis.derek_beach.OperationalizationAuditor.
_audit_causal_implications").get("auto_param_L2714_45", 0.4) + get_parameter_loader().get(
"saaaaaaa.analysis.derek_beach.OperationalizationAuditor._audit_causal_implications").get(
"auto_param_L2714_51", 0.1) * len(node_omissions)),
                'affected_nodes': successors,
                'description': f'Omisiones pueden afectar {len(successors)} nodos
dependientes'
            }

            implications[node_id] = {
                'causal_effects': causal_effects,
                'total_risk': sum(e['probability'] for e in causal_effects.values()) /
max(len(causal_effects), 1)
            }

```

```

return implications

@calibrated_method("saaaaaaa.analysis.derek_beach.OperationalizationAuditor._audit_systemic_risk")
def _audit_systemic_risk(self, nodes: dict[str, MetaNode],
                        graph: nx.DiGraph,
                        direct_evidence: dict[str, dict[str, Any]],
                        causal_implications: dict[str, dict[str, Any]],
                        pdet_alignment: float | None = None) -> dict[str, Any]:
    """
    AUDIT POINT 2.3: Policy Alignment Dual Constraint
    Layer 3: Calculate systemic risk from accumulated omissions
    """

```

Harmonic Front 3 - Enhancement 1: Alignment and Systemic Risk Linkage  
Incorporates Policy Alignment scores (PND, ODS, RRI) as variable in systemic risk.

For D5-Q4 (Riesgos Sistémicos) and D4-Q5 (Alineación):

- If `pdet_alignment ≤ get_parameter_loader().get("saaaaaaa.analysis.derek_beach.OperationalizationAuditor._audit_systemic_risk").get("auto_param_L2740_30", 0.60)`, applies 1.2x multiplier to `risk_score`
- Excelente on D5-Q4 requires `risk_score < get_parameter_loader().get("saaaaaaa.analysis.derek_beach.OperationalizationAuditor._audit_systemic_risk").get("auto_param_L2741_51", 0.10)`

Implements dual constraints integrating macro-micro causality per Lieberman 2015.

```

# Identify critical nodes (high centrality)
if graph.number_of_nodes() > 0:
    try:
        centrality = nx.betweenness_centrality(graph)
    except (nx.NetworkXError, ZeroDivisionError, Exception) as e:
        logging.warning(f"Failed to calculate betweenness centrality: {e}. Using default values.")
        centrality = dict.fromkeys(graph.nodes(), get_parameter_loader().get("saaaa.analysis.derek_beach.OperationalizationAuditor._audit_systemic_risk").get("auto_param_L2752_58", 0.5))
    else:
        centrality = {}

# Calculate P(cascade_failure | omission_set)
critical_omissions = []
for node_id, evidence in direct_evidence.items():
    if evidence['omission_severity'] in ['critical', 'high']:
        node_centrality = centrality.get(node_id, get_parameter_loader().get("saaaa.analysis.derek_beach.OperationalizationAuditor._audit_systemic_risk").get("auto_param_L2760_58", 0.5))
        critical_omissions.append({
            'node_id': node_id,
            'severity': evidence['omission_severity'],
            'centrality': node_centrality,
            'omissions': evidence['omissions']
        })

```

# Calculate systemic risk

```

if critical_omissions:
    # Weighted by centrality
    risk_score = sum(
        (get_parameter_loader().get("saaaaaaa.analysis.derek_beach.OperationalizationAuditor._audit_systemic_risk").get("auto_param_L2772_17", 1.0) if om['severity'] == 'critical' else get_parameter_loader().get("saaaaaaa.analysis.derek_beach.OperationalizationAuditor._audit_systemic_risk").get("auto_param_L2772_58", 0.7)) * (om['centrality'] + get_parameter_loader().get("saaaaaaa.analysis.derek_beach.OperationalizationAuditor._audit_systemic_risk").get("auto_param_L2772_85", 0.1))
        for om in critical_omissions
    ) / len(nodes)
else:

```

```

risk_score = get_parameter_loader().get("saaaaaaa.analysis.derek_beach.Operatio
nalizationAuditor._audit_systemic_risk").get("risk_score", 0.0) # Refactored

# AUDIT POINT 2.3: Policy Alignment Dual Constraint
# If pdet_alignment ≤ get_parameter_loader().get("saaaaaaa.analysis.derek_beach.Ope
rationalizationAuditor._audit_systemic_risk").get("auto_param_L2779_30", 0.60), apply 1.2×
multiplier to risk_score
# This enforces integration between D4-Q5 (Alineación) and D5-Q4 (Riesgos
Sistémicos)
alignment_penalty_applied = False
alignment_threshold = get_parameter_loader().get("saaaaaaa.analysis.derek_beach.Ope
rationalizationAuditor._audit_systemic_risk").get("alignment_threshold", 0.6) # Refactored
alignment_multiplier = 1.2

if pdet_alignment is not None and pdet_alignment <= alignment_threshold:
    original_risk = risk_score
    risk_score = risk_score * alignment_multiplier
    alignment_penalty_applied = True
    self.logger.warning(
        f"ALIGNMENT PENALTY (D5-Q4): pdet_alignment={pdet_alignment:.2f} ≤
{alignment_threshold}, "
        f"risk_score escalated from {original_risk:.3f} to {risk_score:.3f} "
        f"(multiplier: {alignment_multiplier}x). Dual constraint per Lieberman
2015."
    )

# Calculate P(success | current_state)
total_omissions = sum(len(e['omissions']) for e in direct_evidence.values())
total_possible = len(nodes) * 5 # 5 key attributes per node
completeness = get_parameter_loader().get("saaaaaaa.analysis.derek_beach.Opera
lizationAuditor._audit_systemic_risk").get("auto_param_L2798_23", 1.0) - (total_omissions
/ max(total_possible, 1))

# Success probability (simplified Bayesian)
base_success_rate = get_parameter_loader().get("saaaaaaa.analysis.derek_beach.Opera
tionalizationAuditor._audit_systemic_risk").get("base_success_rate", 0.7) # Refactored
success_probability = base_success_rate * completeness

# D5-Q4 quality criteria check (AUDIT POINT 2.3)
# Excellent requires risk_score < get_parameter_loader().get("saaaaaaa.analysis.der
ek_beach.OperationalizationAuditor._audit_systemic_risk").get("auto_param_L2805_42", 0.10)
(matching ODS benchmarks per UN 2020)
d5_q4_quality = 'insuficiente'
risk_threshold_excellent = get_parameter_loader().get("saaaaaaa.analysis.derek_beac
h.OperationalizationAuditor._audit_systemic_risk").get("risk_threshold_excellent", 0.1) #
Refactored
risk_threshold_good = get_parameter_loader().get("saaaaaaa.analysis.derek_beach.Ope
rationalizationAuditor._audit_systemic_risk").get("risk_threshold_good", 0.2) # Refactored
risk_threshold_acceptable = get_parameter_loader().get("saaaaaaa.analysis.derek_be
ch.OperationalizationAuditor._audit_systemic_risk").get("risk_threshold_acceptable", 0.35)
# Refactored

if risk_score < risk_threshold_excellent:
    d5_q4_quality = 'excelente'
elif risk_score < risk_threshold_good:
    d5_q4_quality = 'bueno'
elif risk_score < risk_threshold_acceptable:
    d5_q4_quality = 'aceptable'

# Flag if alignment is causing quality failure
alignment_causing_failure = (
    alignment_penalty_applied and
    original_risk < risk_threshold_excellent and
    risk_score >= risk_threshold_excellent
)

return {
    'risk_score': min(get_parameter_loader().get("saaaaaaa.analysis.derek_beach.Ope

```

```

rationalizationAuditor._audit_systemic_risk").get("auto_param_L2826_30", 1.0),
risk_score),
    'success_probability': success_probability,
    'critical_omissions': critical_omissions,
    'completeness': completeness,
    'total_omissions': total_omissions,
    'pdet_alignment': pdet_alignment,
    'alignment_penalty_applied': alignment_penalty_applied,
    'alignment_threshold': alignment_threshold,
    'alignment_multiplier': alignment_multiplier,
    'alignment_causing_failure': alignment_causing_failure,
    'd5_q4_quality': d5_q4_quality,
    'd4_q5_alignment_score': pdet_alignment,
    'risk_thresholds': {
        'excellent': risk_threshold_excellent,
        'good': risk_threshold_good,
        'acceptable': risk_threshold_acceptable
    }
}

@calibrated_method("saaaaaa.analysis.derek_beach.OperationalizationAuditor._generate_optimal_remediations")
def _generate_optimal_remediations(self,
                                    direct_evidence: dict[str, dict[str, Any]],
                                    causal_implications: dict[str, dict[str, Any]],
                                    systemic_risk: dict[str, Any]) -> list[dict[str,
Any]]:
    """Generate prioritized remediation recommendations"""
    remediations = []

    # Calculate expected value of information for each remediation
    for node_id, evidence in direct_evidence.items():
        if not evidence['omissions']:
            continue

        for omission in evidence['omissions']:
            # Estimate impact
            omission_prob = evidence['omission_probabilities'].get(omission, get_param
eter_loader().get("saaaaaa.analysis.derek_beach.OperationalizationAuditor._generate_optima
l_remediations").get("auto_param_L2860_81", 0.1))
            causal_risk = causal_implications[node_id]['total_risk']

            # Expected value = P(failure_avoided) * Impact
            expected_value = omission_prob * (1 + causal_risk)

            # Effort estimate (simplified)
            effort_map = {
                'baseline': 3, # Moderate effort to research
                'target': 2, # Low effort to define
                'entity': 2, # Low effort to assign
                'budget': 4, # Higher effort to allocate
                'mechanism': 5 # Highest effort to design
            }
            effort = effort_map.get(omission, 3)

            # Priority = Expected Value / Effort
            priority = expected_value / effort

            remediations.append({
                'node_id': node_id,
                'omission': omission,
                'severity': evidence['omission_severity'],
                'expected_value': expected_value,
                'effort': effort,
                'priority': priority,
                'recommendation': self._get_remediation_text(omission, node_id)
            })

```

```

# Sort by priority (descending)
remediations.sort(key=lambda x: x['priority'], reverse=True)

return remediations

@calibrated_method("saaaaaa.analysis.derek_beach.OperationalizationAuditor._get_remediation_text")
def _get_remediation_text(self, omission: str, node_id: str) -> str:
    """Get specific remediation text for an omission"""
    texts = {
        'baseline': f"Definir línea base cuantitativa para {node_id} basada en diagnóstico actual",
        'target': f"Especificar meta cuantitativa alcanzable para {node_id} con horizonte temporal",
        'entity': f"Asignar entidad responsable clara para la ejecución de {node_id}",
        'budget': f"Asignar recursos presupuestarios específicos a {node_id}",
        'mechanism': f"Documentar mecanismo causal (Entidad-Actividad) para {node_id}"
    }
    return texts.get(omission, f"Completar {omission} para {node_id}")

```

```

@calibrated_method("saaaaaa.analysis.derek_beach.OperationalizationAuditor._perform_counterfactual_budget_check")
def _perform_counterfactual_budget_check(self, nodes: dict[str, MetaNode],
                                         graph: nx.DiGraph) -> dict[str, Any]:
    """
    Perform counterfactual budget check for operationalization audit.
    """

```

This method evaluates whether removing budget allocation would prevent goal execution, helping identify necessary vs. superfluous allocations.

Args:

nodes: Dictionary of meta nodes  
graph: Causal graph

Returns:

Dictionary with counterfactual analysis results

```

results = {
    'nodes_analyzed': 0,
    'budget_necessary': [],
    'budget_optional': [],
    'unallocated': []
}

```

```

for node_id, node in nodes.items():
    results['nodes_analyzed'] += 1

```

```

has_budget = node.financial_allocation and node.financial_allocation > 0
has_mechanism = node.entity_activity is not None
has_dependencies = len(list(graph.successors(node_id))) > 0 if
graph.has_node(node_id) else False

```

```

if not has_budget:
    results['unallocated'].append(node_id)
elif has_mechanism and has_dependencies:
    # Budget seems necessary for execution
    results['budget_necessary'].append(node_id)
else:
    # Budget may be optional or disconnected
    results['budget_optional'].append(node_id)

```

```

return results

```

```

class BayesianMechanismInference:
    """

```

AGUJA II: El Modelo Generativo de Mecanismos  
Hierarchical Bayesian model for causal mechanism inference

## F1.2 ARCHITECTURAL REFACTORING:

This class now integrates with refactored Bayesian engine components:

- BayesianPriorBuilder: Construye priors adaptativos (AGUJA I)
- BayesianSamplingEngine: Ejecuta MCMC sampling (AGUJA II)
- NecessitySufficiencyTester: Ejecuta Hoop Tests (AGUJA III)

The refactored components provide:

- Crystal-clear separation of concerns
- Trivial unit testing
- Explicit compliance with Fronts B and C

Legacy methods are preserved for backward compatibility.

```
def __init__(self, config: ConfigLoader, nlp_model: spacy.Language, **kwargs) -> None:
```

Initialize Bayesian Mechanism Inference engine.

Args:

config: Configuration loader instance

nlp\_model: spaCy NLP model for text processing

\*\*kwargs: Accepts additional keyword arguments for backward compatibility.

Unexpected arguments (e.g., 'causal\_hierarchy') are logged and

ignored.

Note:

This function signature has been made defensive to handle unexpected keyword arguments that may be passed due to interface drift.

....

# Log warning if unexpected kwargs are passed

if kwargs:

logging.getLogger(\_\_name\_\_).warning(

f"BayesianMechanismInference.\_\_init\_\_ received unexpected keyword

arguments: {list(kwargs.keys())}."

"These will be ignored. Expected signature: \_\_init\_\_(self, config:

ConfigLoader, nlp\_model: spacy.Language)"

)

self.logger = logging.getLogger(self.\_\_class\_\_.\_\_name\_\_)

self.config = config

self.nlp = nlp\_model

# F1.2: Initialize refactored Bayesian engine adapter if available

if REFACTORED\_BAYESIAN\_AVAILABLE:

try:

self.bayesian\_adapter = BayesianEngineAdapter(config, nlp\_model)

if self.bayesian\_adapter.is\_available():

self.logger.info("✓ Usando motor Bayesiano refactorizado (F1.2)")

self.\_log\_refactored\_components()

else:

self.bayesian\_adapter = None

except Exception as e:

self.logger.warning(f"Error inicializando motor refactorizado: {e}")

self.bayesian\_adapter = None

else:

self.bayesian\_adapter = None

# Load mechanism type hyperpriors from configuration (externalized)

self.mechanism\_type\_priors = {

'administrativo': self.config.get\_mechanism\_prior('administrativo'),

'tecnico': self.config.get\_mechanism\_prior('tecnico'),

'financiero': self.config.get\_mechanism\_prior('financiero'),

'politico': self.config.get\_mechanism\_prior('politico'),

'mixto': self.config.get\_mechanism\_prior('mixto')}

}

# Typical activity sequences by mechanism type

# These could also be externalized if needed for domain-specific customization

```

self.mechanism_sequences = {
    'administrativo': ['planificar', 'coordinar', 'gestionar', 'supervisar'],
    'tecnico': ['diagnosticar', 'diseñar', 'implementar', 'evaluar'],
    'financiero': ['asignar', 'ejecutar', 'auditar', 'reportar'],
    'politico': ['concertar', 'negociar', 'aprobar', 'promulgar']
}

# Track inferred mechanisms
self.inferred_mechanisms: dict[str, dict[str, Any]] = {}

@calibrated_method("saaaaaaa.analysis.derek_beach.BayesianMechanismInference._log_refactored_components")
def _log_refactored_components(self) -> None:
    """Log status of refactored Bayesian components (F1.2)"""
    if self.bayesian_adapter:
        status = self.bayesian_adapter.get_component_status()
        self.logger.info(" - BayesianPriorBuilder: " +
                         ("✓" if status['prior_builder_ready'] else "✗"))
        self.logger.info(" - BayesianSamplingEngine: " +
                         ("✓" if status['sampling_engine_ready'] else "✗"))
        self.logger.info(" - NecessitySufficiencyTester: " +
                         ("✓" if status['necessity_tester_ready'] else "✗"))

@calibrated_method("saaaaaaa.analysis.derek_beach.BayesianMechanismInference.infer_mechanisms")
def infer_mechanisms(self, nodes: dict[str, MetaNode],
                     text: str) -> dict[str, dict[str, Any]]:
    """
    Infer latent causal mechanisms using hierarchical Bayesian modeling

    HARMONIC FRONT 4 ENHANCEMENT:
    - Tracks mean mechanism_type uncertainty for quality criteria
    - Reports uncertainty reduction metrics
    """

    self.logger.info("Iniciando inferencia Bayesiana de mecanismos...")

    # Focus on 'producto' nodes which should have mechanisms
    product_nodes = {nid: n for nid, n in nodes.items() if n.type == 'producto'}

    # Track uncertainties for mean calculation
    mechanism_uncertainties = []

    for node_id, node in product_nodes.items():
        mechanism = self._infer_single_mechanism(node, text, nodes)
        self.inferred_mechanisms[node_id] = mechanism

        # Track mechanism type uncertainty for quality criteria
        if 'uncertainty' in mechanism:
            mech_type_uncertainty = mechanism['uncertainty'].get('mechanism_type', get_parameter_loader().get("saaaaaaa.analysis.derek_beach.BayesianMechanismInference.infer_mechanisms").get("auto_param_L3063_87", 1.0))
            mechanism_uncertainties.append(mech_type_uncertainty)

    # Calculate mean mechanism uncertainty for Harmonic Front 4 quality criteria
    mean_mech_uncertainty = (
        np.mean(mechanism_uncertainties) if mechanism_uncertainties else get_parameter_loader().get("saaaaaaa.analysis.derek_beach.BayesianMechanismInference.infer_mechanisms").get("auto_param_L3068_77", 1.0)
    )

    self.logger.info(f"Mecanismos inferidos: {len(self.inferred_mechanisms)}")
    self.logger.info(f"Mean mechanism_type uncertainty: {mean_mech_uncertainty:.4f}")

    # Store for reporting
    self._mean_mechanism_uncertainty = mean_mech_uncertainty

return self.inferred_mechanisms

```

```

    @calibrated_method("saaaaaaa.analysis.derek_beach.BayesianMechanismInference._infer_single_mechanism")
    def _infer_single_mechanism(self, node: MetaNode, text: str,
                                all_nodes: dict[str, MetaNode]) -> dict[str, Any]:
        """Infer mechanism for a single product node"""
        # Extract observations from text
        observations = self._extract_observations(node, text)

        # Level 3: Sample mechanism type from hyperprior
        mechanism_type_posterior = self._infer_mechanism_type(observations)

        # Level 2: Infer activity sequence given mechanism type
        sequence_posterior = self._infer_activity_sequence(
            observations, mechanism_type_posterior
        )

        # Level 1: Calculate coherence factor
        coherence_score = self._calculate_coherence_factor(
            node, observations, all_nodes
        )

        # Validation tests
        sufficiency = self._test_sufficiency(node, observations)
        necessity = self._test_necessity(node, observations)

        # Quantify uncertainty
        uncertainty = self._quantify_uncertainty(
            mechanism_type_posterior, sequence_posterior, coherence_score
        )

        # Detect gaps
        gaps = self._detect_gaps(node, observations, uncertainty)

    return {
        'mechanism_type': mechanism_type_posterior,
        'activity_sequence': sequence_posterior,
        'coherence_score': coherence_score,
        'sufficiency_test': sufficiency,
        'necessity_test': necessity,
        'uncertainty': uncertainty,
        'gaps': gaps,
        'observations': observations
    }

    @calibrated_method("saaaaaaa.analysis.derek_beach.BayesianMechanismInference._extract_observations")
    def _extract_observations(self, node: MetaNode, text: str) -> dict[str, Any]:
        """Extract textual observations related to the mechanism"""

        # Find node context in text
        node_pattern = re.escape(node.id)
        matches = list(re.finditer(node_pattern, text, re.IGNORECASE))

        observations = {
            'entity_activity': None,
            'verbs': [],
            'entities': [],
            'budget': node.financial_allocation,
            'context_snippets': []
        }

        if node.entity_activity:
            observations['entity_activity'] = {
                'entity': node.entity_activity.entity,
                'activity': node.entity_activity.activity,
                'verb_lemma': node.entity_activity.verb_lemma
            }

        # Extract context around node mentions

```

```

for match in matches[:3]: # Limit to first 3 occurrences
    start = max(0, match.start() - 300)
    end = min(len(text), match.end() + 300)
    context = text[start:end]

    # Process with spaCy
    doc = self.nlp(context)

    # Extract verbs
    verbs = [token.lemma_ for token in doc if token.pos_ == 'VERB']
    observations['verbs'].extend(verbs)

    # Extract entities
    entities = [ent.text for ent in doc.ents if ent.label_ in ['ORG', 'PER']]
    observations['entities'].extend(entities)

    observations['context_snippets'].append(context[:200])

return observations

@calibrated_method("saaaaaa.analysis.derek_beach.BayesianMechanismInference._infer_mec
hanism_type")
def _infer_mechanism_type(self, observations: dict[str, Any]) -> dict[str, float]:
    """Infer mechanism type using Bayesian updating"""
    # Start with hyperprior
    posterior = dict(self.mechanism_type_priors)

    # Get Laplace smoothing parameter from configuration
    laplace_smooth = self.config.get_bayesian_threshold('laplace_smoothing')

    # Update based on observed verbs
    observed_verbs = set(observations.get('verbs', []))

    if observed_verbs:
        for mech_type, typical_verbs in self.mechanism_sequences.items():
            # Count overlap
            overlap = len(observed_verbs.intersection(set(typical_verbs)))
            total = len(typical_verbs)

            if total > 0:
                # Likelihood: proportion of typical verbs observed with Laplace
                # smoothing
                likelihood = (overlap + laplace_smooth) / (total + 2 * laplace_smooth)

                # Bayesian update
                posterior[mech_type] *= likelihood

    # Update based on entity-activity
    if observations.get('entity_activity'):
        verb = observations['entity_activity'].get('verb_lemma', "")
        for mech_type, typical_verbs in self.mechanism_sequences.items():
            if verb in typical_verbs:
                posterior[mech_type] *= 1.5

    # Normalize
    total = sum(posterior.values())
    if total > 0:
        posterior = {k: v / total for k, v in posterior.items()}

return posterior

@calibrated_method("saaaaaa.analysis.derek_beach.BayesianMechanismInference._infer_act
ivity_sequence")
def _infer_activity_sequence(self, observations: dict[str, Any],
                            mechanism_type_posterior: dict[str, float]) -> dict[str,
Any]:
    """Infer activity sequence parameters"""
    # Get most likely mechanism type

```

```

best_type = max(mechanism_type_posterior.items(), key=lambda x: x[1])[0]
expected_sequence = self.mechanism_sequences.get(best_type, [])

observed_verbs = observations.get('verbs', [])

# Calculate transition probabilities (simplified Markov chain)
transitions = {}
for i in range(len(expected_sequence) - 1):
    current = expected_sequence[i]
    next_verb = expected_sequence[i + 1]

    # Check if transition is observed
    if current in observed_verbs and next_verb in observed_verbs:
        transitions[(current, next_verb)] = get_parameter_loader().get("saaaaaaa.an"
alysis.derek_beach.BayesianMechanismInference._infer_activity_sequence").get("auto_param_L"
3222_52", 0.85)
    else:
        transitions[(current, next_verb)] = get_parameter_loader().get("saaaaaaa.an"
alysis.derek_beach.BayesianMechanismInference._infer_activity_sequence").get("auto_param_L"
3224_52", 0.40)

return {
    'expected_sequence': expected_sequence,
    'observed_verbs': observed_verbs,
    'transition_probabilities': transitions,
    'sequence_completeness': len(set(observed_verbs) & set(expected_sequence)) /
max(len(expected_sequence), 1)
}

@calibrated_method("saaaaaaa.analysis.derek_beach.BayesianMechanismInference._calculate
_coherence_factor")
def _calculate_coherence_factor(self, node: MetaNode,
                               observations: dict[str, Any],
                               all_nodes: dict[str, MetaNode]) -> float:
    """Calculate mechanism coherence score"""
    coherence = get_parameter_loader().get("saaaaaaa.analysis.derek_beach.BayesianMecha
nismInference._calculate_coherence_factor").get("coherence", 0.0) # Refactored
    weights = []

    # Factor 1: Entity-Activity presence
    if observations.get('entity_activity'):
        coherence += get_parameter_loader().get("saaaaaaa.analysis.derek_beach.Bayesian
MechanismInference._calculate_coherence_factor").get("auto_param_L3243_25", 0.30)
        weights.append(get_parameter_loader().get("saaaaaaa.analysis.derek_beach.Bayesi
anMechanismInference._calculate_coherence_factor").get("auto_param_L3244_27", 0.30))

    # Factor 2: Budget consistency
    if observations.get('budget'):
        coherence += get_parameter_loader().get("saaaaaaa.analysis.derek_beach.Bayesian
MechanismInference._calculate_coherence_factor").get("auto_param_L3248_25", 0.20)
        weights.append(get_parameter_loader().get("saaaaaaa.analysis.derek_beach.Bayesi
anMechanismInference._calculate_coherence_factor").get("auto_param_L3249_27", 0.20))

    # Factor 3: Verb sequence completeness
    seq_info = observations.get('verbs', [])
    if seq_info:
        verb_score = min(len(seq_info) / 4.0, get_parameter_loader().get("saaaaaaa.anal
ysis.derek_beach.BayesianMechanismInference._calculate_coherence_factor").get("auto_param_
L3254_50", 1.0)) # Expect ~4 verbs
        coherence += verb_score * get_parameter_loader().get("saaaaaaa.analysis.derek_b
each.BayesianMechanismInference._calculate_coherence_factor").get("auto_param_L3255_38",
0.25)
        weights.append(get_parameter_loader().get("saaaaaaa.analysis.derek_beach.Bayesi
anMechanismInference._calculate_coherence_factor").get("auto_param_L3256_27", 0.25))

    # Factor 4: Entity presence
    if observations.get('entities'):
        coherence += get_parameter_loader().get("saaaaaaa.analysis.derek_beach.Bayesian

```

```

MechanismInference._calculate_coherence_factor").get("auto_param_L3260_25", 0.15)
    weights.append(get_parameter_loader().get("saaaaaaa.analysis.derek_beach.Bayesi
anMechanismInference._calculate_coherence_factor").get("auto_param_L3261_27", 0.15))

# Factor 5: Context richness
snippets = observations.get('context_snippets', [])
if snippets:
    coherence += get_parameter_loader().get("saaaaaaa.analysis.derek_beach.Bayesian
MechanismInference._calculate_coherence_factor").get("auto_param_L3266_25", 0.10)
    weights.append(get_parameter_loader().get("saaaaaaa.analysis.derek_beach.Bayesi
anMechanismInference._calculate_coherence_factor").get("auto_param_L3267_27", 0.10))

# Normalize by actual weights used
if weights:
    coherence = coherence / sum(weights) if sum(weights) > 0 else get_parameter_lo
ader().get("saaaaaaa.analysis.derek_beach.BayesianMechanismInference._calculate_coherence_f
actor").get("auto_param_L3271_74", 0.0)

return coherence

@calibrated_method("saaaaaaa.analysis.derek_beach.BayesianMechanismInference._test_suff
iciency")
def _test_sufficiency(self, node: MetaNode,
                      observations: dict[str, Any] -> dict[str, Any]:
    """Test if mechanism is sufficient to produce the outcome"""
    # Check if entity has capability
    has_entity = observations.get('entity_activity') is not None

    # Check if activities are present
    has_activities = len(observations.get('verbs', [])) >= 2

    # Check if resources are allocated
    has_resources = observations.get('budget') is not None

    sufficiency_score = (
        (get_parameter_loader().get("saaaaaaa.analysis.derek_beach.BayesianMechanis
mInference._test_sufficiency").get("auto_param_L3289_17", 0.4) if has_entity else get_para
meter_loader().get("saaaaaaa.analysis.derek_beach.BayesianMechanismInference._test_sufficie
ncy").get("auto_param_L3289_40", 0.0)) +
        (get_parameter_loader().get("saaaaaaa.analysis.derek_beach.BayesianMechanis
mInference._test_sufficiency").get("auto_param_L3290_17", 0.4) if has_activities else get_
parameter_loader().get("saaaaaaa.analysis.derek_beach.BayesianMechanismInference._test_suff
iciency").get("auto_param_L3290_44", 0.0)) +
        (get_parameter_loader().get("saaaaaaa.analysis.derek_beach.BayesianMechanis
mInference._test_sufficiency").get("auto_param_L3291_17", 0.2) if has_resources else get_p
arameter_loader().get("saaaaaaa.analysis.derek_beach.BayesianMechanismInference._test_suffi
ciency").get("auto_param_L3291_43", 0.0))
    )

    return {
        'score': sufficiency_score,
        'is_sufficient': sufficiency_score >= get_parameter_loader().get("saaaaaaa.anal
ysis.derek_beach.BayesianMechanismInference._test_sufficiency").get("auto_param_L3296_50",
        0.6),
        'components': {
            'entity': has_entity,
            'activities': has_activities,
            'resources': has_resources
        }
    }

@calibrated_method("saaaaaaa.analysis.derek_beach.BayesianMechanismInference._test_nece
ssity")
def _test_necessity(self, node: MetaNode,
                   observations: dict[str, Any] -> dict[str, Any]:
    """
AUDIT POINT 2.2: Mechanism Necessity Hoop Test

```

Test if mechanism is necessary by checking documented components:

- Entity (responsable)
- Activity (verb lemma sequence)
- Budget (presupuesto asignado)

Implements Beach 2017 Hoop Tests for necessity verification.

Per Falletti & Lynch 2009, Bayesian-deterministic hybrid boosts mechanism depth.

Returns:

Dict with 'is\_necessary', 'missing\_components', and remediation text

"""

```
# F1.2: Use refactored NecessitySufficiencyTester if available
if self.bayesian_adapter and self.bayesian_adapter.necessity_tester:
    try:
        return self.bayesian_adapter.test_necessity_from_observations(
            node.id,
            observations
        )
    except Exception as e:
        self.logger.warning(f"Error en tester refactorizado: {e}, usando legacy")

# AUDIT POINT 2.2: Enhanced necessity test with documented components
missing_components = []

# 1. Check Entity documentation
entities = observations.get('entities', [])
entity_activity = observations.get('entity_activity')

if not entity_activity or not entity_activity.get('entity'):
    missing_components.append('entity')
else:
    # Verify unique entity (not multiple conflicting entities)
    unique_entity = len(set(entities)) == 1 if entities else False
    if not unique_entity and len(entities) > 1:
        missing_components.append('unique_entity')

# 2. Check Activity documentation (verb lemma sequence)
verbs = observations.get('verbs', [])
if not verbs or len(verbs) < 1:
    missing_components.append('activity')
else:
    # Check for specific action verbs (not just generic ones)
    specific_verbs = [v for v in verbs if v in [
        'implementar', 'ejecutar', 'realizar', 'desarrollar',
        'construir', 'diseñar', 'planificar', 'coordinar',
        'gestionar', 'supervisar', 'controlar', 'auditar'
    ]]
    if not specific_verbs:
        missing_components.append('specific_activity')

# 3. Check Budget documentation
budget = observations.get('budget')
if budget is None or budget <= 0:
    missing_components.append('budget')

# Calculate necessity score
# All three components must be present for necessity=True
is_necessary = len(missing_components) == 0

# Calculate partial score for reporting
max_components = 3 # entity, activity, budget
present_components = max_components - len([
    c for c in missing_components if c in ['entity', 'activity', 'budget']])
necessity_score = present_components / max_components

result = {
    'score': necessity_score,
    'is_necessary': is_necessary,
```

```

'missing_components': missing_components,
'alternatives_likely': not is_necessary,
'hoop_test_passed': is_necessary
}

# Add remediation text if test fails
if not is_necessary:
    result['remediation'] = self._generate_necessity_remediation(node.id,
missing_components)

return result

@calibrated_method("saaaaaa.analysis.derek_beach.BayesianMechanismInference._generate_necessity_remediation")
def _generate_necessity_remediation(self, node_id: str, missing_components: list[str])
-> str:
    """Generate remediation text for failed necessity test"""
    component_descriptions = {
        'entity': 'entidad responsable claramente identificada',
        'unique_entity': 'una única entidad responsable (múltiples entidades detectadas)',
        'activity': 'secuencia de actividades documentada',
        'specific_activity': 'actividades específicas (no genéricas)',
        'budget': 'presupuesto asignado y cuantificado'
    }

    missing_desc = ', '.join([component_descriptions.get(c, c) for c in
missing_components])

    return (
        f'Mecanismo para {node_id} falla Hoop Test de necesidad (D6-Q2). '
        f'Componentes faltantes: {missing_desc}. '
        f'Se requiere documentar estos componentes necesarios para validar '
        f'la cadena causal según Beach 2017.'
    )

@calibrated_method("saaaaaa.analysis.derek_beach.BayesianMechanismInference._quantify_uncertainty")
def _quantify_uncertainty(self, mechanism_type_posterior: dict[str, float],
sequence_posterior: dict[str, Any],
coherence_score: float) -> dict[str, float]:
    """Quantify epistemic uncertainty"""
    # Entropy of mechanism type distribution
    mech_probs = list(mechanism_type_posterior.values())
    if mech_probs:
        mech_entropy = -sum(p * np.log(p + 1e-10) for p in mech_probs if p > 0)
        max_entropy = np.log(len(mech_probs))
        mech_uncertainty = mech_entropy / max_entropy if max_entropy > 0 else get_parameter_loader().get("saaaaaa.analysis.derek_beach.BayesianMechanismInference._quantify_uncertainty").get("auto_param_L3419_82", 1.0)
    else:
        mech_uncertainty = get_parameter_loader().get("saaaaaa.analysis.derek_beach.BayesianMechanismInference._quantify_uncertainty").get("mech_uncertainty", 1.0) # Refactored

    # Sequence completeness uncertainty
    seq_completeness = sequence_posterior.get('sequence_completeness', get_parameter_loader().get("saaaaaa.analysis.derek_beach.BayesianMechanismInference._quantify_uncertainty").get("auto_param_L3424_75", 0.0))
    seq_uncertainty = get_parameter_loader().get("saaaaaa.analysis.derek_beach.BayesianMechanismInference._quantify_uncertainty").get("auto_param_L3425_26", 1.0) -
seq_completeness

    # Coherence uncertainty
    coherence_uncertainty = get_parameter_loader().get("saaaaaa.analysis.derek_beach.BayesianMechanismInference._quantify_uncertainty").get("auto_param_L3428_32", 1.0) -
coherence_score

    # Combined uncertainty

```

```

total_uncertainty = (
    mech_uncertainty * get_parameter_loader().get("saaaaaaa.analysis.derek_beac
h.BayesianMechanismInference._quantify_uncertainty").get("auto_param_L3432_35", 0.4) +
    seq_uncertainty * get_parameter_loader().get("saaaaaaa.analysis.derek_beach
.BayesianMechanismInference._quantify_uncertainty").get("auto_param_L3433_34", 0.3) +
    coherence_uncertainty * get_parameter_loader().get("saaaaaaa.analysis.derek
_beach.BayesianMechanismInference._quantify_uncertainty").get("auto_param_L3434_40", 0.3)
)

return {
    'total': total_uncertainty,
    'mechanism_type': mech_uncertainty,
    'sequence': seq_uncertainty,
    'coherence': coherence_uncertainty
}
}

@calibrated_method("saaaaaaa.analysis.derek_beach.BayesianMechanismInference._detect_gaps")
def _detect_gaps(self, node: MetaNode, observations: dict[str, Any],
                 uncertainty: dict[str, float]) -> list[dict[str, str]]:
    """Detect documentation gaps based on uncertainty"""
    gaps = []

    # High total uncertainty
    if uncertainty['total'] > get_parameter_loader().get("saaaaaaa.analysis.derek_beach
.BayesianMechanismInference._detect_gaps").get("auto_param_L3451_34", 0.6):
        gaps.append({
            'type': 'high_uncertainty',
            'severity': 'high',
            'message': f'Mecanismo para {node.id} tiene alta incertidumbre
({uncertainty["total"]:.2f})',
            'suggestion': "Se requiere más documentación sobre el mecanismo causal"
        })

    # Missing entity
    if not observations.get('entity_activity'):
        gaps.append({
            'type': 'missing_entity',
            'severity': 'high',
            'message': f'No se especifica entidad responsable para {node.id}',
            'suggestion': "Especificar qué entidad ejecutará las actividades"
        })

    # Insufficient activities
    if len(observations.get('verbs', [])) < 2:
        gaps.append({
            'type': 'insufficient_activities',
            'severity': 'medium',
            'message': f'Pocas actividades documentadas para {node.id}',
            'suggestion': "Detallar las actividades necesarias para lograr el
producto"
        })

    # Missing budget
    if not observations.get('budget'):
        gaps.append({
            'type': 'missing_budget',
            'severity': 'medium',
            'message': f'Sin asignación presupuestaria para {node.id}',
            'suggestion': "Asignar recursos financieros al producto"
        })

    return gaps
}

@calibrated_method("saaaaaaa.analysis.derek_beach.BayesianMechanismInference._aggregate
_bayesian_confidence")
def _aggregate_bayesian_confidence(self, confidences: list[float]) -> float:
    """

```

Aggregate multiple Bayesian confidence values.

Args:

    confidences: List of confidence values to aggregate

Returns:

    Aggregated confidence value

"""

if not confidences:

```
    return get_parameter_loader().get("saaaaaaa.analysis.derek_beach.BayesianMechanismInference._aggregate_bayesian_confidence").get("auto_param_L3500_19", 0.5) # Default neutral confidence
```

```
return float(np.mean(confidences))
```

```
@calibrated_method("saaaaaaa.analysis.derek_beach.BayesianMechanismInference._build_transition_matrix")
```

```
def _build_transition_matrix(self, mechanism_type: str) -> np.ndarray:
```

"""

Build transition matrix for activity sequences.

Args:

    mechanism\_type: Type of mechanism

Returns:

    Transition probability matrix

"""

# Get typical sequence for this mechanism type

```
sequence = self.mechanism_sequences.get(mechanism_type, ['planificar', 'ejecutar', 'evaluar'])
```

```
n = len(sequence)
```

# Create a simple sequential transition matrix

```
matrix = np.zeros((n, n))
```

```
for i in range(n - 1):
```

```
    matrix[i, i + 1] = get_parameter_loader().get("saaaaaaa.analysis.derek_beach.BayesianMechanismInference._build_transition_matrix").get("auto_param_L3521_31", 0.7) #
```

High probability of next step

```
    matrix[i, i] = get_parameter_loader().get("saaaaaaa.analysis.derek_beach.BayesianMechanismInference._build_transition_matrix").get("auto_param_L3522_27", 0.2) #
```

Some probability of staying in same step

```
    if i < n - 2:
```

```
        matrix[i, i + 2] = get_parameter_loader().get("saaaaaaa.analysis.derek_beach.BayesianMechanismInference._build_transition_matrix").get("auto_param_L3524_35", 0.1) #
```

```
Small probability of skipping
```

```
    matrix[n - 1, n - 1] = get_parameter_loader().get("saaaaaaa.analysis.derek_beach.BayesianMechanismInference._build_transition_matrix").get("auto_param_L3525_31", 1.0) #
```

Final state is absorbing

```
return matrix
```

```
@calibrated_method("saaaaaaa.analysis.derek_beach.BayesianMechanismInference._calculate_type_transition_prior")
```

```
def _calculate_type_transition_prior(self, from_type: str, to_type: str) -> float:
```

"""

Calculate prior probability of transitioning between mechanism types.

Args:

    from\_type: Source mechanism type

    to\_type: Target mechanism type

Returns:

    Prior probability of transition

"""

# Same type has high probability

```
if from_type == to_type:
```

```
    return get_parameter_loader().get("saaaaaaa.analysis.derek_beach.BayesianMechanismInference._calculate_type_transition_prior").get("auto_param_L3543_19", 0.7)
```

```

# Related types have medium probability
related_pairs = [
    ('administrativo', 'politico'),
    ('tecnico', 'financiero'),
    ('financiero', 'administrativo'),
]
if (from_type, to_type) in related_pairs or (to_type, from_type) in related_pairs:
    return get_parameter_loader().get("saaaaaaa.analysis.derek_beach.BayesianMechanismInference._calculate_type_transition_prior").get("auto_param_L3552_19", 0.2)

# Unrelated types have low probability
return get_parameter_loader().get("saaaaaaa.analysis.derek_beach.BayesianMechanismInference._calculate_type_transition_prior").get("auto_param_L3555_15", 0.1)

@calibrated_method("saaaaaaa.analysis.derek_beach.BayesianMechanismInference._classify_mechanism_type")
def _classify_mechanism_type(self, observations: dict[str, Any]) -> str:
    """
    Classify mechanism type based on observations.

    Args:
        observations: Observed features

    Returns:
        Classified mechanism type
    """
    # Extract features
    verbs = observations.get('verbs', [])
    entities = observations.get('entities', [])
    budget = observations.get('budget')

    # Score each mechanism type
    scores = {}
    for mech_type, typical_verbs in self.mechanism_sequences.items():
        score = get_parameter_loader().get("saaaaaaa.analysis.derek_beach.BayesianMechanismInference._classify_mechanism_type").get("score", 0.0) # Refactored
        # Count matching verbs
        for verb in verbs:
            if any(tv in verb.lower() for tv in typical_verbs):
                score += get_parameter_loader().get("saaaaaaa.analysis.derek_beach.BayesianMechanismInference._classify_mechanism_type").get("auto_param_L3580_29", 1.0)
        scores[mech_type] = score

    # Adjust for budget presence (indicates financial mechanism)
    if budget and budget > 0:
        scores['financiero'] = scores.get('financiero', 0) + 2.0

    # Adjust for political/administrative entities
    for entity in entities:
        entity_lower = entity.lower()
        if any(word in entity_lower for word in ['alcaldía', 'consejo', 'gobernación']):
            scores['politico'] = scores.get('politico', 0) + get_parameter_loader().get("saaaaaaa.analysis.derek_beach.BayesianMechanismInference._classify_mechanism_type").get("auto_param_L3591_65", 1.0)
        if any(word in entity_lower for word in ['secretaría', 'dirección', 'oficina']):
            scores['administrativo'] = scores.get('administrativo', 0) + get_parameter_loader().get("saaaaaaa.analysis.derek_beach.BayesianMechanismInference._classify_mechanism_type").get("auto_param_L3593_77", 1.0)

    # Return type with highest score, or 'mixto' if tie
    if not scores or all(s == 0 for s in scores.values()):
        return 'mixto'

    max_score = max(scores.values())
    max_types = [t for t, s in scores.items() if s == max_score]

```

```

if len(max_types) > 1:
    return 'mixto'
return max_types[0]

class CausalInferenceSetup:
    """Prepare model for causal inference"""

    def __init__(self, config: ConfigLoader) -> None:
        self.logger = logging.getLogger(self.__class__.__name__)
        self.config = config
        self.goal_classification = config.get('lexicons.goal_classification', {})
        self.admin_keywords = config.get('lexicons.administrative_keywords', [])
        self.contextual_factors = config.get('lexicons.contextual_factors', [])

    @calibrated_method("saaaaaa.analysis.derek_beach.CausalInferenceSetup.classify_goal_dynamics")
    def classify_goal_dynamics(self, nodes: dict[str, MetaNode]) -> None:
        """Classify dynamics for each goal"""
        for node in nodes.values():
            text_lower = node.text.lower()

            for keyword, dynamics in self.goal_classification.items():
                if keyword in text_lower:
                    node.dynamics = cast("DynamicsType", dynamics)
                    self.logger.debug(f"Meta {node.id} clasificada como {node.dynamics}")
                    break

    @calibrated_method("saaaaaa.analysis.derek_beach.CausalInferenceSetup.assign_probative_value")
    def assign_probative_value(self, nodes: dict[str, MetaNode]) -> None:
        """Assign probative test types to nodes"""
        # Import INDICATOR_STRUCTURE from financiero_viviabilidad_tablas
        try:
            from financiero_viviabilidad_tablas import ColombianMunicipalContext
            indicator_structure = ColombianMunicipalContext.INDICATOR_STRUCTURE
        except ImportError:
            indicator_structure = {
                'resultado': ['línea_base', 'meta', 'año_base', 'año_meta', 'fuente',
'responsable'],
                'producto': ['indicador', 'fórmula', 'unidad_medida', 'línea_base',
'meta', 'periodicidad'],
                'gestión': ['eficacia', 'eficiencia', 'economía', 'costo_beneficio']
            }

        for node in nodes.values():
            text_lower = node.text.lower()

            # Cross-reference with INDICATOR_STRUCTURE to classify critical requirements
            # as Hoop Tests or Smoking Guns
            indicator_structure.get(node.type, [])

            # Check if node has all critical DNP requirements (D3-Q1 indicators)
            has_linea_base = bool(
                node.baseline and str(node.baseline).upper() not in ['ND', 'POR DEFINIR',
'N/A', 'NONE'])
            has_meta = bool(node.target and str(node.target).upper() not in ['ND', 'POR
DEFINIR', 'N/A', 'NONE'])
            has_fuente = 'fuente' in text_lower or 'fuente de información' in text_lower

            # Perfect Hoop Test: Missing any critical requirement = total hypothesis
failure
            # This applies to producto nodes with D3-Q1 indicators
            if node.type == 'producto':
                if has_linea_base and has_meta and has_fuente:
                    # Perfect indicators trigger Hoop Test classification
                    node.test_type = 'hoop_test'
                    self.logger.debug(f"Meta {node.id} classified as hoop_test (perfect
D3-Q1 compliance)")

```

```

        elif not has_linea_base or not has_meta:
            # Missing critical requirements - still Hoop Test but will fail
            node.test_type = 'hoop_test'
            node.audit_flags.append('hoop_test_failure')
            self.logger.warning(f"Meta {node.id} FAILS hoop_test (missing D3-Q1
critical fields)")
        else:
            node.test_type = 'straw_in_wind'
    # Check for administrative/regulatory nature (Hoop Test)
    elif any(keyword in text_lower for keyword in self.admin_keywords):
        node.test_type = 'hoop_test'
    # Check for highly specific outcomes (Smoking Gun)
    elif node.type == 'resultado' and node.target and node.baseline:
        try:
            float(str(node.target).replace(',', '').replace('%', ''))
            # Smoking Gun: rare, highly specific evidence with strong inferential
power
            node.test_type = 'smoking_gun'
        except (ValueError, TypeError):
            node.test_type = 'straw_in_wind'
    # Double decisive for critical impact goals
    elif node.type == 'impacto' and node.rigor_status == 'fuerte':
        node.test_type = 'doubly_decisional'
    else:
        node.test_type = 'straw_in_wind'

    self.logger.debug(f"Meta {node.id} assigned test type: {node.test_type}")

```

```

@calibrated_method("saaaaaa.analysis.derek_beach.CausalInferenceSetup.identify_failure
_points")
def identify_failure_points(self, graph, text: str) -> set[str]:
    """Identify single points of failure in causal chain

```

Harmonic Front 3 - Enhancement 2: Contextual Failure Point Detection  
 Expands risk\_pattern to explicitly include localized contextual factors from  
 rubrics:

- restricciones territoriales
- patrones culturales machistas
- limitación normativa

For D6-Q5 (Enfoque Diferencial/Restricciones): Excelente requires ≥3 distinct  
 contextual factors correctly mapped to nodes, satisfying enfoque\_diferencial  
 and análisis\_contextual criteria.

"""

failure\_points = set()

```

# Find nodes with high out-degree (many dependencies)
for node_id in graph.nodes():
    out_degree = graph.out_degree(node_id)
    node_type = graph.nodes[node_id].get('type')

    if node_type == 'producto' and out_degree >= 3:
        failure_points.add(node_id)
        self.logger.warning(f"Uniquo punto de falla identificado: {node_id} "
                           f"(grado de salida: {out_degree})")

```

```

# HARMONIC FRONT 3 - Enhancement 2: Expand contextual factors
# Add specific rubric factors for D6-Q5 compliance
extended_contextual_factors = list(self.contextual_factors) +
    ['restricciones territoriales',
     'restricción territorial',
     'limitación territorial',
     'patrones culturales machistas',
     'machismo',
     'inequidad de género',
     'violencia de género',
     'limitación normativa',
     'limitación legal',

```

```

'restricción legal',
'barriera institucional',
'restricción presupuestal',
'ausencia de capacidad técnica',
'baja capacidad institucional',
'conflicto armado',
'desplazamiento forzado',
'población dispersa',
'ruralidad dispersa',
'acceso vial limitado',
'conectividad deficiente'
]

# Extract contextual risks from text
risk_pattern = '|'.join(re.escape(factor) for factor in
extended_contextual_factors)
risk_regex = re.compile(rf'\b({risk_pattern})\b', re.IGNORECASE)

# Track distinct contextual factors for D6-Q5 quality criteria
contextual_factors_detected = set()
node_contextual_map = defaultdict(set)

# Find risk mentions and associate with nodes
for match in risk_regex.finditer(text):
    risk_text = match.group()
    contextual_factors_detected.add(risk_text.lower())

    context_start = max(0, match.start() - 200)
    context_end = min(len(text), match.end() + 200)
    context = text[context_start:context_end]

    # Try to find node mentions in risk context
    for node_id in graph.nodes():
        if node_id in context:
            failure_points.add(node_id)
            if 'contextual_risks' not in graph.nodes[node_id]:
                graph.nodes[node_id]['contextual_risks'] = []
            graph.nodes[node_id]['contextual_risks'].append(risk_text)
            node_contextual_map[node_id].add(risk_text.lower())

# D6-Q5 quality criteria assessment
distinct_factors_count = len(contextual_factors_detected)
d6_q5_quality = 'insuficiente'
if distinct_factors_count >= 3:
    d6_q5_quality = 'excelente'
elif distinct_factors_count >= 2:
    d6_q5_quality = 'bueno'
elif distinct_factors_count >= 1:
    d6_q5_quality = 'aceptable'

# Store D6-Q5 metrics in graph attributes
graph.graph['d6_q5_contextual_factors'] = list(contextual_factors_detected)
graph.graph['d6_q5_distinct_count'] = distinct_factors_count
graph.graph['d6_q5_quality'] = d6_q5_quality
graph.graph['d6_q5_node_mapping'] = dict(node_contextual_map)

self.logger.info(f"Puntos de falla identificados: {len(failure_points)}")
self.logger.info(
    f"D6-Q5: {distinct_factors_count} factores contextuales distintos detectados - "
    f"{d6_q5_quality}")

return failure_points

@calibrated_method("saaaaaa.analysis.derek_beach.CausalInferenceSetup._get_dynamics_pa
ttern")
def _get_dynamics_pattern(self, dynamics_type: str) -> str:
    """
    Get the pattern associated with a dynamics type.

```

```

Args:
    dynamics_type: Type of dynamics (suma, decreciente, constante, indefinido)

Returns:
    Pattern string for the dynamics type
"""

patterns = {
    'suma': 'suma|total|agregado|consolidado',
    'decreciente': 'reducir|disminuir|decrementar|bajar',
    'constante': 'mantener|sostener|preservar|conservar',
    'indefinido': 'por definir|sin especificar|indefinido'
}
return patterns.get(dynamics_type, "")

class ReportingEngine:
    """Generate visualizations and reports"""

    def __init__(self, config: ConfigLoader, output_dir: Path) -> None:
        self.logger = logging.getLogger(self.__class__.__name__)
        self.config = config
        self.output_dir = output_dir
        self.output_dir.mkdir(parents=True, exist_ok=True)

    @calibrated_method("saaaaaaa.analysis.derek_beach.ReportingEngine.generate_causal_diagram")
    def generate_causal_diagram(self, graph: nx.DiGraph, policy_code: str) -> Path:
        """Generate causal diagram visualization"""
        dot = Dot(graph_type='digraph', rankdir='TB')
        dot.set_name(f'{policy_code}_causal_model')
        dot.set_node_defaults(
            shape='box',
            style='rounded,filled',
            fontname='Arial',
            fontsize='10'
        )
        dot.set_edge_defaults(
            fontsize='8',
            fontname='Arial'
        )

        # Add nodes with rigor coloring
        for node_id in graph.nodes():
            node_data = graph.nodes[node_id]

            # Determine color based on rigor status and audit flags
            rigor = node_data.get('rigor_status', 'sin_evaluar')
            audit_flags = node_data.get('audit_flags', [])
            financial = node_data.get('financial_allocation')

            if rigor == 'débil' or not financial:
                color = 'lightcoral' # Red
            elif audit_flags:
                color = 'lightyellow' # Yellow
            else:
                color = 'lightgreen' # Green

            # Create label
            node_type = node_data.get('type', 'programa')
            text = node_data.get('text', "")[:80]
            label = f"{node_id}\n{node_type.upper()}\n{text}..."

            entity = node_data.get('responsible_entity')
            if entity:
                label += f"\n👤 {entity[:30]}"

            if financial:
                label += f"\n💰 ${financial:.0f}"
```

```

dot_node = Node(
    node_id,
    label=label,
    fillcolor=color
)
dot.add_node(dot_node)

# Add edges with causal logic
for source, target in graph.edges():
    edge_data = graph.edges[source, target]
    keyword = edge_data.get('keyword', '')
    strength = edge_data.get('strength', get_parameter_loader().get("saaaaaaa.analysis.derek_beach.ReportingEngine.generate_causal_diagram").get("auto_param_L3870_49", 0.5))

    # Determine edge style based on strength
    style = 'solid' if strength > get_parameter_loader().get("saaaaaaa.analysis.derek_beach.ReportingEngine.generate_causal_diagram").get("auto_param_L3873_42", 0.7) else 'dashed'

    dot_edge = Edge(
        source,
        target,
        label=keyword[:20],
        style=style
    )
    dot.add_edge(dot_edge)

# Save files
# Delegate to factory for I/O operation
from .factory import write_text_file

dot_path = self.output_dir / f"{policy_code}_causal_diagram.dot"
png_path = self.output_dir / f"{policy_code}_causal_diagram.png"

try:
    write_text_file(dot.to_string(), dot_path)
    self.logger.info(f"Diagrama DOT guardado en: {dot_path}")

    # Try to render PNG
    try:
        dot.write_png(str(png_path))
        self.logger.info(f"Diagrama PNG renderizado en: {png_path}")
    except Exception as e:
        self.logger.warning(f"No se pudo renderizar PNG (¿Graphviz instalado?): {e}")
except Exception as e:
    self.logger.error(f"Error guardando diagrama: {e}")

return png_path

@calibrated_method("saaaaaaa.analysis.derek_beach.ReportingEngine.generate_accountability_matrix")
def generate_accountability_matrix(self, graph: nx.DiGraph,
                                    policy_code: str) -> Path:
    """Generate accountability matrix in Markdown"""
    md_path = self.output_dir / f"{policy_code}_accountability_matrix.md"

    # Group by impact goals
    impact_goals = [n for n in graph.nodes()
                    if graph.nodes[n].get('type') == 'impacto']

    content = [f"# Matriz de Responsabilidades - {policy_code}\n"]
    content.append("**Generado automáticamente por CDAF v2.0**\n")
    content.append("...\n\n")

    for impact in impact_goals:
        impact_data = graph.nodes[impact]

```

```

content.append(f"## Meta de Impacto: {impact}\n")
content.append(f"***Descripción:** {impact_data.get('text', 'N/A')}\n\n")

# Find all predecessor chains
predecessors = list(nx.ancestors(graph, impact))

if predecessors:
    content.append("| Meta | Tipo | Entidad Responsable | Actividad Clave | Presupuesto |\n")
    content.append("-----|-----|-----|-----|\n")

    for pred in predecessors:
        pred_data = graph.nodes[pred]
        meta_type = pred_data.get('type', 'N/A')
        entity = pred_data.get('responsible_entity', 'No asignado')

        ea = pred_data.get('entity_activity')
        activity = 'N/A'
        if ea and isinstance(ea, dict):
            activity = ea.get('activity', 'N/A')

        budget = pred_data.get('financial_allocation')
        budget_str = f"${budget:,0f}" if budget else "Sin presupuesto"

        content.append(f"| {pred} | {meta_type} | {entity} | {activity} | {budget_str} |\n")

    content.append("\n")
else:
    content.append("*No se encontraron metas intermedias.*\n\n")

content.append("\n--\n")
content.append("### Leyenda\n")
content.append("- **Meta de Impacto:** Resultado final esperado\n")
content.append("- **Meta de Resultado:** Cambio intermedio observable\n")
content.append("- **Meta de Producto:** Entrega tangible del programa\n")

# Delegate to factory for I/O operation
from .factory import write_text_file

try:
    write_text_file(".join(content), md_path)
    self.logger.info(f"Matriz de responsabilidades guardada en: {md_path}")
except Exception as e:
    self.logger.error(f"Error guardando matriz de responsabilidades: {e}")

return md_path

@calibrated_method("saaaaaa.analysis.derek_beach.ReportingEngine.generate_confidence_report")
def generate_confidence_report(self,
                               nodes: dict[str, MetaNode],
                               graph: nx.DiGraph,
                               causal_chains: list[CausalLink],
                               audit_results: dict[str, AuditResult],
                               financial_auditor: FinancialAuditor,
                               sequence_warnings: list[str],
                               policy_code: str) -> Path:
    """Generate extraction confidence report"""
    json_path = self.output_dir / f"{policy_code}{EXTRACTION_REPORT_SUFFIX}"

    # Calculate metrics
    total_metas = len(nodes)

    metas_with_ea = sum(1 for n in nodes.values() if n.entity_activity)
    metas_with_ea_pct = (metas_with_ea / total_metas * 100) if total_metas > 0 else 0

```

```

enlaces_with_logic = sum(1 for link in causal_chains if link.get('logic'))
total_edges = graph.number_of_edges()
enlaces_with_logic_pct = (enlaces_with_logic / total_edges * 100) if total_edges >
0 else 0

metas_passed_audit = sum(1 for r in audit_results.values() if r['passed'])
metas_with_traceability_pct = (metas_passed_audit / total_metas * 100) if
total_metas > 0 else 0

metas_with_financial = sum(1 for n in nodes.values() if n.financial_allocation)
metas_with_financial_pct = (metas_with_financial / total_metas * 100) if
total_metas > 0 else 0

# Node type distribution
type_distribution = defaultdict(int)
for node in nodes.values():
    type_distribution[node.type] += 1

# Rigor distribution
rigor_distribution = defaultdict(int)
for node in nodes.values():
    rigor_distribution[node.rigor_status] += 1

report = {
    "metadata": {
        "policy_code": policy_code,
        "framework_version": "2.get_parameter_loader().get("saaaaaaaa.analysis.derek
_beach.ReportingEngine.generate_confidence_report").get("auto_param_L4008_40", 0.0)",
        "total_nodes": total_metas,
        "total_edges": total_edges
    },
    "extraction_metrics": {
        "total_metas_identificadas": total_metas,
        "metas_con_EA_extraido": metas_with_ea,
        "metas_con_EA_extraido_pct": round(metas_with_ea_pct, 2),
        "enlaces_con_logica_causal": enlaces_with_logic,
        "enlaces_con_logica_causal_pct": round(enlaces_with_logic_pct, 2),
        "metas_con_trazabilidad_evidencia": metas_passed_audit,
        "metas_con_trazabilidad_evidencia_pct": round(metas_with_traceability_pct,
2),
        "metas_con_trazabilidad_financiera": metas_with_financial,
        "metas_con_trazabilidad_financiera_pct": round(metas_with_financial_pct,
2)
    },
    "financial_audit": {
        "tablas_financieras_parseadas_exitosamente": financial_auditor.successful_parses,
        "tablas_financieras_fallidas": financial_auditor.failed_parses,
        "asignaciones_presupuestarias_rastreadas": len(financial_auditor.financial_data)
    },
    "sequence_audit": {
        "alertas_secuencia_logica": len(sequence_warnings),
        "detalles": sequence_warnings
    },
    "type_distribution": dict(type_distribution),
    "rigor_distribution": dict(rigor_distribution),
    "audit_summary": {
        "total audited": len(audit_results),
        "passed": sum(1 for r in audit_results.values() if r['passed']),
        "failed": sum(1 for r in audit_results.values() if not r['passed']),
        "total_warnings": sum(len(r['warnings']) for r in audit_results.values()),
        "total_errors": sum(len(r['errors']) for r in audit_results.values())
    },
    "quality_score": self._calculate_quality_score(
        metas_with_traceability_pct,
        metas_with_financial_pct,
        enlaces_with_logic_pct,

```

```

        metas_with_ea_pct
    )
}

# Delegate to factory for I/O operation
from .factory import save_json

try:
    save_json(report, json_path)
    self.logger.info(f"Reporte de confianza guardado en: {json_path}")
except Exception as e:
    self.logger.error(f"Error guardando reporte de confianza: {e}")

return json_path

@calibrated_method("saaaaaaa.analysis.derek_beach.ReportingEngine._calculate_quality_score")
def _calculate_quality_score(self, traceability: float, financial: float,
                             logic: float, ea: float) -> float:
    """Calculate overall quality score (0-100)"""
    weights = {'traceability': get_parameter_loader().get("saaaaaaa.analysis.derek_beach.ReportingEngine._calculate_quality_score").get("auto_param_L4064_35", 0.35),
               'financial': get_parameter_loader().get("saaaaaaa.analysis.derek_beach.ReportingEngine._calculate_quality_score").get("auto_param_L4064_54", 0.25), 'logic': get_parameter_loader().get("saaaaaaa.analysis.derek_beach.ReportingEngine._calculate_quality_score").get("auto_param_L4064_69", 0.25), 'ea': get_parameter_loader().get("saaaaaaa.analysis.derek_beach.ReportingEngine._calculate_quality_score").get("auto_param_L4064_81", 0.15)}
    score = (traceability * weights['traceability'] +
             financial * weights['financial'] +
             logic * weights['logic'] +
             ea * weights['ea'])
    return round(score, 2)

@calibrated_method("saaaaaaa.analysis.derek_beach.ReportingEngine.generate_causal_model_json")
def generate_causal_model_json(self, graph: nx.DiGraph, nodes: dict[str, MetaNode],
                               policy_code: str) -> Path:
    """Generate structured JSON export of causal model"""
    json_path = self.output_dir / f"{policy_code}{CAUSAL_MODEL_SUFFIX}"

    # Prepare node data
    nodes_data = {}
    for node_id, node in nodes.items():
        node_dict = asdict(node)
        # Convert NamedTuple to dict
        if node.entity_activity:
            node_dict['entity_activity'] = node.entity_activity._asdict()
        nodes_data[node_id] = node_dict

    # Prepare edge data
    edges_data = []
    for source, target in graph.edges():
        edge_dict = {
            'source': source,
            'target': target,
            **graph.edges[source, target]
        }
        edges_data.append(edge_dict)

    model_data = {
        "policy_code": policy_code,
        "framework_version": "2.get_parameter_loader().get("saaaaaaa.analysis.derek_beach.ReportingEngine.generate_causal_model_json").get("auto_param_L4098_36", 0.0)",
        "nodes": nodes_data,
        "edges": edges_data,
        "statistics": {
            "total_nodes": len(nodes_data),
            "total_edges": len(edges_data),
    
```

```

        "node_types": {
            node_type: sum(1 for n in nodes.values() if n.type == node_type)
            for node_type in ['programa', 'producto', 'resultado', 'impacto']
        }
    }

# Delegate to factory for I/O operation
from .factory import save_json

try:
    save_json(model_data, json_path)
    self.logger.info(f"Modelo causal JSON guardado en: {json_path}")
except Exception as e:
    self.logger.error(f"Error guardando modelo causal: {e}")

return json_path

class CDAFFramework:
    """Main orchestrator for the CDAF pipeline"""

    def __init__(self, config_path: Path, output_dir: Path, log_level: str = "INFO") ->
None:
        self.logger = logging.getLogger(self.__class__.__name__)
        self.logger.setLevel(getattr(logging, log_level.upper()))

        # Initialize components
        self.config = ConfigLoader(config_path)
        self.output_dir = output_dir

        # Initialize retry handler for external dependencies
        try:
            from retry_handler import DependencyType, get_retry_handler
            self.retry_handler = get_retry_handler()
            retry_enabled = True
        except ImportError:
            self.logger.warning("RetryHandler no disponible, funcionando sin retry logic")
            self.retry_handler = None
            retry_enabled = False

        # Load spaCy model with retry logic
        # Delegate to factory for I/O operation
        from .factory import load_spacy_model

        if retry_enabled and self.retry_handler:
            @self.retry_handler.with_retry(
                DependencyType.SPACY_MODEL,
                operation_name="load_spacy_model",
                exceptions=(OSError, IOError, ImportError)
            )
        def load_spacy_with_retry():
            try:
                nlp = load_spacy_model("es_core_news_lg")
                self.logger.info("Modelo spaCy cargado: es_core_news_lg")
                return nlp
            except OSError:
                self.logger.warning("Modelo es_core_news_lg no encontrado. Intentando es_core_news_sm...")
                nlp = load_spacy_model("es_core_news_sm")
                return nlp

            try:
                self.nlp = load_spacy_with_retry()
            except OSError:
                self.logger.error("No se encontró ningún modelo de spaCy en español. "
                                 "Ejecute: python -m spacy download es_core_news_lg")
                sys.exit(1)
        else:

```

```

# Fallback to original logic without retry
try:
    self.nlp = load_spacy_model("es_core_news_lg")
    self.logger.info("Modelo spaCy cargado: es_core_news_lg")
except OSError:
    self.logger.warning("Modelo es_core_news_lg no encontrado. Intentando
es_core_news_sm...")
    try:
        self.nlp = load_spacy_model("es_core_news_sm")
    except OSError:
        self.logger.error("No se encontró ningún modelo de spaCy en español. "
                          "Ejecute: python -m spacy download es_core_news_lg")
        sys.exit(1)

# Initialize modules (pass retry_handler to PDF processor)
self.pdf_processor = PDFProcessor(self.config, retry_handler=self.retry_handler if
retry_enabled else None)
self.causal_extractor = CausalExtractor(self.config, self.nlp)
self.mechanism_extractor = MechanismPartExtractor(self.config, self.nlp)
self.bayesian_mechanism = BayesianMechanismInference(self.config, self.nlp)
self.financial_auditor = FinancialAuditor(self.config)
self.op_auditor = OperationalizationAuditor(self.config)
self.inference_setup = CausalInferenceSetup(self.config)
self.reporting_engine = ReportingEngine(self.config, output_dir)

# Initialize DNP validator if available
self.dnp_validator = None
if DNP_AVAILABLE:
    self.dnp_validator = ValidadorDNP(es_municipio_pdet=False) # Can be
configured
    self.logger.info("Validador DNP inicializado")

@calibrated_method("saaaaaa.analysis.derek_beach.CDAFFramework.process_document")
def process_document(self, pdf_path: Path, policy_code: str) -> bool:
    """Main processing pipeline"""
    self.logger.info(f"Iniciando procesamiento de documento: {pdf_path}")

    try:
        # Step 1: Load and extract PDF
        if not self.pdf_processor.load_document(pdf_path):
            return False

        text = self.pdf_processor.extract_text()
        tables = self.pdf_processor.extract_tables()
        self.pdf_processor.extract_sections()

        # Step 2: Extract causal hierarchy
        self.logger.info("Extrayendo jerarquía causal...")
        graph = self.causal_extractor.extract_causal_hierarchy(text)
        nodes = self.causal_extractor.nodes

        # Step 3: Extract Entity-Activity pairs
        self.logger.info("Extrayendo tuplas Entidad-Actividad...")
        for node in nodes.values():
            if node.type == 'producto':
                ea = self.mechanism_extractor.extract_entity_activity(node.text)
                if ea:
                    node.entity_activity = ea
                    graph.nodes[node.id]['entity_activity'] = ea._asdict()

        # Step 4: Financial traceability
        self.logger.info("Auditando trazabilidad financiera...")
        self.financial_auditor.trace_financial_allocation(tables, nodes, graph)

        # Step 4.5: Bayesian Mechanism Inference (AGUJA II)
        self.logger.info("Infiriendo mecanismos causales con modelo Bayesiano...")
        inferred_mechanisms = self.bayesian_mechanism.infer_mechanisms(nodes, text)

    except Exception as e:
        self.logger.error(f"Error during processing: {e}")
        return False

    return True

```

```

# Step 5: Operationalization audit
self.logger.info("Auditando operacionalización...")
audit_results = self.op_auditor.audit_evidence_traceability(nodes)
sequence_warnings = self.op_auditor.audit_sequence_logic(graph)

# Step 5.5: Bayesian Counterfactual Audit (AGUJA III)
# Note: pdet_alignment should be calculated separately if needed via
financiero_viability_tablas
# For now, using None as placeholder - can be enhanced by integrating
PDETMunicipalPlanAnalyzer
self.logger.info("Ejecutando auditoría contrafactual Bayesiana...")
counterfactual_audit = self.op_auditor.bayesian_counterfactual_audit(nodes,
graph, pdet_alignment=None)

# Step 6: Causal inference setup
self.logger.info("Preparando para inferencia causal...")
self.inference_setup.classify_goal_dynamics(nodes)
self.inference_setup.assign_probative_value(nodes)
self.inference_setup.identify_failure_points(graph, text)

# Step 7: DNP Standards Validation (if available)
if self.dnp_validator:
    self.logger.info("Validando cumplimiento de estándares DNP...")
    self._validate_dnp_compliance(nodes, graph, policy_code)

# Step 8: Generate reports
self.logger.info("Generando reportes y visualizaciones...")
self.reporting_engine.generate_causal_diagram(graph, policy_code)
self.reporting_engine.generate_accountability_matrix(graph, policy_code)
self.reporting_engine.generate_confidence_report(
    nodes, graph, self.causal_extractor.causal_chains,
    audit_results, self.financial_auditor, sequence_warnings, policy_code
)
self.reporting_engine.generate_causal_model_json(graph, nodes, policy_code)

# Step 8: Generate Bayesian inference reports
self.logger.info("Generando reportes de inferencia Bayesiana...")
self._generate_bayesian_reports(
    inferred_mechanisms, counterfactual_audit, policy_code
)

# Step 9: Self-reflective learning from audit results (frontier paradigm)
if self.config.validated_config and
self.config.validated_config.self_reflection.enable_prior_learning:
    self.logger.info("Actualizando priors con retroalimentación del
análisis...")
    feedback_data = self._extract_feedback_from_audit(
        inferred_mechanisms, counterfactual_audit, audit_results
    )
    self.config.update_priors_from_feedback(feedback_data)

# HARMONIC FRONT 4: Check uncertainty reduction criterion
if hasattr(self.bayesian_mechanism, '_mean_mechanism_uncertainty'):
    uncertainty_check = self.config.check_uncertainty_reduction_criterion(
        self.bayesian_mechanism._mean_mechanism_uncertainty
    )
    self.logger.info(
        f"Uncertainty criterion check: {uncertainty_check['status']} "
        f"({uncertainty_check['iterations_tracked']}/10 iterations, "
        f"{uncertainty_check['reduction_percent']:.2f}% reduction)"
    )

    self.logger.info(f"✓ Procesamiento completado exitosamente para
{policy_code}")
    return True

except CDAFException as e:
    # Structured error handling with custom exceptions

```

```

self.logger.error(f"Error CDAF: {e.message}")
self.logger.error(f"Detalles: {json.dumps(e.to_dict(), indent=2)}")
if not erecoverable:
    raise
return False
except Exception as e:
    # Wrap unexpected errors in CDAFProcessingError
    raise CDAFProcessingError(
        "Error crítico en el procesamiento",
        details={'error': str(e), 'type': type(e).__name__},
        stage="document_processing",
        recoverable=False
    ) from e

@calibrated_method("saaaaaa.analysis.derek_beach.CDAFFramework._extract_feedback_from_
audit")
def _extract_feedback_from_audit(self, inferred_mechanisms: dict[str, dict[str, Any]],
                                 counterfactual_audit: dict[str, Any],
                                 audit_results: dict[str, AuditResult]) -> dict[str,
Any]:
    """
    Extract feedback data from audit results for self-reflective prior updating

    This implements the frontier paradigm of learning from audit results
    to improve future inference accuracy.

    HARMONIC FRONT 4 ENHANCEMENT:
    - Reduces mechanism_type_priors for mechanisms with implementation_failure flags
    - Tracks necessity/sufficiency test failures
    - Penalizes "miracle" mechanisms that fail counterfactual tests
    """

    feedback = {}

    # Extract mechanism type frequencies from successful inferences
    mechanism_frequencies = defaultdict(float)
    failure_frequencies = defaultdict(float) # NEW: Track failures
    total_mechanisms = 0
    total_failures = 0

    # Get causal implications from audit
    causal_impressions = counterfactual_audit.get('causal_impressions', {})

    for node_id, mechanism in inferred_mechanisms.items():
        mechanism_type_dist = mechanism.get('mechanism_type', {})
        # Weight by confidence (coherence score)
        confidence = mechanism.get('coherence_score', get_parameter_loader().get("sa
aaa.analysis.derek_beach.CDAFFramework._extract_feedback_from_audit").get("auto_param_L434
0_58", 0.5))

        # Check for implementation_failure flags in audit results
        node_impressions = causal_impressions.get(node_id, {})
        causal_effects = node_impressions.get('causal_effects', {})
        hasImplementationFailure = 'implementation_failure' in causal_effects

        # Check necessity/sufficiency test results
        necessity_test = mechanism.get('necessity_test', {})
        sufficiency_test = mechanism.get('sufficiency_test', {})
        failed_necessity = not necessity_test.get('is_necessary', True)
        failed_sufficiency = not sufficiency_test.get('is_sufficient', True)

        # If mechanism failed tests or has implementation_failure flag
        if hasImplementationFailure or failed_necessity or failed_sufficiency:
            total_failures += 1
            # Track which mechanism types are associated with failures
            for mech_type, prob in mechanism_type_dist.items():
                failure_frequencies[mech_type] += prob * confidence
        else:
            # Only count successes for positive reinforcement

```

```

        for mech_type, prob in mechanism_type_dist.items():
            mechanism_frequencies[mech_type] += prob * confidence
            total_mechanisms += confidence

    # Normalize frequencies
    if total_mechanisms > 0:
        mechanism_frequencies = {
            k: v / total_mechanisms
            for k, v in mechanism_frequencies.items()
        }
        feedback['mechanism_frequencies'] = dict(mechanism_frequencies)

    # NEW: Calculate penalty factors for failed mechanism types
    if total_failures > 0:
        failure_frequencies = {
            k: v / total_failures
            for k, v in failure_frequencies.items()
        }
        feedback['failure_frequencies'] = dict(failure_frequencies)

    # Calculate penalty: reduce priors for frequently failing types
    penalty_factors = {}
    for mech_type, failure_freq in failure_frequencies.items():
        # Higher failure frequency = stronger penalty (get_parameter_loader().get(
        "aaaaaaaa.analysis.derek_beach.CDAFFramework._extract_feedback_from_audit").get("auto_param_L4384_63", 0.7) to get_parameter_loader().get("aaaaaaaa.analysis.derek_beach.CDAFFramework._extract_feedback_from_audit").get("auto_param_L4384_70", 0.95) reduction)
        penalty_factors[mech_type] = get_parameter_loader().get("aaaaaaaa.analysis.derek_beach.CDAFFramework._extract_feedback_from_audit").get("auto_param_L4385_45", 0.95) - (failure_freq * get_parameter_loader().get("aaaaaaaa.analysis.derek_beach.CDAFFramework._extract_feedback_from_audit").get("auto_param_L4385_68", 0.25))
    feedback['penalty_factors'] = penalty_factors

    # Add audit quality metrics for future reference
    feedback['audit_quality'] = {
        'total_nodes audited': len(audit_results),
        'passed_count': sum(1 for r in audit_results.values() if r['passed']),
        'success_rate': sum(1 for r in audit_results.values() if r['passed']) / max(len(audit_results), 1),
        'failure_count': total_failures, # NEW
        'failure_rate': total_failures / max(len(inferred_mechanisms), 1) # NEW
    }

    # Track necessity/sufficiency failures for iterative validation loop
    necessity_failures = sum(1 for m in inferred_mechanisms.values()
        if not m.get('necessity_test', {}).get('is_necessary',
        True))
    sufficiency_failures = sum(1 for m in inferred_mechanisms.values()
        if not m.get('sufficiency_test',
        {}).get('is_sufficient', True))

    feedback['test_failures'] = {
        'necessity_failures': necessity_failures,
        'sufficiency_failures': sufficiency_failures
    }

    return feedback

@calibrated_method("aaaaaaaa.analysis.derek_beach.CDAFFramework._validate_dnp_compliance")
def _validate_dnp_compliance(self, nodes: dict[str, MetaNode],
    graph: nx.DiGraph, policy_code: str) -> None:
    """
    Validate DNP compliance for all nodes/projects
    Generates DNP compliance report
    """
    if not self.dnp_validator:
        return

```

```

# Build project list from nodes
proyectos = []
for node_id, node in nodes.items():
    # Extract sector from responsible entity or type
    sector = "general"
    if node.responsible_entity:
        entity_lower = node.responsible_entity.lower()
        if "educaci" in entity_lower or "edu" in entity_lower:
            sector = "educacion"
        elif "salud" in entity_lower:
            sector = "salud"
        elif "agua" in entity_lower or "acueducto" in entity_lower:
            sector = "agua_potable_saneamiento"
        elif (
            "via" in entity_lower or "vial" in entity_lower or "transporte" in
entity_lower or "infraestructura" in entity_lower):
            sector = "vias_transporte"
        elif "agr" in entity_lower or "rural" in entity_lower:
            sector = "desarrollo_agropecuario"

    # Infer indicators from node type
    indicadores = []
    if node.type == "producto":
        # Map to MGA product indicators based on sector
        if sector == "educacion":
            indicadores = ["EDU-020", "EDU-021"]
        elif sector == "salud":
            indicadores = ["SAL-020", "SAL-021"]
        elif sector == "agua_potable_saneamiento":
            indicadores = ["APS-020", "APS-021"]
    elif node.type == "resultado":
        # Map to MGA result indicators
        if sector == "educacion":
            indicadores = ["EDU-001", "EDU-002"]
        elif sector == "salud":
            indicadores = ["SAL-001", "SAL-002"]
        elif sector == "agua_potable_saneamiento":
            indicadores = ["APS-001", "APS-002"]

    proyectos.append({
        "nombre": node_id,
        "sector": sector,
        "descripcion": node.text[:200] if node.text else "",
        "indicadores": indicadores,
        "presupuesto": node.financial_allocation or get_parameter_loader().get("sa
aaaaa.analysis.derek_beach.CDAFFramework._validate_dnp_compliance").get("auto_param_L4463_
60", 0.0),
        "es_rural": "rural" in node.text.lower() if node.text else False,
        "poblacion_victimas": "v ctima" in node.text.lower() if node.text else
False
    })
}

# Validate each project
dnp_results = []
for proyecto in proyectos:
    resultado = self.dnp_validator.validar_proyecto_integral(
        sector=proyecto["sector"],
        descripcion=proyecto["descripcion"],
        indicadores_propuestos=proyecto["indicadores"],
        presupuesto=proyecto["presupuesto"],
        es_rural=proyecto["es_rural"],
        poblacion_victimas=proyecto["poblacion_victimas"]
    )
    dnp_results.append({
        "proyecto": proyecto["nombre"],
        "resultado": resultado
    })
}

```

```

# Generate DNP compliance report
self._generate_dnp_report(dnp_results, policy_code)

@calibrated_method("saaaaaa.analysis.derek_beach.CDAFFramework._generate_dnp_report")
def _generate_dnp_report(self, dnp_results: list[dict], policy_code: str) -> None:
    """Generate comprehensive DNP compliance report"""
    report_path = self.output_dir / f"{policy_code}{DNP_REPORT_SUFFIX}"

    total_proyectos = len(dnp_results)
    if total_proyectos == 0:
        return

    # Calculate aggregate statistics
    proyectos_excelente = sum(1 for r in dnp_results
                               if r["resultado"].nivel_cumplimiento.value ==
                               "excelente")
    proyectos_bueno = sum(1 for r in dnp_results
                          if r["resultado"].nivel_cumplimiento.value == "bueno")
    proyectos_aceptable = sum(1 for r in dnp_results
                              if r["resultado"].nivel_cumplimiento.value ==
                               "aceptable")
    proyectos_insuficiente = sum(1 for r in dnp_results
                                  if r["resultado"].nivel_cumplimiento.value ==
                                   "insuficiente")

    score_promedio = sum(r["resultado"].score_total for r in dnp_results) /
    total_proyectos

    # Build report
    lines = []
    lines.append("=" * 100)
    lines.append("REPORTE DE CUMPLIMIENTO DE ESTÁNDARES DNP")
    lines.append(f"Código de Política: {policy_code}")
    lines.append("=" * 100)
    lines.append("")

    lines.append("RESUMEN EJECUTIVO")
    lines.append("-" * 100)
    lines.append(f"Total de Proyectos/Metas Analizados: {total_proyectos}")
    lines.append(f"Score Promedio de Cumplimiento: {score_promedio:.1f}/100")
    lines.append("")
    lines.append("Distribución por Nivel de Cumplimiento:")
    lines.append(
        f" • Excelente (>90%): {proyectos_excelente:3d} ({proyectos_excelente / total_proyectos * 100:.1f}%)")
    lines.append(
        f" • Bueno (75-90%): {proyectos_bueno:3d} ({proyectos_bueno / total_proyectos * 100:.1f}%)")
    lines.append(
        f" • Aceptable (60-75%): {proyectos_aceptable:3d} ({proyectos_aceptable / total_proyectos * 100:.1f}%)")
    lines.append(
        f" • Insuficiente (<60%): {proyectos_insuficiente:3d} ({proyectos_insuficiente / total_proyectos * 100:.1f}%)")
    lines.append("")

    # Detailed validation per project
    lines.append("VALIDACIÓN DETALLADA POR PROYECTO/META")
    lines.append("=" * 100)

    for i, result_data in enumerate(dnp_results, 1):
        proyecto = result_data["proyecto"]
        resultado = result_data["resultado"]

        lines.append("")
        lines.append(f"{i}. {proyecto}")
        lines.append("-" * 100)

```

```

lines.append(
    f" Score: {resultado.score_total:.1f}/100 | Nivel:
{resultado.nivel_cumplimiento.value.upper()}")


# Competencies
comp_status = "✓" if resultado.cumple_competencias else "✗"
lines.append(f" Competencias Municipales: {comp_status}")
if resultado.competencias_validadas:
    lines.append(f" - Aplicables: {',
'.join(resultado.competencias_validadas[:3])}")


# MGA Indicators
mga_status = "✓" if resultado.cumple_mga else "✗"
lines.append(f" Indicadores MGA: {mga_status}")
if resultado.indicadores_mga_usados:
    lines.append(f" - Usados: {',
'.join(resultado.indicadores_mga_usados)}")
if resultado.indicadores_mga_faltantes:
    lines.append(f" - Recomendados: {',
'.join(resultado.indicadores_mga_faltantes)}")


# PDET (if applicable)
if resultado.es_municipio_pdet:
    pdet_status = "✓" if resultado.cumple_pdet else "✗"
    lines.append(f" Lineamientos PDET: {pdet_status}")
    if resultado.lineamientos_pdet_cumplidos:
        lines.append(f" - Cumplidos:
{len(resultado.lineamientos_pdet_cumplidos)}")


# Critical alerts
if resultado.alertas_criticas:
    lines.append(" △ ALERTAS CRÍTICAS:")
    for alerta in resultado.alertas_criticas:
        lines.append(f" - {alerta}")


# Recommendations
if resultado.recomendaciones:
    lines.append(" ↲ RECOMENDACIONES:")
    for rec in resultado.recomendaciones[:3]: # Top 3
        lines.append(f" - {rec}")


lines.append("")
lines.append("=" * 100)
lines.append("NORMATIVA DE REFERENCIA")
lines.append("-" * 100)
lines.append("• Competencias Municipales: Ley 136/1994, Ley 715/2001, Ley
1551/2012")
lines.append("• Indicadores MGA: DNP - Metodología General Ajustada")
lines.append("• PDET: Decreto 893/2017, Acuerdo Final de Paz")
lines.append("=" * 100)


# Write report
# Delegate to factory for I/O operation
from .factory import write_text_file


try:
    write_text_file('\n'.join(lines), report_path)
    self.logger.info(f"Reporte de cumplimiento DNP guardado en: {report_path}")
except Exception as e:
    self.logger.error(f"Error guardando reporte DNP: {e}")


@calibrated_method("saaaaaa.analysis.derek_beach.CDAFFramework._audit_causal_coherence")
def _audit_causal_coherence(self, graph: nx.DiGraph, nodes: dict[str, MetaNode]) ->
dict[str, Any]:
"""
Audit causal coherence of the extracted model.

```

Args:

- graph: Causal graph
- nodes: Dictionary of nodes

Returns:

- Dictionary with coherence audit results

```
"""
audit = {
    'total_nodes': len(nodes),
    'total_edges': graph.number_of_edges(),
    'disconnected_nodes': [],
    'cycles': [],
    'coherence_score': get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CD
AFFramework._audit_causal_coherence").get("auto_param_L4615_31", 0.0)
}
# Check for disconnected nodes
for node_id in nodes:
    if graph.has_node(node_id) and graph.degree(node_id) == 0:
        audit['disconnected_nodes'].append(node_id)

# Check for cycles (should not exist in causal DAG)
try:
    cycles = list(nx.simple_cycles(graph))
    audit['cycles'] = cycles
except:
    pass

# Calculate coherence score
connected_ratio = get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CDAFFra
mework._audit_causal_coherence").get("auto_param_L4631_26", 1.0) -
(len(audit['disconnected_nodes']) / max(len(nodes), 1))
acyclic_score = get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CDAFFrame
work._audit_causal_coherence").get("auto_param_L4632_24", 1.0) if len(audit['cycles']) ==
0 else get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CDAFFramework._audit_causa
l_coherence").get("auto_param_L4632_62", 0.5)
audit['coherence_score'] = (connected_ratio + acyclic_score) / 2.0

return audit

@calibrated_method("saaaaaaa.analysis.derek_beach.CDAFFramework._generate_causal_model_
json")
def _generate_causal_model_json(self, graph: nx.DiGraph, nodes: dict[str, MetaNode],
                                policy_code: str) -> None:
"""
Generate JSON representation of causal model.

Args:
    graph: Causal graph
    nodes: Dictionary of nodes
    policy_code: Policy code for filename
"""
model = {
    'policy_code': policy_code,
    'nodes': [],
    'edges': []
}

# Add nodes
for node_id, node in nodes.items():
    model['nodes'].append({
        'id': node_id,
        'text': node.text,
        'type': node.type,
        'baseline': str(node.baseline) if node.baseline else None,
        'target': str(node.target) if node.target else None
    })
}
```

```

# Add edges
for source, target in graph.edges():
    edge_data = graph.get_edge_data(source, target)
    model['edges'].append({
        'source': source,
        'target': target,
        'logic': edge_data.get('logic', 'unknown'),
        'strength': edge_data.get('strength', get_parameter_loader().get("saaaaaa.
analysis.derek_beach.CDAFFramework._generate_causal_model_json").get("auto_param_L4671_54"
, 0.5))
    })

# Write to file
output_path = self.output_dir / f"{policy_code}{CAUSAL_MODEL_SUFFIX}"
try:
    with open(output_path, 'w', encoding='utf-8') as f:
        json.dump(model, f, indent=2, ensure_ascii=False)
    self.logger.info(f"Causal model JSON saved to: {output_path}")
except Exception as e:
    self.logger.error(f"Error saving causal model JSON: {e}")

@calibrated_method("saaaaaa.analysis.derek_beach.CDAFFramework._generate_dnp_compliance_report")
def _generate_dnp_compliance_report(self, nodes: dict[str, MetaNode],
                                    policy_code: str) -> dict[str, Any]:
    """
    Generate DNP compliance report.

    Args:
        nodes: Dictionary of nodes
        policy_code: Policy code

    Returns:
        Compliance report dictionary
    """
    report = {
        'policy_code': policy_code,
        'total_products': 0,
        'compliant_products': 0,
        'compliance_rate': get_parameter_loader().get("saaaaaa.analysis.derek_beach.CD
AFFFramework._generate_dnp_compliance_report").get("auto_param_L4700_31", 0.0),
        'gaps': []
    }

    # Check products for DNP compliance
    for node_id, node in nodes.items():
        if node.type == 'producto':
            report['total_products'] += 1

            # Check required fields
            has_baseline = node.baseline is not None
            has_target = node.target is not None
            has_indicator = len(node.text) > 10 # Simple check

            is_compliant = has_baseline and has_target and has_indicator

            if is_compliant:
                report['compliant_products'] += 1
            else:
                gaps = []
                if not has_baseline:
                    gaps.append('missing_baseline')
                if not has_target:
                    gaps.append('missing_target')
                if not has_indicator:
                    gaps.append('missing_indicator')

            report['gaps'].append({

```

```

        'node_id': node_id,
        'issues': gaps
    })

if report['total_products'] > 0:
    report['compliance_rate'] = report['compliant_products'] /
report['total_products']

return report

@calibrated_method("saaaaaa.analysis.derek_beach.CDAFFramework._generate_extraction_report")
def _generate_extraction_report(self, nodes: dict[str, MetaNode],
                                graph: nx.DiGraph,
                                policy_code: str) -> None:
    """
    Generate extraction confidence report.

    Args:
        nodes: Dictionary of nodes
        graph: Causal graph
        policy_code: Policy code
    """

    report = {
        'policy_code': policy_code,
        'extraction_summary': {
            'total_nodes': len(nodes),
            'total_edges': graph.number_of_edges(),
            'nodes_by_type': {}
        },
        'node_confidence': []
    }

    # Count nodes by type
    for node in nodes.values():
        node_type = node.type
        report['extraction_summary']['nodes_by_type'][node_type] = \
            report['extraction_summary']['nodes_by_type'].get(node_type, 0) + 1

    # Add confidence scores
    for node_id, node in nodes.items():
        confidence = get_parameter_loader().get("saaaaaa.analysis.derek_beach.CDAFFramework._generate_extraction_report").get("confidence", 0.8) # Refactored
        if hasattr(node, 'rigor_status'):
            if node.rigor_status == 'fuerte':
                confidence = get_parameter_loader().get("saaaaaa.analysis.derek_beach.CDAFFramework._generate_extraction_report").get("confidence", 0.9) # Refactored
            elif node.rigor_status == 'débil':
                confidence = get_parameter_loader().get("saaaaaa.analysis.derek_beach.CDAFFramework._generate_extraction_report").get("confidence", 0.6) # Refactored

        report['node_confidence'].append({
            'node_id': node_id,
            'confidence': confidence
        })

    # Write report
    output_path = self.output_dir / f"{policy_code}{EXTRACTION_REPORT_SUFFIX}"
    try:
        with open(output_path, 'w', encoding='utf-8') as f:
            json.dump(report, f, indent=2, ensure_ascii=False)
        self.logger.info(f"Extraction report saved to: {output_path}")
    except Exception as e:
        self.logger.error(f"Error saving extraction report: {e}")

# =====#
# AGUJA I: PRIOR ADAPTATIVO (EVIDENCIA → BAYES)
# =====#

```

```

class BayesFactorTable:
    """Tabla fija de Bayes Factors por tipo de test evidencial (Beach & Pedersen 2019)"""
    FACTORS = {
        'straw': (get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CDAFFramework._generate_extraction_report").get("auto_param_L4795_18", 1.0), 1.5),      # STRAW_IN_WIND:
        Weak evidence
        'hoop': (3.0, 5.0),      # HOOP TEST: Necessary but not sufficient
        'smoking': (1get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CDAFFramework._generate_extraction_report").get("auto_param_L4797_21", 0.0), 3get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CDAFFramework._generate_extraction_report").get("auto_param_L4797_27", 0.0)), # SMOKING GUN: Sufficient but not necessary
        'doubly': (5get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CDAFFramework._generate_extraction_report").get("auto_param_L4798_20", 0.0), 10get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CDAFFramework._generate_extraction_report").get("auto_param_L4798_27", 0.0)) # DOUBLY DECISIVE: Necessary AND sufficient
    }

    @classmethod
    def get_bayes_factor(cls, test_type: str) -> float:
        """Obtiene BF medio para tipo de test"""
        if test_type not in cls.FACTORS:
            return 1.5 # Default straw-in-wind
        min_bf, max_bf = cls.FACTORS[test_type]
        return (min_bf + max_bf) / 2.0

    @classmethod
    def get_version(cls) -> str:
        """Version de tabla BF para trazabilidad"""
        return "Beach2019_vget_parameter_loader().get("saaaaaaa.analysis.derek_beach.CDAFFramework._generate_extraction_report").get("auto_param_L4812_27", 1.0)"

class AdaptivePriorCalculator:
    """
    AGUJA I - Prior Adaptativo con Bayes Factor y calibración
    PROMPT I-1: Ponderación evidencial con BF y calibración
    Mapea test_type→BayesFactor, calcula likelihood adaptativo combinando
    dominios {semantic, temporal, financial, structural} con pesos normalizados.
    PROMPT I-2: Sensibilidad, OOD y ablation evidencial
    Perturba cada componente ±10% y reporta ∂p/∂component top-3.
    PROMPT I-3: Trazabilidad y reproducibilidad
    Con semilla fija, guarda bf_table_version, weights_version, snippets.
    QUALITY CRITERIA:
    - BrierScore ≤ get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CDAFFramework._generate_extraction_report").get("auto_param_L4829_19", 0.20) en validación sintética
    - ACE ∈ [-get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CDAFFramework._generate_extraction_report").get("auto_param_L4830_14", 0.02), get_parameter_loader().get("saaaaaaa.analysis.derek_beach.CDAFFramework._generate_extraction_report").get("auto_param_L4830_20", 0.02)] (Average Calibration Error)
    - Cobertura CI95% ∈ [92%, 98%]
    - Monotonidad: ↑ señales → ↓ p_mechanism
    """

```

```

def __init__(self, calibration_params: dict[str, float] | None = None) -> None:
    self.logger = logging.getLogger(self.__class__.__name__)
    self.bf_table = BayesFactorTable()

    # Calibration params: logit⁻¹(α + β·score)
    self.calibration = calibration_params or {
        'alpha': -2.0, # Intercept
        'beta': 4.0   # Slope
    }

    # Domain weights (normalized)

```

```

    self.default_domain_weights = {
        'semantic': get_parameter_loader().get("saaaaaaa.analysis.derek_beach.AdaptivePriorCalculator.__init__").get("auto_param_L4847_24", 0.35),
        'temporal': get_parameter_loader().get("saaaaaaa.analysis.derek_beach.AdaptivePriorCalculator.__init__").get("auto_param_L4848_24", 0.25),
        'financial': get_parameter_loader().get("saaaaaaa.analysis.derek_beach.AdaptivePriorCalculator.__init__").get("auto_param_L4849_25", 0.25),
        'structural': get_parameter_loader().get("saaaaaaa.analysis.derek_beach.AdaptivePriorCalculator.__init__").get("auto_param_L4850_26", 0.15)
    }
}

```

```

def calculate_likelihood_adaptativo(
    self,
    evidence_dict: dict[str, Any],
    test_type: str = 'hoop'
) -> dict[str, Any]:
    """
    """

```

PROMPT I-1: Calcula likelihood adaptativo con BF y dominios

Args:

evidence\_dict: Evidencia por caso {semantic, temporal, financial, structural}  
test\_type: Tipo de test evidencial (straw, hoop, smoking, doubly)

Returns:

Dict con p\_mechanism, BF\_used, domain\_weights, triangulation\_bonus, etc.

"""

# 1. Obtener Bayes Factor para test\_type

bf\_used = self.bf\_table.get\_bayes\_factor(test\_type)

# 2. Extraer scores por dominio

domain\_scores = {

'semantic': evidence\_dict.get('semantic', {}).get('score', get\_parameter\_loader().get("saaaaaaa.analysis.derek\_beach.AdaptivePriorCalculator.\_\_init\_\_").get("auto\_param\_L4873\_71", 0.0)),

'temporal': evidence\_dict.get('temporal', {}).get('score', get\_parameter\_loader().get("saaaaaaa.analysis.derek\_beach.AdaptivePriorCalculator.\_\_init\_\_").get("auto\_param\_L4874\_71", 0.0)),

'financial': evidence\_dict.get('financial', {}).get('score', get\_parameter\_loader().get("saaaaaaa.analysis.derek\_beach.AdaptivePriorCalculator.\_\_init\_\_").get("auto\_param\_L4875\_73", 0.0)),

'structural': evidence\_dict.get('structural', {}).get('score', get\_parameter\_loader().get("saaaaaaa.analysis.derek\_beach.AdaptivePriorCalculator.\_\_init\_\_").get("auto\_param\_L4876\_75", 0.0))

}

# 3. Ajustar pesos si falta dominio (baja peso a 0, reparte)

adjusted\_weights = self.\_adjust\_domain\_weights(domain\_scores)

# 4. Calcular score combinado normalizado

combined\_score = sum(

domain\_scores[domain] \* adjusted\_weights[domain]  
for domain in domain\_scores

)

# 5. Aplicar multiplicador BF normalizado

all bfs = [np.mean(bf\_range) for bf\_range in self.bf\_table.FACTORS.values()]

mean\_bf = np.mean(all\_bfs)

bf\_multiplier = bf\_used / mean\_bf

adapted\_score = combined\_score \* bf\_multiplier

# 6. Bonus de triangulación si ≥3 dominios activos

active\_domains = sum(1 for s in domain\_scores.values() if s > get\_parameter\_loader().get("saaaaaaa.analysis.derek\_beach.AdaptivePriorCalculator.\_\_init\_\_").get("auto\_param\_L4895\_70", 0.1))

triangulation\_bonus = get\_parameter\_loader().get("saaaaaaa.analysis.derek\_beach.AdaptivePriorCalculator.\_\_init\_\_").get("auto\_param\_L4896\_30", 0.05) if active\_domains >= 3  
else get\_parameter\_loader().get("saaaaaaa.analysis.derek\_beach.AdaptivePriorCalculator.\_\_init\_\_").get("auto\_param\_L4896\_63", 0.0)

```

    final_score = min(get_parameter_loader().get("saaaaaaa.analysis.derek_beach.AdaptivePriorCalculator.__init__").get("auto_param_L4898_26", 1.0), adapted_score + triangulation_bonus)

# 7. Transformar a probabilidad con logit inverso: p = 1/(1+exp(-(α+β·score)))
alpha = self.calibration['alpha']
beta = self.calibration['beta']
logit_value = alpha + beta * final_score
p_mechanism = get_parameter_loader().get("saaaaaaa.analysis.derek_beach.AdaptivePriorCalculator.__init__").get("auto_param_L4904_22", 1.0) / (get_parameter_loader().get("saaaaaaa.analysis.derek_beach.AdaptivePriorCalculator.__init__").get("auto_param_L4904_29", 1.0) + np.exp(-logit_value))

# 8. Clip [1e-6, 1-1e-6]
p_mechanism = np.clip(p_mechanism, 1e-6, 1 - 1e-6)

return {
    'p_mechanism': float(p_mechanism),
    'BF_used': bf_used,
    'domain_weights': adjusted_weights,
    'triangulation_bonus': triangulation_bonus,
    'calibration_params': self.calibration,
    'test_type': test_type,
    'combined_score': combined_score,
    'active_domains': active_domains
}

@calibrated_method("saaaaaaa.analysis.derek_beach.AdaptivePriorCalculator._adjust_domain_weights")
def _adjust_domain_weights(self, domain_scores: dict[str, float]) -> dict[str, float]:
    """Ajusta pesos si falta dominio: baja a 0 y reparte"""
    adjusted = self.default_domain_weights.copy()

    # Identificar dominios faltantes (score ≤ 0)
    missing_domains = [d for d, s in domain_scores.items() if s <= 0]

    if missing_domains:
        # Bajar peso a 0 para dominios faltantes
        total_missing_weight = sum(adjusted[d] for d in missing_domains)
        for d in missing_domains:
            adjusted[d] = get_parameter_loader().get("saaaaaaa.analysis.derek_beach.AdaptivePriorCalculator._adjust_domain_weights").get("auto_param_L4932_30", 0.0)

        # Repartir peso entre dominios activos
        active_domains = [d for d in adjusted if adjusted[d] > 0]
        if active_domains:
            bonus_per_domain = total_missing_weight / len(active_domains)
            for d in active_domains:
                adjusted[d] += bonus_per_domain

        # Renormalizar para asegurar suma = 1.0
        total = sum(adjusted.values())
        if total > 0:
            adjusted = {k: v / total for k, v in adjusted.items()}

    return adjusted

def sensitivity_analysis(
    self,
    evidence_dict: dict[str, Any],
    test_type: str = 'hoop',
    perturbation: float = get_parameter_loader().get("saaaaaaa.analysis.derek_beach.AdaptivePriorCalculator._adjust_domain_weights").get("auto_param_L4952_30", 0.10)
) -> dict[str, Any]:
    """
    """

```

## PROMPT I-2: Sensibilidad, OOD y ablation evidencial

Perturba cada componente  $\pm 10\%$  y reporta  $\partial p / \partial \text{component}$  top-3.

Ejecuta ablaciones: sólo textual, sólo financiero, sólo estructural.

CRITERIA:

- $|\delta_p|_{\text{max}} \leq \text{get\_parameter\_loader().get("saaaaaaa.analysis.derek_b each.AdaptivePriorCalculator._adjust\_domain\_weights").get("auto\_param\_L4961\_38", 0.15)}$
- $\text{sign\_concordance} \geq 2/3$
- $\text{OOD\_drop} \leq \text{get\_parameter\_loader().get("saaaaaaa.analysis.derek\_beach.AdaptivePriorCalculator._adjust\_domain\_weights").get("auto\_param\_L4963\_21", 0.10)}$

""

```
# Baseline
baseline_result = self.calculate_likelihood_adaptativo(evidence_dict, test_type)
baseline_p = baseline_result['p_mechanism']
```

# 1. Sensibilidad por componente

```
sensitivity_map = {}
for domain in ['semantic', 'temporal', 'financial', 'structural']:
    if domain in evidence_dict and isinstance(evidence_dict[domain], dict) and
'score' in evidence_dict[domain]:
        # Perturbar +10%
        perturbed_evidence = self._perturb_evidence(evidence_dict, domain,
perturbation)
        perturbed_result =
self.calculate_likelihood_adaptativo(perturbed_evidence, test_type)
        delta_p = perturbed_result['p_mechanism'] - baseline_p

        sensitivity_map[domain] = {
            'delta_p': delta_p,
            'relative_change': delta_p / max(baseline_p, 1e-6)
        }
```

# Top-3 por magnitud

```
top_3 = sorted(
    sensitivity_map.items(),
    key=lambda x: abs(x[1]['delta_p']),
    reverse=True
)[:3]
```

# 2. Ablaciones: sólo un dominio

```
ablation_results = {}
for domain in ['semantic', 'financial', 'structural']:
    ablated_evidence = {
        domain: evidence_dict.get(domain, {'score': get_parameter_loader().get("sa
aaaaaa.analysis.derek_beach.AdaptivePriorCalculator._adjust_domain_weights").get("auto_para
m_L4994_60", 0.0)})
    }
    if ablated_evidence[domain].get('score', 0) > 0:
        abl_result = self.calculate_likelihood_adaptativo(ablated_evidence,
test_type)
        ablation_results[f'only_{domain}'] = {
            'p_mechanism': abl_result['p_mechanism'],
            'sign_match': (abl_result['p_mechanism'] > get_parameter_loader().get(
'saaaaaaaa.analysis.derek_beach.AdaptivePriorCalculator._adjust_domain_weights').get("auto_p
aram_L5000_63", 0.5)) == (baseline_p > get_parameter_loader().get("aaaaaaaa.analysis.derek_
beach.AdaptivePriorCalculator._adjust_domain_weights").get("auto_param_L5000_85", 0.5))
        }
```

# Sign concordance

```
sign_concordance = sum(
    1 for r in ablation_results.values() if r['sign_match']
) / max(len(ablation_results), 1)
```

# 3. OOD con ruido

```
ood_evidence = self._add_ood_noise(evidence_dict)
ood_result = self.calculate_likelihood_adaptativo(ood_evidence, test_type)
ood_drop = abs(baseline_p - ood_result['p_mechanism'])
```

```

# 4. Evaluación de criterios
max_sensitivity = max((abs(item[1]['delta_p']) for item in top_3), default=get_parameter_loader().get("saaaaaa.analysis.derek_beach.AdaptivePriorCalculator._adjust_domain_weights").get("auto_param_L5014_83", 0.0))
criteria_met = {
    'max_sensitivity_ok': max_sensitivity <= get_parameter_loader().get("saaaaaa.analysis.derek_beach.AdaptivePriorCalculator._adjust_domain_weights").get("auto_param_L5016_53", 0.15),
    'sign_concordance_ok': sign_concordance >= 2/3,
    'ood_drop_ok': ood_drop <= get_parameter_loader().get("saaaaaa.analysis.derek_beach.AdaptivePriorCalculator._adjust_domain_weights").get("auto_param_L5018_39", 0.10)
}

```

# Determinar si caso es frágil  
is\_fragile = not all(criteria\_met.values())

```

return {
    'influence_top3': [(domain, data['delta_p']) for domain, data in top_3],
    'delta_p_sensitivity': max_sensitivity,
    'sign_concordance': sign_concordance,
    'OOD_drop': ood_drop,
    'ablation_results': ablation_results,
    'criteria_met': criteria_met,
    'is_fragile': is_fragile,
    'recommendation': 'downgrade' if is_fragile else 'accept'
}

```

```

def _perturb_evidence(
    self,
    evidence_dict: dict[str, Any],
    domain: str,
    perturbation: float
) -> dict[str, Any]:
    """Perturba un dominio específico"""
    import copy
    perturbed = copy.deepcopy(evidence_dict)
    if domain in perturbed and isinstance(perturbed[domain], dict) and 'score' in perturbed[domain]:
        perturbed[domain]['score'] *= (get_parameter_loader().get("saaaaaa.analysis.derek_beach.AdaptivePriorCalculator._adjust_domain_weights").get("auto_param_L5045_43", 1.0) + perturbation)
        perturbed[domain]['score'] = min(get_parameter_loader().get("saaaaaa.analysis.derek_beach.AdaptivePriorCalculator._adjust_domain_weights").get("auto_param_L5046_45", 1.0), perturbed[domain]['score'])
    return perturbed

```

```

@calibrated_method("saaaaaa.analysis.derek_beach.AdaptivePriorCalculator._add_ood_noise")
def _add_ood_noise(self, evidence_dict: dict[str, Any]) -> dict[str, Any]:
    """Genera set OOD con ruido semántico y tablas malformadas"""
    import copy
    ood = copy.deepcopy(evidence_dict)

    # Agregar ruido gaussiano a todos los scores
    for domain in ood:
        if isinstance(ood[domain], dict) and 'score' in ood[domain]:
            noise = np.random.normal(0, get_parameter_loader().get("saaaaaa.analysis.derek_beach.AdaptivePriorCalculator._add_ood_noise").get("auto_param_L5058_44", 0.05)) # 5% noise
            ood[domain]['score'] = np.clip(ood[domain]['score'] + noise, get_parameter_loader().get("saaaaaa.analysis.derek_beach.AdaptivePriorCalculator._add_ood_noise").get("auto_param_L5059_77", 0.0), get_parameter_loader().get("saaaaaa.analysis.derek_beach.AdaptivePriorCalculator._add_ood_noise").get("auto_param_L5059_82", 1.0))

    return ood

```

```

def generate_traceability_record()

```

```

self,
evidence_dict: dict[str, Any],
test_type: str,
result: dict[str, Any],
seed: int = 42
) -> dict[str, Any]:
"""
PROMPT I-3: Trazabilidad y reproducibilidad

Con semilla fija, guarda bf_table_version, weights_version,
snippets textuales con offsets, campos financieros usados.

METRICS:
- Re-ejecución con misma semilla produce hash_result idéntico
- trace_completeness ≥ get_parameter_loader().get("saaaaaaa.analysis.derek_beach.AdaptivePriorCalculator._add_ood_noise").get("auto_param_L5078_31", 0.95)
"""
# Fijar semilla para reproducibilidad
np.random.seed(seed)

# Construir evidence trace
evidence_trace = []
for domain, data in evidence_dict.items():
    if isinstance(data, dict) and 'score' in data:
        trace_item = {
            'source': domain,
            'line_span': data.get('line_span', 'unknown'),
            'transform_before': data.get('raw_value', None),
            'transform_after': data['score'],
            'snippet': data.get('snippet', "")[:100] # Primeros 100 chars
        }
        evidence_trace.append(trace_item)

# Config hash
config_str = json.dumps({
    'bf_table_version': self.bf_table.get_version(),
    'calibration_params': self.calibration,
    'domain_weights': self.default_domain_weights,
    'test_type': test_type,
    'seed': seed
}, sort_keys=True)

config_hash = hashlib.sha256(config_str.encode()).hexdigest()[:16]

# Result hash
result_str = json.dumps(result, sort_keys=True)
result_hash = hashlib.sha256(result_str.encode()).hexdigest()[:16]

# Trace completeness
factors_in_trace = len(evidence_trace)
total_factors = len([d for d in evidence_dict if isinstance(evidence_dict.get(d), dict)])
trace_completeness = factors_in_trace / max(total_factors, 1)

return {
    'evidence_trace': evidence_trace,
    'hash_config': config_hash,
    'hash_result': result_hash,
    'seed': seed,
    'bf_table_version': self.bf_table.get_version(),
    'weights_version': 'default_vget_parameter_loader().get("saaaaaaa.analysis.derek_beach.AdaptivePriorCalculator._add_ood_noise").get("auto_param_L5122_41", 1.0)',
    'trace_completeness': trace_completeness,
    'reproducibility_guaranteed': trace_completeness >= get_parameter_loader().get("saaaaaaa.analysis.derek_beach.AdaptivePriorCalculator._add_ood_noise").get("auto_param_L5124_64", 0.95)
}

```

```

@calibrated_method("saaaaaaa.analysis.derek_beach.AdaptivePriorCalculator.validate_quality_criteria")
def validate_quality_criteria(self, validation_samples: list[dict[str, Any]]) ->
dict[str, Any]:
"""
    Valida criterios de calidad en conjunto de validación sintética

    QUALITY CRITERIA:
    - BrierScore ≤ get_parameter_loader().get("saaaaaaa.analysis.derek_beach.AdaptivePriorCalculator.validate_quality_criteria").get("auto_param_L5133_23", 0.20)
    - ACE ∈ [-get_parameter_loader().get("saaaaaaa.analysis.derek_beach.AdaptivePriorCalculator.validate_quality_criteria").get("auto_param_L5134_18", 0.02), get_parameter_loader().get("saaaaaaa.analysis.derek_beach.AdaptivePriorCalculator.validate_quality_criteria").get("auto_param_L5134_24", 0.02)]
    - Cobertura CI95% ∈ [92%, 98%]
    - Monotonicidad verificada
"""

predictions = []
actuals = []

for sample in validation_samples:
    evidence = sample.get('evidence', {})
    actual_label = sample.get('actual_label', get_parameter_loader().get("saaaaaaa.analysis.derek_beach.AdaptivePriorCalculator.validate_quality_criteria").get("auto_param_L5143_54", 0.5))
    test_type = sample.get('test_type', 'hoop')

    result = self.calculate_likelihood_adaptativo(evidence, test_type)
    predictions.append(result['p_mechanism'])
    actuals.append(actual_label)

predictions = np.array(predictions)
actuals = np.array(actuals)

# 1. Brier Score
brier_score = np.mean((predictions - actuals) ** 2)
brier_ok = brier_score <= get_parameter_loader().get("saaaaaaa.analysis.derek_beach.AdaptivePriorCalculator.validate_quality_criteria").get("auto_param_L5155_34", 0.20)

# 2. ACE (Average Calibration Error)
# Dividir en bins
n_bins = 10
bin_boundaries = np.linspace(0, 1, n_bins + 1)
ace = get_parameter_loader().get("saaaaaaa.analysis.derek_beach.AdaptivePriorCalculator.validate_quality_criteria").get("ace", 0.0) # Refactored

for i in range(n_bins):
    bin_mask = (predictions >= bin_boundaries[i]) & (predictions < bin_boundaries[i + 1])
    if bin_mask.sum() > 0:
        bin_accuracy = actuals[bin_mask].mean()
        bin_confidence = predictions[bin_mask].mean()
        ace += abs(bin_accuracy - bin_confidence) / n_bins

ace_ok = -get_parameter_loader().get("saaaaaaa.analysis.derek_beach.AdaptivePriorCalculator.validate_quality_criteria").get("auto_param_L5170_18", 0.02) <= ace <= get_parameter_loader().get("saaaaaaa.analysis.derek_beach.AdaptivePriorCalculator.validate_quality_criteria").get("auto_param_L5170_33", 0.02)

# 3. Cobertura CI95%
# Simular con bootstrap
n_bootstrap = 100
coverage_count = 0

for _ in range(n_bootstrap):
    idx = np.random.choice(len(predictions), size=len(predictions), replace=True)
    boot_preds = predictions[idx]
    boot_actuals = actuals[idx]

```

```

# Calcular CI95%
ci_low = np.percentile(boot_preds, 2.5)
ci_high = np.percentile(boot_preds, 97.5)

# Verificar si mean actual está dentro
actual_mean = boot_actuals.mean()
if ci_low <= actual_mean <= ci_high:
    coverage_count += 1

coverage = coverage_count / n_bootstrap
coverage_ok = get_parameter_loader().get("saaaaaaa.analysis.derek_beach.AdaptivePriorCalculator.validate_quality_criteria").get("auto_param_L5192_22", 0.92) <= coverage <= get_parameter_loader().get("saaaaaaa.analysis.derek_beach.AdaptivePriorCalculator.validate_quality_criteria").get("auto_param_L5192_42", 0.98)

# 4. Monotonicidad: verificar que ↑ señales → ↘ p_mechanism
monotonicity_violations = 0

for i in range(len(validation_samples) - 1):
    current_total = sum(
        validation_samples[i]['evidence'].get(d, {}).get('score', 0)
        for d in ['semantic', 'temporal', 'financial', 'structural']
    )
    next_total = sum(
        validation_samples[i + 1]['evidence'].get(d, {}).get('score', 0)
        for d in ['semantic', 'temporal', 'financial', 'structural']
    )

    if next_total > current_total and predictions[i + 1] < predictions[i]:
        monotonicity_violations += 1

monotonicity_ok = monotonicity_violations == 0

return {
    'brier_score': float(brier_score),
    'brier_ok': brier_ok,
    'ace': float(ace),
    'ace_ok': ace_ok,
    'ci95_coverage': float(coverage),
    'coverage_ok': coverage_ok,
    'monotonicity_violations': monotonicity_violations,
    'monotonicity_ok': monotonicity_ok,
    'all_criteria_met': brier_ok and ace_ok and coverage_ok and monotonicity_ok,
    'quality_grade': 'EXCELLENT' if (brier_ok and ace_ok and coverage_ok and monotonicity_ok) else 'NEEDS IMPROVEMENT'
}

```

```

# =====
# AGUJA II: MODELO GENERATIVO JERÁRQUICO
# =====

```

class HierarchicalGenerativeModel:

"""

AGUJA II - Modelo Generativo Jerárquico con inferencia MCMC

PROMPT II-1: Inferencia jerárquica con incertidumbre  
Estima posterior(mechanism\_type, activity\_sequence | obs) con MCMC.

PROMPT II-2: Posterior Predictive Checks + Ablation  
Genera datos simulados desde posterior y compara con observados.

PROMPT II-3: Independencias y parsimonia  
Verifica d-separaciones y calcula ΔWAIC.

QUALITY CRITERIA:

- R-hat ≤ 1.10
- ESS ≥ 200

```

- entropy/entropy_max < get_parameter_loader().get("saaaaaaa.analysis.derek_beach.AdaptivePriorCalculator.validate_quality_criteria").get("auto_param_L5245_28", 0.7) para certeza
- ppd_p_value ∈ [get_parameter_loader().get("saaaaaaa.analysis.derek_beach.AdaptivePriorCalculator.validate_quality_criteria").get("auto_param_L5246_21", 0.1), get_parameter_loader().get("saaaaaaa.analysis.derek_beach.AdaptivePriorCalculator.validate_quality_criteria").get("auto_param_L5246_26", 0.9)]
- ΔWAIC ≤ -2 para preferir jerárquico
"""

```

```

def __init__(self, mechanism_priors: dict[str, float] | None = None) -> None:
    self.logger = logging.getLogger(self.__class__.__name__)

    # Priors débiles para mechanism_type si no se proveen
    self.mechanism_priors = mechanism_priors or {
        'administrativo': get_parameter_loader().get("saaaaaaa.analysis.derek_beach.HierarchicalGenerativeModel.__init__").get("auto_param_L5255_30", 0.30),
        'técnico': get_parameter_loader().get("saaaaaaa.analysis.derek_beach.HierarchicalGenerativeModel.__init__").get("auto_param_L5256_23", 0.25),
        'financiero': get_parameter_loader().get("saaaaaaa.analysis.derek_beach.HierarchicalGenerativeModel.__init__").get("auto_param_L5257_26", 0.20),
        'político': get_parameter_loader().get("saaaaaaa.analysis.derek_beach.HierarchicalGenerativeModel.__init__").get("auto_param_L5258_24", 0.15),
        'mixto': get_parameter_loader().get("saaaaaaa.analysis.derek_beach.HierarchicalGenerativeModel.__init__").get("auto_param_L5259_21", 0.10)
    }

    # Validar que suman ~get_parameter_loader().get("saaaaaaa.analysis.derek_beach.HierarchicalGenerativeModel.__init__").get("auto_param_L5262_29", 1.0)
    prior_sum = sum(self.mechanism_priors.values())
    if abs(prior_sum - get_parameter_loader().get("saaaaaaa.analysis.derek_beach.HierarchicalGenerativeModel.__init__").get("auto_param_L5264_27", 1.0)) > get_parameter_loader().get("saaaaaaa.analysis.derek_beach.HierarchicalGenerativeModel.__init__").get("auto_param_L5264_34", 0.01):
        self.logger.warning(f"Mechanism priors sum to {prior_sum:.3f}, normalizing...")
    self.mechanism_priors = {
        k: v / prior_sum for k, v in self.mechanism_priors.items()
    }

```

```

def infer_mechanism_posterior(
    self,
    observations: dict[str, Any],
    n_iter: int = 500,
    burn_in: int = 100,
    n_chains: int = 2
) -> dict[str, Any]:
"""

```

PROMPT II-1: Inferencia jerárquica con MCMC

Estima posterior(mechanism\_type, activity\_sequence | obs) usando MCMC.

Args:

```

observations: Dict con {verbos, co_ocurrencias, coherence, structural_signals}
n_iter: Iteraciones MCMC (≥500)
burn_in: Burn-in iterations (≥100)
n_chains: Número de cadenas para R-hat (≥2)

```

Returns:

```

Dict con type_posterior, sequence_mode, coherence_score, entropy, CI95, R-hat,
ESS
"""

```

```

self.logger.info(f"Starting MCMC inference: {n_iter} iter, {burn_in} burn-in,
{n_chains} chains")

```

# Validar observaciones mínimas

if not observations or 'coherence' not in observations:

```

self.logger.warning("Missing observations, using weak priors")

```

```

observations = observations or []
observations.setdefault('coherence', get_parameter_loader().get("saaaaaaa.analysis.derek_beach.HierarchicalGenerativeModel.__init__").get("auto_param_L5297_49", 0.5))

# Ejecutar múltiples cadenas para diagnóstico
chains = []
for chain_idx in range(n_chains):
    chain_samples = self._run_mcmc_chain(
        observations, n_iter, burn_in, seed=42 + chain_idx
    )
    chains.append(chain_samples)
    self.logger.debug(f"Chain {chain_idx + 1}/{n_chains} completed: {len(chain_samples)} samples")

# Agregar samples de todas las cadenas
all_samples = []
for chain in chains:
    all_samples.extend(chain)

# 1. Type posterior (frecuencias de mechanism_type)
type_counts = dict.fromkeys(self.mechanism_priors.keys(), 0)
for sample in all_samples:
    mtype = sample.get('mechanism_type', 'mixto')
    if mtype in type_counts:
        type_counts[mtype] += 1

total_samples = len(all_samples)
type_posterior = {
    mtype: count / max(total_samples, 1)
    for mtype, count in type_counts.items()
}

# 2. Sequence mode (secuencia más frecuente)
sequence_mode = self._get_mode_sequence(all_samples)

# 3. Coherence score (estadísticas)
coherence_scores = [s.get('coherence', get_parameter_loader().get("saaaaaaa.analysis.derek_beach.HierarchicalGenerativeModel.__init__").get("auto_param_L5330_47", 0.5)) for s in all_samples]
coherence_mean = float(np.mean(coherence_scores))
coherence_std = float(np.std(coherence_scores))

# 4. Entropy del posterior
posterior_probs = list(type_posterior.values())
entropy_posterior = -sum(p * np.log(p + 1e-10) for p in posterior_probs if p > 0)
max_entropy = np.log(len(self.mechanism_priors))
normalized_entropy = entropy_posterior / max_entropy if max_entropy > 0 else get_parameter_loader().get("saaaaaaa.analysis.derek_beach.HierarchicalGenerativeModel.__init__").get("auto_param_L5338_85", 0.0)

# 5. CI95 para coherence
ci95_low = float(np.percentile(coherence_scores, 2.5))
ci95_high = float(np.percentile(coherence_scores, 97.5))

# 6. R-hat aproximado (between-chain variance / within-chain variance)
r_hat = self._calculate_r_hat(chains)

# 7. ESS (Effective Sample Size)
ess = self._calculate_ess(all_samples)

# 8. Verificar criterios de calidad
is_uncertain = normalized_entropy > get_parameter_loader().get("saaaaaaa.analysis.derek_beach.HierarchicalGenerativeModel.__init__").get("auto_param_L5351_44", 0.7)
criteria_met = {
    'r_hat_ok': r_hat <= 1.10,
    'ess_ok': ess >= 200,
    'entropy_ok': not is_uncertain
}

```

```

# Warning si alta incertidumbre
warning = None
if is_uncertain:
    warning = f"HIGH_UNCERTAINTY: entropy/entropy_max = {normalized_entropy:.3f} >
get_parameter_loader().get("saaaaaa.analysis.derek_beach.HierarchicalGenerativeModel.__in
it__").get("auto_param_L5361_91", 0.7)"
    self.logger.warning(warning)

return {
    'type_posterior': type_posterior,
    'sequence_mode': sequence_mode,
    'coherence_score': coherence_mean,
    'coherence_std': coherence_std,
    'entropy_posterior': float(entropy_posterior),
    'normalized_entropy': float(normalized_entropy),
    'CI95': (ci95_low, ci95_high),
    'CI95_width': ci95_high - ci95_low,
    'R_hat': float(r_hat),
    'ESS': float(ess),
    'n_samples': total_samples,
    'is_uncertain': is_uncertain,
    'criteria_met': criteria_met,
    'warning': warning
}

def _run_mcmc_chain(
    self,
    observations: dict[str, Any],
    n_iter: int,
    burn_in: int,
    seed: int
) -> list[dict[str, Any]]:
    """Ejecuta una cadena MCMC con Metropolis-Hastings"""
    np.random.seed(seed)
    samples = []

    # Estado inicial: sample desde prior
    current_type = np.random.choice(
        list(self.mechanism_priors.keys()),
        p=list(self.mechanism_priors.values())
    )
    current_coherence = observations.get('coherence', get_parameter_loader().get("saaa
aaa.analysis.derek_beach.HierarchicalGenerativeModel.__init__").get("auto_param_L5397_58",
0.5))

    for i in range(n_iter):
        # Proponer nuevo mechanism_type
        proposed_type = np.random.choice(list(self.mechanism_priors.keys()))

        # Calcular likelihood ratio
        current_likelihood = self._calculate_likelihood(current_type, observations)
        proposed_likelihood = self._calculate_likelihood(proposed_type, observations)

        # Prior ratio
        prior_ratio = self.mechanism_priors[proposed_type] /
max(self.mechanism_priors[current_type], 1e-10)

        # Acceptance probability (Metropolis-Hastings)
        likelihood_ratio = proposed_likelihood / max(current_likelihood, 1e-10)
        acceptance_prob = min(get_parameter_loader().get("saaaaaa.analysis.derek_beach
.HierarchicalGenerativeModel.__init__").get("auto_param_L5412_34", 1.0), likelihood_ratio
* prior_ratio)

        # Accept/reject
        if np.random.random() < acceptance_prob:
            current_type = proposed_type

```

```

# Simular coherence con ruido
simulated_coherence = current_coherence + np.random.normal(0, get_parameter_loader().get("saaaaaa.analysis.derek_beach.HierarchicalGenerativeModel.__init__").get("auto_param_L5419_74", 0.05))
simulated_coherence = np.clip(simulated_coherence, get_parameter_loader().get("saaaaaa.analysis.derek_beach.HierarchicalGenerativeModel.__init__").get("auto_param_L5420_63", 0.0), get_parameter_loader().get("saaaaaa.analysis.derek_beach.HierarchicalGenerativeModel.__init__").get("auto_param_L5420_68", 1.0))

# Almacenar sample (después de burn-in)
if i >= burn_in:
    sample = {
        'mechanism_type': current_type,
        'coherence': float(simulated_coherence),
        'iteration': i - burn_in,
        'chain_seed': seed
    }
    samples.append(sample)

return samples

def _calculate_likelihood(
    self,
    mechanism_type: str,
    observations: dict[str, Any]
) -> float:
    """Calcula likelihood de observations dado mechanism_type"""
    # Likelihood basado en coherence y structural signals
    coherence = observations.get('coherence', get_parameter_loader().get("saaaaaa.analysis.derek_beach.HierarchicalGenerativeModel.__init__").get("auto_param_L5441_50", 0.5))
    structural_signals = observations.get('structural_signals', {})

    # Base likelihood desde prior
    prior = self.mechanism_priors.get(mechanism_type, get_parameter_loader().get("saaa.analysis.derek_beach.HierarchicalGenerativeModel.__init__").get("auto_param_L5445_58", 0.1))

    # Ajuste por coherence (mayor coherence → mayor likelihood)
    coherence_factor = get_parameter_loader().get("saaaaaaaa.analysis.derek_beach.HierarchicalGenerativeModel.__init__").get("auto_param_L5448_27", 1.0) + coherence

    # Ajuste por señales estructurales específicas del tipo
    structural_match = get_parameter_loader().get("saaaaaaaa.analysis.derek_beach.HierarchicalGenerativeModel.__init__").get("structural_match", 0.0) # Refactored
    if mechanism_type == 'administrativo' and structural_signals.get('admin_keywords', 0) > 0:
        structural_match = get_parameter_loader().get("saaaaaaaa.analysis.derek_beach.HierarchicalGenerativeModel.__init__").get("structural_match", 0.2) # Refactored
    elif mechanism_type == 'financiero' and structural_signals.get('budget_data', 0) > 0:
        structural_match = get_parameter_loader().get("saaaaaaaa.analysis.derek_beach.HierarchicalGenerativeModel.__init__").get("structural_match", 0.3) # Refactored
    elif mechanism_type == 'técnico' and structural_signals.get('technical_terms', 0) > 0:
        structural_match = get_parameter_loader().get("saaaaaaaa.analysis.derek_beach.HierarchicalGenerativeModel.__init__").get("structural_match", 0.25) # Refactored

    likelihood = prior * coherence_factor * (get_parameter_loader().get("saaaaaaaa.analysis.derek_beach.HierarchicalGenerativeModel.__init__").get("auto_param_L5459_49", 1.0) + structural_match)
    return likelihood

@calibrated_method("saaaaaaaa.analysis.derek_beach.HierarchicalGenerativeModel._get_mode_sequence")
def _get_mode_sequence(self, samples: list[dict[str, Any]]) -> str:
    """Obtiene secuencia modal (tipo más frecuente)"""
    type_counts = {}
    for s in samples:

```

```

mtype = s.get('mechanism_type', 'mixto')
type_counts[mtype] = type_counts.get(mtype, 0) + 1

if type_counts:
    return max(type_counts.items(), key=lambda x: x[1])[0]
return 'mixto'

@calibrated_method("saaaaaaa.analysis.derek_beach.HierarchicalGenerativeModel._calculat
e_r_hat")
def _calculate_r_hat(self, chains: list[list[dict[str, Any]]]) -> float:
    """Calcula Gelman-Rubin R-hat para diagnóstico de convergencia"""
    if len(chains) < 2:
        return get_parameter_loader().get("saaaaaaa.analysis.derek_beach.HierarchicalGe
nerativeModel._calculate_r_hat").get("auto_param_L5478_19", 1.0)

    # Extraer coherence de cada cadena
    chain_means = []
    chain_vars = []

    for chain in chains:
        coherences = [s.get('coherence', get_parameter_loader().get("saaaaaaa.analysis.
derek_beach.HierarchicalGenerativeModel._calculate_r_hat").get("auto_param_L5485_45",
0.5)) for s in chain]
        if len(coherences) > 0:
            chain_means.append(np.mean(coherences))
            chain_vars.append(np.var(coherences, ddof=1))

    if len(chain_means) < 2:
        return get_parameter_loader().get("saaaaaaa.analysis.derek_beach.HierarchicalGe
nerativeModel._calculate_r_hat").get("auto_param_L5491_19", 1.0)

    # Between-chain variance (B)
    n = len(chains[0]) # samples per chain
    B = np.var(chain_means, ddof=1) * n

    # Within-chain variance (W)
    W = np.mean(chain_vars)

    # R-hat estimator
    if W > 0:
        var_plus = ((n - 1) / n) * W + (1 / n) * B
        r_hat = np.sqrt(var_plus / W)
    else:
        r_hat = get_parameter_loader().get("saaaaaaa.analysis.derek_beach.HierarchicalG
enerativeModel._calculate_r_hat").get("r_hat", 1.0) # Refactored

    return float(r_hat)

@calibrated_method("saaaaaaa.analysis.derek_beach.HierarchicalGenerativeModel._calculat
e_ess")
def _calculate_ess(self, samples: list[dict[str, Any]]) -> float:
    """Calcula Effective Sample Size (simplificado)"""
    n = len(samples)

    # Estimar autocorrelación
    coherences = np.array([s.get('coherence', get_parameter_loader().get("saaaaaaa.anal
ysis.derek_beach.HierarchicalGenerativeModel._calculate_ess").get("auto_param_L5515_50",
0.5)) for s in samples])

    if len(coherences) < 2:
        return n

    # Lag-1 autocorrelation
    mean_coh = np.mean(coherences)
    var_coh = np.var(coherences)

    if var_coh > 0:
        lag1_autocorr = np.mean(

```

```

(coherences[:-1] - mean_coh) * (coherences[1:] - mean_coh)
) / var_coh
else:
    lag1_autocorr = get_parameter_loader().get("saaaaaaa.analysis.derek_beach.Hiera
rchicalGenerativeModel._calculate_ess").get("lag1_autocorr", 0.0) # Refactored

# ESS approximation
ess = n / (1 + 2 * max(0, lag1_autocorr))
return float(ess)

def posterior_predictive_check(
    self,
    posterior_samples: list[dict[str, Any]],
    observed_data: dict[str, Any]
) -> dict[str, Any]:
    """
    PROMPT II-2: Posterior Predictive Checks + Ablation

    Genera datos simulados desde posterior y compara con observados.
    Realiza ablation de pasos de secuencia.

    Args:
        posterior_samples: Samples del posterior MCMC
        observed_data: Datos observados reales

    Returns:
        Dict con ppd_p_value, distance_metric, ablation_curve, criteria_met
    """
    self.logger.info("Running posterior predictive checks...")

    # 1. Generar datos predictivos desde posterior
    n_ppd_samples = min(100, len(posterior_samples))
    ppd_samples = []

    for _i in range(n_ppd_samples):
        sample_idx = np.random.randint(0, len(posterior_samples))
        posterior_sample = posterior_samples[sample_idx]

        # Simular coherence desde distribución posterior
        simulated_coherence = posterior_sample.get('coherence', get_parameter_loader()
            .get("saaaaaaa.analysis.derek_beach.HierarchicalGenerativeModel._calculate_ess")
            .get("auto_param_L5564_68", 0.5)) + np.random.normal(0, get_parameter_loader()
            .get("saaaaaaa.analysis.derek_beach.HierarchicalGenerativeModel._calculate_ess")
            .get("auto_param_L5564_95", 0.05))

        simulated_coherence = np.clip(simulated_coherence, get_parameter_loader()
            .get("saaaaaaa.analysis.derek_beach.HierarchicalGenerativeModel._calculate_ess")
            .get("auto_param_L5565_63", 0.0), get_parameter_loader()
            .get("saaaaaaa.analysis.derek_beach.HierarchicalGenerativeModel._calculate_ess")
            .get("auto_param_L5565_68", 1.0))
        ppd_samples.append(simulated_coherence)

    ppd_samples = np.array(ppd_samples)

    # 2. Comparar con observado usando KS test
    observed_coherence = observed_data.get('coherence', get_parameter_loader()
        .get("saaaaaaa.analysis.derek_beach.HierarchicalGenerativeModel._calculate_ess")
        .get("auto_param_L571_60", 0.5))

    # KS test: comparar distribución PPD con punto observado
    from scipy.stats import kstest
    ks_stat, ppd_p_value = kstest(ppd_samples, lambda x: 0 if x < observed_coherence
    else 1)
    ppd_p_value = float(ppd_p_value)

    # 3. Ablation de secuencia
    ablation_curve = self._ablation_analysis(posterior_samples, observed_data)

    # 4. Verificar criterios
    ppd_ok = get_parameter_loader()
        .get("saaaaaaa.analysis.derek_beach.HierarchicalGenerativeModel._calculate_ess")
        .get("auto_param_L5582_17", 0.1) <= ppd_p_value <= get_parameter_loader()
        .get("saaaaaaa.analysis.derek_beach.HierarchicalGenerativeModel._calculate_ess")
        .get("auto_param_L5582_95", 0.9)

```

```

er_loader().get("saaaaaaa.analysis.derek_beach.HierarchicalGenerativeModel._calculate_ess")
.get("auto_param_L5582_39", 0.9)
ablation_ok = all(delta >= -get_parameter_loader().get("saaaaaaa.analysis.derek_bea
ch.HierarchicalGenerativeModel._calculate_ess").get("auto_param_L5583_36", 0.05) for delta
in ablation_curve.values()) # Tolerancia -5%

```

criteria\_met = {  
 'ppd\_p\_value\_ok': ppd\_ok,  
 'ablation\_ok': ablation\_ok  
}

# Recomendación  
if ppd\_ok and ablation\_ok:  
 recommendation = 'accept'  
else:  
 recommendation = 'rebaja\_posterior'  
self.logger.warning(f"PPC failed: ppd\_p={ppd\_p\_value:.3f},  
ablation\_ok={ablation\_ok}")

```

return {
    'ppd_p_value': ppd_p_value,
    'ppd_samples_mean': float(np.mean(ppd_samples)),
    'ppd_samples_std': float(np.std(ppd_samples)),
    'distance_metric': 'KS',
    'ks_statistic': float(ks_stat),
    'ablation_curve': ablation_curve,
    'criteria_met': criteria_met,
    'recommendation': recommendation
}

```

def \_ablation\_analysis(  
 self,  
 posterior\_samples: list[dict[str, Any]],  
 observed\_data: dict[str, Any]
) -> dict[str, float]:  
 """Mide caída en coherence al quitar pasos de secuencia"""
 baseline\_coherence = np.mean([s.get('coherence', get\_parameter\_loader().get("saaaa
aa.analysis.derek\_beach.HierarchicalGenerativeModel.\_calculate\_ess").get("auto\_param\_L5614
\_57", 0.5)) for s in posterior\_samples])

# Simular ablación de pasos clave  
# En práctica real, esto requeriría re-ejecutar modelo sin ciertos steps  
ablation\_deltas = {  
 'remove\_step\_diagnostic': baseline\_coherence - (baseline\_coherence \* get\_param
eter\_loader().get("saaaaaaa.analysis.derek\_beach.HierarchicalGenerativeModel.\_calculate\_ess
").get("auto\_param\_L5619\_81", 0.95)), # -5%
 'remove\_step\_planning': baseline\_coherence - (baseline\_coherence \* get\_paramet
er\_loader().get("saaaaaaa.analysis.derek\_beach.HierarchicalGenerativeModel.\_calculate\_ess")
.get("auto\_param\_L5620\_79", 0.85)), # -15%
 'remove\_step\_execution': baseline\_coherence - (baseline\_coherence \* get\_param
eter\_loader().get("saaaaaaa.analysis.derek\_beach.HierarchicalGenerativeModel.\_calculate\_ess"
).get("auto\_param\_L5621\_80", 0.90)), # -10%
 'remove\_step\_monitoring': baseline\_coherence - (baseline\_coherence \* get\_param
eter\_loader().get("saaaaaaa.analysis.derek\_beach.HierarchicalGenerativeModel.\_calculate\_ess
").get("auto\_param\_L5622\_81", 0.97)) # -3%
}

return ablation\_deltas

```

def verify_conditional_independence(
    self,
    dag: nx.DiGraph,
    independence_tests: list[tuple[str, str, list[str]]] | None = None
) -> dict[str, Any]:
    """

```

PROMPT II-3: Independencias y parsimonia

Verifica d-separaciones implicadas por el DAG.

Calcula  $\Delta$ WAIC entre modelo jerárquico vs. nulo.

Args:

dag: NetworkX DiGraph del modelo causal  
independence\_tests: Lista de tuplas (X, Y, Z) para test  $X \perp\!\!\!\perp Y | Z$

Returns:

Dict con independence\_tests, delta\_waic, model\_preference, criteria\_met

"""

self.logger.info("Verifying conditional independencies...")

# 1. Tests de independencia (d-separación)

test\_results = []

if independence\_tests is None:

# Generar tests automáticamente si no se proveen

independence\_tests = self.\_generate\_independence\_tests(dag)

for x, y, z\_set in independence\_tests:

try:

# Verificar d-separación en DAG

is\_independent = nx.d\_separated(dag, {x}, {y}, set(z\_set))

test\_results.append({

'test': f'{x} ⊥\!\!\!\perp {y} | {{', '.join(z\_set)}}}',

'x': x,

'y': y,

'z': z\_set,

'passed': is\_independent

})

except Exception as e:

self.logger.warning(f"Independence test failed: {x} ⊥\!\!\!\perp {y} | {z\_set} -

{e}")

test\_results.append({

'test': f'{x} ⊥\!\!\!\perp {y} | {{', '.join(z\_set)}}}',

'x': x,

'y': y,

'z': z\_set,

'passed': False,

'error': str(e)

})

tests\_passed = sum(1 for t in test\_results if t['passed'])

# 2. Calcular  $\Delta$ WAIC (simplificado)

# En práctica real: usar librería como arviz para WAIC calculation  
delta\_waic = self.\_calculate\_waic\_difference(dag)

# 3. Verificar criterios

independence\_ok = tests\_passed >= 2

waic\_ok = delta\_waic <= -2.0

# 4. Preferencia de modelo

if independence\_ok and waic\_ok:

model\_preference = 'hierarchical'

elif not waic\_ok:

model\_preference = 'inconclusive'

else:

model\_preference = 'null'

criteria\_met = {

'independence\_ok': independence\_ok,

'waic\_ok': waic\_ok

}

return {

'independence\_tests': test\_results,

'tests\_passed': tests\_passed,

'tests\_total': len(test\_results),

```

'delta_waic': float(delta_waic),
'model_preference': model_preference,
'criteria_met': criteria_met
}

def _generate_independence_tests(
    self,
    dag: nx.DiGraph,
    n_tests: int = 3
) -> list[tuple[str, str, list[str]]]:
    """Genera tests de independencia automáticamente desde DAG"""
    tests = []
    nodes = list(dag.nodes())

    if len(nodes) < 3:
        return tests

    # Generar tests de forma heurística
    for _ in range(min(n_tests, len(nodes) - 2)):
        # Seleccionar nodos aleatorios
        x, y = np.random.choice(nodes, size=2, replace=False)

        # Z: padres comunes o mediadores
        z_candidates = set(dag.predecessors(x)) | set(dag.predecessors(y))
        z_set = list(z_candidates)[:2] # Máximo 2 nodos en conditioning set

        if x != y:
            tests.append((x, y, z_set))

    return tests

@calibrated_method("saaaaaa.analysis.derek_beach.HierarchicalGenerativeModel._calculate_waic_difference")
def _calculate_waic_difference(self, dag: nx.DiGraph) -> float:
    """
    Calcula ΔWAIC = WAIC_hierarchical - WAIC_null (simplificado)

    En producción: usar arviz.waic() con trace real de PyMC/Stan
    """

    # Heurística: modelos jerárquicos con más estructura (edges) son preferidos
    n_edges = dag.number_of_edges()
    dag.number_of_nodes()

    # Penalización por complejidad
    complexity_penalty = n_edges * get_parameter_loader().get("saaaaaa.analysis.derek_beach.HierarchicalGenerativeModel._calculate_waic_difference").get("auto_param_L5746_39", 0.5)

    # WAIC aproximado
    waic_hierarchical = -5get_parameter_loader().get("saaaaaa.analysis.derek_beach.HierarchicalGenerativeModel._calculate_waic_difference").get("auto_param_L5749_30", 0.0) - n_edges * 2 # Mejor fit con más estructura
    waic_null = -45.0 # Modelo nulo sin estructura

    delta_waic = waic_hierarchical - waic_null + complexity_penalty

    return delta_waic

# =====
# AGUJA III: AUDITOR CONTRAFACTUAL BAYESIANO
# =====

class BayesianCounterfactualAuditor:
    """
    AGUJA III - Auditor Contrafactual con SCM y do-calculus

    PROMPT III-1: Construcción de SCM y queries gemelas
    Construye SCM={DAG, f_i} y responde omission_impact, sufficiency_test, necessity_test.
    """

```

PROMPT III-2: Riesgo sistémico y priorización  
Agrega riesgos, propaga incertidumbre, calcula priority.

PROMPT III-3: Refutación, negativos y cordura do(.)  
Ejecuta controles negativos, pruebas placebo, sanity checks.

QUALITY CRITERIA:

- Consistencia de signos factual/contrafactual
  - effect\_stability:  $\Delta \text{effect} \leq \text{get\_parameter\_loader}().\text{get}("saaaaaa.analysis.derek_beach.HierarchicalGenerativeModel._calculate_waic_difference").\text{get}("auto_param_L5775_34", 0.15)$  al variar priors  $\pm 10\%$
  - negative\_controls: mediana  $|\text{efecto}| \leq \text{get\_parameter\_loader}().\text{get}("saaaaaa.analysis.derek_beach.HierarchicalGenerativeModel._calculate_waic_difference").\text{get}("auto_param_L5776_44", 0.05)$
  - sanity\_violations: 0
- """

```
def __init__(self) -> None:  
    self.logger = logging.getLogger(self.__class__.__name__)  
    self.scm: dict[str, Any] | None = None
```

```
def construct_scm(  
    self,  
    dag: nx.DiGraph,  
    structural_equations: dict[str, callable] | None = None  
) -> dict[str, Any]:  
    """
```

PROMPT III-1: Construcción de SCM

Construye SCM = {DAG, f\_i} desde grafo y ecuaciones estructurales.

Args:

dag: NetworkX DiGraph (debe ser acíclico)  
structural\_equations: Dict {node: function} para f\_i

Returns:

SCM con DAG validado y funciones estructurales

Raises:

ValueError: Si DAG no es acíclico

"""

```
self.logger.info(f"Constructing SCM with {dag.number_of_nodes()} nodes,  
{dag.number_of_edges()} edges")
```

# 1. Validar que DAG es acíclico

```
if not nx.is_directed_acyclic_graph(dag):
```

```
    raise ValueError("DAG must be acyclic for SCM construction. Use cycle  
detection first.")
```

# 2. Crear ecuaciones por defecto si no se proveen

```
if structural_equations is None:
```

```
    structural_equations = self._create_default_equations(dag)  
    self.logger.info(f"Created {len(structural_equations)} default structural  
equations")
```

# 3. Construir SCM

```
scm = {  
    'dag': dag,  
    'equations': structural_equations,  
    'nodes': list(dag.nodes()),  
    'edges': list(dag.edges()),  
    'topological_order': list(nx.topological_sort(dag))  
}
```

```
self.scm = scm  
self.logger.info("✓ SCM constructed successfully")  
return scm
```

```

@calibrated_method("saaaaaa.analysis.derek_beach.BayesianCounterfactualAuditor._create
_default_equations")
def _create_default_equations(self, dag: nx.DiGraph) -> dict[str, callable]:
    """Crea ecuaciones estructurales lineales por defecto"""
    equations = {}

    for node in dag.nodes():
        parents = list(dag.predecessors(node))

        if not parents:
            # Nodo raíz: variable exógena U
            def root_eq(noise=get_parameter_loader().get("saaaaaa.analysis.derek_beach
.BayesianCounterfactualAuditor._create_default_equations").get("auto_param_L5838_34",
0.0), node_name=node):
                return get_parameter_loader().get("saaaaaa.analysis.derek_beach.Bayesi
anCounterfactualAuditor._create_default_equations").get("auto_param_L5839_27", 0.5) +
noise # Prior neutral + ruido
            equations[node] = root_eq
        else:
            # Nodo con padres: función lineal
            def child_eq(parent_values, noise=get_parameter_loader().get("saaaaaa.anal
ysis.derek_beach.BayesianCounterfactualAuditor._create_default_equations").get("auto_param
_L5843_50", 0.0), node_name=node, n_parents=len(parents)):
                if isinstance(parent_values, dict):
                    return sum(parent_values.values()) / max(n_parents, 1) + noise
                return get_parameter_loader().get("saaaaaa.analysis.derek_beach.Bayesi
anCounterfactualAuditor._create_default_equations").get("auto_param_L5846_27", 0.5) +
noise
            equations[node] = child_eq

    return equations

```

```

def counterfactual_query(
    self,
    intervention: dict[str, float],
    target: str,
    evidence: dict[str, float] | None = None
) -> dict[str, Any]:
    """
    PROMPT III-1: Queries gemelas (omission, sufficiency, necessity)

```

Evaluá:

- Factual:  $P(Y | \text{evidence})$
- Counterfactual:  $P(Y | \text{do}(X=x), \text{evidence})$
- Causal effect, sufficiency, necessity

Args:

intervention: {nodo: valor} para do(.) operation  
 target: Nodo objetivo Y  
 evidence: Evidencia observada (opcional)

Returns:

Dict con p\_factual, p\_counterfactual, causal\_effect, is\_sufficient,  
 is\_necessary

if self.scm is None:

raise ValueError("SCM must be constructed first. Call construct\_scm().")

evidence = evidence or {}

self.logger.debug(f"Counterfactual query: intervention={intervention},  
 target={target}")

# 1. Factual:  $P(Y | \text{evidence})$

p\_factual = self.\_evaluate\_factual(target, evidence)

# 2. Counterfactual:  $P(Y | \text{do}(X=x), \text{evidence})$

```

p_counterfactual = self._evaluate_counterfactual(target, intervention, evidence)

# 3. Causal effect
causal_effect = p_counterfactual - p_factual

# 4. Sufficiency test:  $\text{do}(X=1) \rightarrow Y=1$ ?
intervention_node = list(intervention.keys())[0] if intervention else None
if intervention_node:
    p_y_given_do_x1 = self._evaluate_counterfactual(target, {intervention_node: ge
t_parameter_loader().get("saaaaaaa.analysis.derek_beach.BayesianCounterfactualAuditor._crea
te_default_equations").get("auto_param_L5892_88", 1.0)}, {})
    is_sufficient = p_y_given_do_x1 > get_parameter_loader().get("saaaaaaa.analysis
.derek_beach.BayesianCounterfactualAuditor._create_default_equations").get("auto_param_L58
93_46", 0.7)
else:
    is_sufficient = False

# 5. Necessity test:  $\text{do}(X=0) \rightarrow Y=0$ ?
if intervention_node:
    p_y_given_do_x0 = self._evaluate_counterfactual(target, {intervention_node: ge
t_parameter_loader().get("saaaaaaa.analysis.derek_beach.BayesianCounterfactualAuditor._crea
te_default_equations").get("auto_param_L5899_88", 0.0)}, {})
    is_necessary = p_y_given_do_x0 < get_parameter_loader().get("saaaaaaa.analysis.
derek_beach.BayesianCounterfactualAuditor._create_default_equations").get("auto_param_L590
0_45", 0.3)
else:
    is_necessary = False

# 6. Consistencia de signos
signs_consistent = (
    (causal_effect >= 0 and p_counterfactual >= p_factual) or
    (causal_effect < 0 and p_counterfactual < p_factual)
)

# 7. Effect stability
stability = self._test_effect_stability(intervention, target, evidence)

return {
    'p_factual': float(np.clip(p_factual, get_parameter_loader().get("saaaaaaa.anal
ysis.derek_beach.BayesianCounterfactualAuditor._create_default_equations").get("auto_param
_L5914_50", 0.0), get_parameter_loader().get("saaaaaaa.analysis.derek_beach.BayesianCounter
factualAuditor._create_default_equations").get("auto_param_L5914_55", 1.0))),
    'p_counterfactual': float(np.clip(p_counterfactual, get_parameter_loader().get
("saaaaaaa.analysis.derek_beach.BayesianCounterfactualAuditor._create_default_equations").g
et("auto_param_L5915_64", 0.0), get_parameter_loader().get("saaaaaaa.analysis.derek_beach.B
ayesianCounterfactualAuditor._create_default_equations").get("auto_param_L5915_69",
1.0))),
    'causal_effect': float(causal_effect),
    'is_sufficient': is_sufficient,
    'is_necessary': is_necessary,
    'signs_consistent': signs_consistent,
    'effect_stability': float(stability),
    'effect_stable': stability <= get_parameter_loader().get("saaaaaaa.analysis.der
ek_beach.BayesianCounterfactualAuditor._create_default_equations").get("auto_param_L5921_4
2", 0.15)
}

def _evaluate_factual(
    self,
    target: str,
    evidence: dict[str, float]
) -> float:
    """Evalúa P(target | evidence) propagando hacia adelante en DAG"""
    if target in evidence:
        return evidence[target]

dag = self.scm['dag']
equations = self.scm['equations']

```

```

topological_order = self.scm['topological_order']

# Evaluar nodos en orden topológico
computed_values = evidence.copy()

for node in topological_order:
    if node in computed_values:
        continue

    parents = list(dag.predecessors(node))

    if not parents:
        # Nodo raíz
        computed_values[node] = equations[node](noise=get_parameter_loader().get("saaaaaa.analysis.derek_beach.BayesianCounterfactualAuditor._create_default_equations").get("auto_param_L5948_62", 0.0))
    else:
        # Evaluar padres primero
        parent_values = {}
        for parent in parents:
            if parent not in computed_values:
                computed_values[parent] = self._evaluate_factual(parent, evidence)
                parent_values[parent] = computed_values[parent]

        # Aplicar ecuación estructural
        try:
            computed_values[node] = equations[node](parent_values, noise=get_parameter_loader().get("saaaaaa.analysis.derek_beach.BayesianCounterfactualAuditor._create_default_equations").get("auto_param_L5959_81", 0.0))
        except:
            # Fallback
            computed_values[node] = sum(parent_values.values()) / max(len(parent_values), 1)

    return float(np.clip(computed_values.get(target, get_parameter_loader().get("saaaaaa.analysis.derek_beach.BayesianCounterfactualAuditor._create_default_equations").get("auto_param_L5964_57", 0.5)), get_parameter_loader().get("saaaaaa.analysis.derek_beach.BayesianCounterfactualAuditor._create_default_equations").get("auto_param_L5964_63", 0.0), get_parameter_loader().get("saaaaaa.analysis.derek_beach.BayesianCounterfactualAuditor._create_default_equations").get("auto_param_L5964_68", 1.0)))

def _evaluate_counterfactual(
    self,
    target: str,
    intervention: dict[str, float],
    evidence: dict[str, float]
) -> float:
    """Evalúa P(target | do(intervention), evidence) con DAG mutilado"""
    # Crear DAG mutilado: quitar aristas hacia nodos intervenidos
    dag_mutilated = self.scm['dag'].copy()

    for node in intervention:
        in_edges = list(dag_mutilated.in_edges(node))
        dag_mutilated.remove_edges_from(in_edges)

    # Guardar SCM original
    original_scm = self.scm.copy()

    # Crear SCM mutilado temporalmente
    self.scm = {
        'dag': dag_mutilated,
        'equations': self.scm['equations'],
        'nodes': self.scm['nodes'],
        'edges': list(dag_mutilated.edges()),
        'topological_order': list(nx.topological_sort(dag_mutilated))
    }

    # Combinar evidence con intervention (intervention tiene prioridad)

```

```

combined_evidence = {**evidence, **intervention}

# Evaluar en SCM mutilado
result = self._evaluate_factual(target, combined_evidence)

# Restaurar SCM original
self.scm = original_scm

return result

def _test_effect_stability(
    self,
    intervention: dict[str, float],
    target: str,
    evidence: dict[str, float] | None,
    n_perturbations: int = 5
) -> float:
    """Testa estabilidad al variar priors/ecuaciones ±10%"""
    evidence = evidence or {}

    # Efecto baseline
    baseline_result = self.counterfactual_query(intervention, target, evidence)
    baseline_effect = baseline_result['causal_effect']

    # Perturbar y medir variación
    perturbed_effects = []

    for _ in range(n_perturbations):
        perturbation_factor = np.random.uniform(get_parameter_loader().get("saaaaaa.analysis.derek_beach.BayesianCounterfactualAuditor._create_default_equations").get("auto_param_L6021_52", 0.9), 1.1) # ±10%
        perturbed_evidence = {
            k: v * perturbation_factor for k, v in evidence.items()
        }

        # Re-evaluar
        try:
            result = self.counterfactual_query(intervention, target, perturbed_evidence)
            perturbed_effects.append(result['causal_effect'])
        except:
            perturbed_effects.append(baseline_effect)

    # Máxima variación
    max_variation = max(abs(e - baseline_effect) for e in perturbed_effects) if perturbed_effects else get_parameter_loader().get("saaaaaa.analysis.derek_beach.BayesianCounterfactualAuditor._create_default_equations").get("auto_param_L6036_107", 0.0)

    return max_variation

def aggregate_risk_and_prioritize(
    self,
    omission_score: float,
    insufficiency_score: float,
    unnecessity_score: float,
    causal_effect: float,
    feasibility: float = get_parameter_loader().get("saaaaaa.analysis.derek_beach.BayesianCounterfactualAuditor._create_default_equations").get("auto_param_L6046_29", 0.8),
    cost: float = get_parameter_loader().get("saaaaaa.analysis.derek_beach.BayesianCounterfactualAuditor._create_default_equations").get("auto_param_L6047_22", 1.0)
) -> dict[str, Any]:
    """
    """

```

PROMPT III-2: Riesgo sistémico y priorización con incertidumbre

Fórmulas:

- risk = get\_parameter\_loader().get("saaaaaa.analysis.derek\_beach.BayesianCounterf

```

actualAuditor._create_default_equations").get("auto_param_L6053_17", 0.50)·omission + get_
parameter_loader().get("saaaaaaa.analysis.derek_beach.BayesianCounterfactualAuditor._create
_default_equations").get("auto_param_L6053_33", 0.35)·insufficiency + get_parameter_loader
().get("saaaaaaa.analysis.derek_beach.BayesianCounterfactualAuditor._create_default_equatio
ns").get("auto_param_L6053_54", 0.15)·unnecessity
- priority = |effect|·feasibility/(cost+ε)·(1–uncertainty)

```

Args:

- omission\_score: Riesgo de omisión de mecanismo [0,1]
- insufficiency\_score: Insuficiencia del mecanismo [0,1]
- unnecessity\_score: Mecanismo innecesario [0,1]
- causal\_effect: Efecto causal estimado
- feasibility: Factibilidad de intervención [0,1]
- cost: Costo relativo (>0)

Returns:

Dict con risk\_score, success\_probability, priority, recommendations

=====

# 1. Componentes de riesgo

risk\_components = {

```

    'omission': float(np.clip(omission_score, get_parameter_loader().get("saaaaaaa.
analysis.derek_beach.BayesianCounterfactualAuditor._create_default_equations").get("auto_p
aram_L6069_54", 0.0), get_parameter_loader().get("saaaaaaa.analysis.derek_beach.BayesianCou
nterfactualAuditor._create_default_equations").get("auto_param_L6069_59", 1.0))),
    'insufficiency': float(np.clip(insufficiency_score, get_parameter_loader().get
("saaaaaaa.analysis.derek_beach.BayesianCounterfactualAuditor._create_default_equations").g
et("auto_param_L6070_64", 0.0), get_parameter_loader().get("saaaaaaa.analysis.derek_beach.B
ayesianCounterfactualAuditor._create_default_equations").get("auto_param_L6070_69",
1.0))),
    'unnecessity': float(np.clip(unnecessity_score, get_parameter_loader().get("sa
aaaaaa.analysis.derek_beach.BayesianCounterfactualAuditor._create_default_equations").get(
"auto_param_L6071_60", 0.0), get_parameter_loader().get("saaaaaaa.analysis.derek_beach.Bayes
ianCounterfactualAuditor._create_default_equations").get("auto_param_L6071_65", 1.0)))
}
```

# 2. Riesgo agregado

risk\_score = (

```

    get_parameter_loader().get("saaaaaaa.analysis.derek_beach.BayesianCounterfactua
lAuditor._create_default_equations").get("auto_param_L6076_12", 0.50) *
risk_components['omission'] +

```

```

    get_parameter_loader().get("saaaaaaa.analysis.derek_beach.BayesianCounterfactua
lAuditor._create_default_equations").get("auto_param_L6077_12", 0.35) *

```

```

risk_components['insufficiency'] +

```

```

    get_parameter_loader().get("saaaaaaa.analysis.derek_beach.BayesianCounterfactua
lAuditor._create_default_equations").get("auto_param_L6078_12", 0.15) *

```

```

risk_components['unnecessity']
)
```

```

    risk_score = float(np.clip(risk_score, get_parameter_loader().get("saaaaaaa.analysi
s.derek_beach.BayesianCounterfactualAuditor._create_default_equations").get("auto_param_L6
080_47", 0.0), get_parameter_loader().get("saaaaaaa.analysis.derek_beach.BayesianCounterfac
tualAuditor._create_default_equations").get("auto_param_L6080_52", 1.0)))

```

# 3. Success probability con incertidumbre

```

success_mean = get_parameter_loader().get("saaaaaaa.analysis.derek_beach.BayesianCo
unterfactualAuditor._create_default_equations").get("auto_param_L6083_23", 1.0) -
risk_score

```

# Incertidumbre: mayor riesgo → mayor uncertainty

```

success_std = get_parameter_loader().get("saaaaaaa.analysis.derek_beach.BayesianCou
nterfactualAuditor._create_default_equations").get("auto_param_L6086_22", 0.05) + get_para
meter_loader().get("saaaaaaa.analysis.derek_beach.BayesianCounterfactualAuditor._create_def
ault_equations").get("auto_param_L6086_29", 0.10) * risk_score # Entre 5% y 15%

```

# CI95 para success

```

ci95_low = max(get_parameter_loader().get("saaaaaaa.analysis.derek_beach.BayesianCo
unterfactualAuditor._create_default_equations").get("auto_param_L6089_23", 0.0),
success_mean - 1.96 * success_std)

```

```

ci95_high = min(get_parameter_loader().get("saaaaaaa.analysis.derek_beach.BayesianC

```

```

counterfactualAuditor._create_default_equations").get("auto_param_L6090_24", 1.0),
success_mean + 1.96 * success_std)

success_probability = {
    'mean': float(success_mean),
    'std': float(success_std),
    'CI95': (float(ci95_low), float(ci95_high))
}

# 4. Prioridad
uncertainty = success_std
epsilon = 1e-6

priority = (
    abs(causal_effect) *
    feasibility /
    (cost + epsilon) *
    (get_parameter_loader().get("saaaaaaa.analysis.derek_beach.BayesianCounterfactu
alAuditor._create_default_equations").get("auto_param_L6106_13", 1.0) - uncertainty)
)
priority = float(priority)

# 5. Recomendaciones ordenadas
recommendations = []

if risk_score > get_parameter_loader().get("saaaaaaa.analysis.derek_beach.BayesianC
ounterfactualAuditor._create_default_equations").get("auto_param_L6113_24", 0.7):
    recommendations.append("CRITICAL_RISK: Immediate intervention required")
elif risk_score > get_parameter_loader().get("saaaaaaa.analysis.derek_beach.Bayesia
nCounterfactualAuditor._create_default_equations").get("auto_param_L6115_26", 0.4):
    recommendations.append("MEDIUM_RISK: Close monitoring required")
else:
    recommendations.append("LOW_RISK: Routine surveillance")

if risk_components['omission'] > get_parameter_loader().get("saaaaaaa.analysis.dere
k_beach.BayesianCounterfactualAuditor._create_default_equations").get("auto_param_L6120_41
", 0.6):
    recommendations.append("HIGH_OMISSION_RISK: Key mechanism may be missing")

if risk_components['insufficiency'] > get_parameter_loader().get("saaaaaaa.analysis
.derek_beach.BayesianCounterfactualAuditor._create_default_equations").get("auto_param_L61
23_46", 0.5):
    recommendations.append("INSUFFICIENCY_DETECTED: Mechanism alone insufficient")

if priority > get_parameter_loader().get("saaaaaaa.analysis.derek_beach.BayesianCou
nterfactualAuditor._create_default_equations").get("auto_param_L6126_22", 0.5):
    recommendations.append("HIGH_PRIORITY: Optimal intervention candidate")
elif priority < get_parameter_loader().get("saaaaaaa.analysis.derek_beach.BayesianC
ounterfactualAuditor._create_default_equations").get("auto_param_L6128_24", 0.2):
    recommendations.append("LOW_PRIORITY: Consider alternative interventions")

# 6. Verificar criterios de calidad
ci95_valid = get_parameter_loader().get("saaaaaaa.analysis.derek_beach.BayesianCoun
terfactualAuditor._create_default_equations").get("auto_param_L6132_21", 0.0) <= ci95_low
<= ci95_high <= get_parameter_loader().get("saaaaaaa.analysis.derek_beach.BayesianCounterfa
ctualAuditor._create_default_equations").get("auto_param_L6132_53", 1.0)
priority_monotonic = priority >= 0
risk_in_range = get_parameter_loader().get("saaaaaaa.analysis.derek_beach.BayesianCoun
terfactualAuditor._create_default_equations").get("auto_param_L6134_24", 0.0) <=
risk_score <= get_parameter_loader().get("saaaaaaa.analysis.derek_beach.BayesianCounterfact
ualAuditor._create_default_equations").get("auto_param_L6134_45", 1.0)

criteria_met = {
    'ci95_valid': ci95_valid,
    'priority_monotonic': priority_monotonic,
    'risk_in_range': risk_in_range
}

```

```

return {
    'risk_components': risk_components,
    'risk_score': risk_score,
    'success_probability': success_probability,
    'priority': priority,
    'recommendations': sorted(recommendations, reverse=True),
    'criteria_met': criteria_met
}

```

```

def refutation_and_sanity_checks(
    self,
    dag: nx.DiGraph,
    target: str,
    treatment: str,
    confounders: list[str] | None = None
) -> dict[str, Any]:
    """
    PROMPT III-3: Refutación, negativos y cordura do(.)

```

Ejecuta:

1. Controles negativos: nodos irrelevantes →  $|efecto| \leq \text{get_parameter_loader}().get("saaaaaaa.analysis.derek_beach.BayesianCounterfactualAuditor._create_default_equations").get("auto_param_L6162_64", 0.05)$
2. Pruebas placebo: permuta edges no causales
3. Sanity checks: añadir cofactores no reduce  $P(Y|do(X=1))$

Args:

dag: Grafo causal  
 target: Nodo objetivo Y  
 treatment: Nodo de tratamiento X  
 confounders: Lista de cofactores

Returns:

Dict con negative\_controls, placebo\_effect, sanity\_violations, recommendation

"""

confounders = confounders or []

self.logger.info("Running refutation and sanity checks...")

# 1. CONTROLES NEGATIVOS: nodos irrelevantes

```

irrelevant_nodes = [
    n for n in dag.nodes()
    if n not in (target, treatment) and not nx.has_path(dag, n, target)
]

```

negative\_effects = []

for node in irrelevant\_nodes[:5]: # Máximo 5 controles

try:

```

intervention = {node: get_parameter_loader().get("saaaaaaa.analysis.derek_beach.BayesianCounterfactualAuditor._create_default_equations").get("auto_param_L6188_38", 1.0)}

```

result = self.counterfactual\_query(intervention, target, {})

effect = abs(result['causal\_effect'])

negative\_effects.append(effect)

except Exception as e:

self.logger.warning(f"Negative control failed for {node}: {e}")

median\_negative\_effect = float(np.median(negative\_effects)) if negative\_effects

```

else get_parameter_loader().get("saaaaaaa.analysis.derek_beach.BayesianCounterfactualAuditor._create_default_equations").get("auto_param_L6195_93", 0.0)

```

```

negative_controls_ok = median_negative_effect <= get_parameter_loader().get("saaaaaaa.analysis.derek_beach.BayesianCounterfactualAuditor._create_default_equations").get("auto_param_L6196_57", 0.05)

```

# 2. PRUEBA PLACEBO: permuta edges no causales

placebo\_dag = dag.copy()

non\_causal\_edges = [

(u, v) for u, v in dag.edges()

```

if u != treatment and v != target
]

placebo_effect = get_parameter_loader().get("saaaaaaa.analysis.derek_beach.Bayesian
CounterfactualAuditor._create_default_equations").get("placebo_effect", 0.0) # Refactored
if non_causal_edges:
    # Permutar una arista
    edge_to_remove = non_causal_edges[0]
    placebo_dag.remove_edge(*edge_to_remove)

    # Medir efecto en DAG permutado
    scm_backup = self.scm
    try:
        self.construct_scm(placebo_dag)
        result = self.counterfactual_query({treatment: get_parameter_loader().get(
"saaaaaaa.analysis.derek_beach.BayesianCounterfactualAuditor._create_default_equations").ge
t("auto_param_L6215_63", 1.0)}, target, {})
        placebo_effect = abs(result['causal_effect'])
    except Exception as e:
        self.logger.warning(f"Placebo test failed: {e}")
    finally:
        self.scm = scm_backup

placebo_ok = placebo_effect <= get_parameter_loader().get("saaaaaaa.analysis.derek_
beach.BayesianCounterfactualAuditor._create_default_equations").get("auto_param_L6222_39",
0.05)

# 3. SANITY CHECKS: añadir cofactores activos no debe reducir P(Y|do(X=1))
sanity_violations = []

# Baseline: do(X=1)
try:
    baseline_result = self.counterfactual_query({treatment: get_parameter_loader()
.get("saaaaaaa.analysis.derek_beach.BayesianCounterfactualAuditor._create_default_equations
").get("auto_param_L6229_68", 1.0)}, target, {})
    baseline_p = baseline_result['p_counterfactual']

    # Con cofactores
    for confounder in confounders[:2]: # Máximo 2
        if confounder in dag.nodes():
            result_with_conf = self.counterfactual_query(
                {treatment: get_parameter_loader().get("saaaaaaa.analysis.derek_be
ach.BayesianCounterfactualAuditor._create_default_equations").get("auto_param_L6236_36",
1.0)},
                target,
                {confounder: get_parameter_loader().get("saaaaaaa.analysis.derek_be
ach.BayesianCounterfactualAuditor._create_default_equations").get("auto_param_L6238_37",
1.0)})
            p_with_conf = result_with_conf['p_counterfactual']

            # Verificar que no reduce significativamente
            if p_with_conf < baseline_p - get_parameter_loader().get("saaaaaaa.anal
ysis.derek_beach.BayesianCounterfactualAuditor._create_default_equations").get("auto_param
_L6243_50", 0.10):
                sanity_violations.append({
                    'confounder': confounder,
                    'baseline_p': float(baseline_p),
                    'p_with_confounder': float(p_with_conf),
                    'violation': f"Adding {confounder} reduced P(Y|do(X)) by
{baseline_p - p_with_conf:.3f}"})
except Exception as e:
    self.logger.error(f"Sanity checks failed: {e}")

sanity_ok = len(sanity_violations) == 0

# 4. DECISIÓN FINAL

```

```

all_checks_passed = negative_controls_ok and placebo_ok and sanity_ok

if not all_checks_passed:
    recommendation = "DEGRADE_ALL: Require DAG revision - observación prioritaria"
    self.logger.error(recommendation)
else:
    recommendation = "ACCEPT: All refutation tests passed"
    self.logger.info(recommendation)

return {
    'negative_controls': {
        'effects': [float(e) for e in negative_effects],
        'median': median_negative_effect,
        'passed': negative_controls_ok,
        'criterion': ' $\leq$  get_parameter_loader().get("saaaaaa.analysis.derek_beach.B
aysianCounterfactualAuditor._create_default_equations").get("auto_param_L6270_32", 0.05)'
    },
    'placebo_effect': {
        'effect': float(placebo_effect),
        'passed': placebo_ok,
        'criterion': ' $\approx 0$ '
    },
    'sanity_violations': sanity_violations,
    'sanity_passed': sanity_ok,
    'all_checks_passed': all_checks_passed,
    'recommendation': recommendation
}

```

```

def main() -> int:
    """CLI entry point"""
    parser = argparse.ArgumentParser(
        description="CDAF v2.0 - Framework de Deconstrucción y Auditoría Causal",
        formatter_class=argparse.RawDescriptionHelpFormatter,
        epilog="")
    Ejemplo de uso:
    python cdaf_framework.py documento.pdf --output-dir resultados/ --policy-code P1

```

Configuración:  
El framework busca config.yaml en el directorio actual.  
Use --config-file para especificar una ruta alternativa.

```

    """
)
parser.add_argument(
    "pdf_path",
    type=Path,
    help="Ruta al archivo PDF del Plan de Desarrollo Territorial"
)
parser.add_argument(
    "--output-dir",
    type=Path,
    default=Path("resultados_analisis"),
    help="Directorio de salida para los artefactos (default: resultados_analisis)"
)
parser.add_argument(
    "--policy-code",
    type=str,
    required=True,
    help="Código de política para nombrar los artefactos (ej: P1, PDT_2024)"
)
parser.add_argument(
    "--config-file",
    type=Path,
    default=Path(DEFAULT_CONFIG_FILE),
    help=f"Ruta al archivo de configuración YAML (default: {DEFAULT_CONFIG_FILE})"
)
```

```

)
parser.add_argument(
    "--log-level",
    choices=["DEBUG", "INFO", "WARNING", "ERROR"],
    default="INFO",
    help="Nivel de logging (default: INFO)"
)
parser.add_argument(
    "--pdet",
    action="store_true",
    help="Indica si el municipio es PDET (activa validación especial)"
)
args = parser.parse_args()

# Validate inputs
if not args.pdf_path.exists():
    print(f"ERROR: Archivo PDF no encontrado: {args.pdf_path}")
    return 1

# Initialize framework
try:
    framework = CDAFFramework(args.config_file, args.output_dir, args.log_level)

    # Configure PDET if specified
    if args.pdet and framework.dnp_validator:
        framework.dnp_validator.es_municipio_pdet = True
        framework.logger.info("Modo PDET activado - Validación especial habilitada")
except Exception as e:
    print(f"ERROR: No se pudo inicializar el framework: {e}")
    return 1

# Process document
success = framework.process_document(args.pdf_path, args.policy_code)

return 0 if success else 1

# =====
# PRODUCER CLASS - Registry Exposure
# =====

class DerekBeachProducer:
"""
    Producer wrapper for Derek Beach causal analysis with registry exposure

    Provides public API methods for orchestrator integration without exposing
    internal implementation details or summarization logic.

    Version: get_parameter_loader().get("saaaaaa.analysis.derek_beach.BayesianCounterfactu
alAuditor._create_default_equations").get("auto_param_L6373_13", 1.0).0
    Producer Type: Causal Mechanism Analysis
"""

    def __init__(self) -> None:
        """Initialize producer"""
        self.logger = logging.getLogger(self.__class__.__name__)
        self.logger.info("DerekBeachProducer initialized")

    # =====
    # EVIDENTIAL TESTS API
    # =====

@calibrated_method("saaaaaa.analysis.derek_beach.DerekBeachProducer.classify_test_type")
def classify_test_type(self, necessity: float, sufficiency: float) -> TestType:
    """Classify evidential test type based on necessity and sufficiency"""

```

```

return BeachEvidentialTest.classify_test(necessity, sufficiency)

def apply_test_logic(
    self,
    test_type: TestType,
    evidence_found: bool,
    prior: float,
    bayes_factor: float
) -> tuple[float, str]:
    """Apply Beach test-specific logic to Bayesian updating"""
    return BeachEvidentialTest.apply_test_logic(
        test_type, evidence_found, prior, bayes_factor
    )

@calibrated_method("saaaaaaa.analysis.derek_beach.DerekBeachProducer.is_hoop_test")
def is_hoop_test(self, test_type: TestType) -> bool:
    """Check if test is hoop test"""
    return test_type == "hoop_test"

@calibrated_method("saaaaaaa.analysis.derek_beach.DerekBeachProducer.is_smoking_gun")
def is_smoking_gun(self, test_type: TestType) -> bool:
    """Check if test is smoking gun"""
    return test_type == "smoking_gun"

@calibrated_method("saaaaaaa.analysis.derek_beach.DerekBeachProducer.is_doubly_decisional")
def is_doubly_decisional(self, test_type: TestType) -> bool:
    """Check if test is doubly decisional"""
    return test_type == "doubly_decisional"

@calibrated_method("saaaaaaa.analysis.derek_beach.DerekBeachProducer.is_straw_in_wind")
def is_straw_in_wind(self, test_type: TestType) -> bool:
    """Check if test is straw in wind"""
    return test_type == "straw_in_wind"

# =====
# HIERARCHICAL GENERATIVE MODEL API
# =====

def create_hierarchical_model(
    self,
    mechanism_priors: dict[str, float] | None = None
) -> HierarchicalGenerativeModel:
    """Create hierarchical generative model"""
    return HierarchicalGenerativeModel(mechanism_priors)

def infer_mechanism_posterior(
    self,
    model: HierarchicalGenerativeModel,
    observations: dict[str, Any],
    n_iter: int = 500,
    burn_in: int = 100,
    n_chains: int = 2
) -> dict[str, Any]:
    """Infer mechanism posterior using MCMC"""
    return model.infer_mechanism_posterior(
        observations, n_iter, burn_in, n_chains
    )

@calibrated_method("saaaaaaa.analysis.derek_beach.DerekBeachProducer.get_type_posterior")
def get_type_posterior(self, inference: dict[str, Any]) -> dict[str, float]:
    """Extract type posterior from inference"""
    return inference.get("type_posterior", {})

@calibrated_method("saaaaaaa.analysis.derek_beach.DerekBeachProducer.get_sequence_mode")
def get_sequence_mode(self, inference: dict[str, Any]) -> str:

```

```

"""Extract sequence mode from inference"""
return inference.get("sequence_mode", "")

@calibrated_method("saaaaaaa.analysis.derek_beach.DerekBeachProducer.get_coherence_score")
def get_coherence_score(self, inference: dict[str, Any]) -> float:
    """Extract coherence score from inference"""
    return inference.get("coherence_score", get_parameter_loader().get("saaaaaaa.analysis.derek_beach.DerekBeachProducer.get_coherence_score").get("auto_param_L6460_48", 0.0))

@calibrated_method("saaaaaaa.analysis.derek_beach.DerekBeachProducer.get_r_hat")
def get_r_hat(self, inference: dict[str, Any]) -> float:
    """Extract R-hat convergence diagnostic"""
    return inference.get("R_hat", get_parameter_loader().get("saaaaaaa.analysis.derek_beach.DerekBeachProducer.get_r_hat").get("auto_param_L6465_38", 1.0))

@calibrated_method("saaaaaaa.analysis.derek_beach.DerekBeachProducer.get_ess")
def get_ess(self, inference: dict[str, Any]) -> float:
    """Extract effective sample size"""
    return inference.get("ESS", get_parameter_loader().get("saaaaaaa.analysis.derek_beach.DerekBeachProducer.get_ess").get("auto_param_L6470_36", 0.0))

@calibrated_method("saaaaaaa.analysis.derek_beach.DerekBeachProducer.is_inference_uncertain")
def is_inference_uncertain(self, inference: dict[str, Any]) -> bool:
    """Check if inference has high uncertainty"""
    return inference.get("is_uncertain", False)

# =====#
# POSTERIOR PREDICTIVE CHECKS API
# =====#

def posterior_predictive_check(
    self,
    model: HierarchicalGenerativeModel,
    posterior_samples: list[dict[str, Any]],
    observed_data: dict[str, Any]
) -> dict[str, Any]:
    """Run posterior predictive checks"""
    return model.posterior_predictive_check(posterior_samples, observed_data)

@calibrated_method("saaaaaaa.analysis.derek_beach.DerekBeachProducer.get_ppd_p_value")
def get_ppd_p_value(self, ppc: dict[str, Any]) -> float:
    """Extract posterior predictive p-value"""
    return ppc.get("ppd_p_value", get_parameter_loader().get("saaaaaaa.analysis.derek_beach.DerekBeachProducer.get_ppd_p_value").get("auto_param_L6493_38", 0.0))

@calibrated_method("saaaaaaa.analysis.derek_beach.DerekBeachProducer.get_ablation_curve")
def get_ablation_curve(self, ppc: dict[str, Any]) -> dict[str, float]:
    """Extract ablation curve from PPC"""
    return ppc.get("ablation_curve", {})

@calibrated_method("saaaaaaa.analysis.derek_beach.DerekBeachProducer.get_ppc_recommendation")
def get_ppc_recommendation(self, ppc: dict[str, Any]) -> str:
    """Extract recommendation from PPC"""
    return ppc.get("recommendation", "")

# =====#
# CONDITIONAL INDEPENDENCE API
# =====#

def verify_conditional_independence(
    self,
    model: HierarchicalGenerativeModel,
    dag: nx.DiGraph,
    independence_tests: list[tuple[str, str, list[str]]] | None = None
)

```

```

) -> dict[str, Any]:
    """Verify conditional independencies in DAG"""
    return model.verify_conditional_independence(dag, independence_tests)

@calibrated_method("saaaaaaa.analysis.derek_beach.DerekBeachProducer.get_independence_t
ests")
def get_independence_tests(self, verification: dict[str, Any]) -> list[dict[str,
Any]]:
    """Extract independence tests from verification"""
    return verification.get("independence_tests", [])

@calibrated_method("saaaaaaa.analysis.derek_beach.DerekBeachProducer.get_delta_waic")
def get_delta_waic(self, verification: dict[str, Any]) -> float:
    """Extract delta WAIC from verification"""
    return verification.get("delta_waic", get_parameter_loader().get("saaaaaaa.analysis
.derek_beach.DerekBeachProducer.get_delta_waic").get("auto_param_L6526_46", 0.0))

@calibrated_method("saaaaaaa.analysis.derek_beach.DerekBeachProducer.get_model_preference")
def get_model_preference(self, verification: dict[str, Any]) -> str:
    """Extract model preference from verification"""
    return verification.get("model_preference", "inconclusive")

# =====
# COUNTERFACTUAL AUDITOR API
# =====

@calibrated_method("saaaaaa.analysis.derek_beach.DerekBeachProducer.create_auditor")
def create_auditor(self) -> BayesianCounterfactualAuditor:
    """Create Bayesian counterfactual auditor"""
    return BayesianCounterfactualAuditor()

def construct_scm(
    self,
    auditor: BayesianCounterfactualAuditor,
    dag: nx.DiGraph,
    structural_equations: dict[str, callable] | None = None
) -> dict[str, Any]:
    """Construct structural causal model"""
    return auditor.construct_scm(dag, structural_equations)

def counterfactual_query(
    self,
    auditor: BayesianCounterfactualAuditor,
    intervention: dict[str, float],
    target: str,
    evidence: dict[str, float] | None = None
) -> dict[str, Any]:
    """Execute counterfactual query"""
    return auditor.counterfactual_query(intervention, target, evidence)

@calibrated_method("saaaaaaa.analysis.derek_beach.DerekBeachProducer.get_causal_effect")
def get_causal_effect(self, query: dict[str, Any]) -> float:
    """Extract causal effect from query"""
    return query.get("causal_effect", get_parameter_loader().get("saaaaaaa.analysis.der
ek_beach.DerekBeachProducer.get_causal_effect").get("auto_param_L6564_42", 0.0))

@calibrated_method("saaaaaaa.analysis.derek_beach.DerekBeachProducer.is_sufficient")
def is_sufficient(self, query: dict[str, Any]) -> bool:
    """Check if mechanism is sufficient"""
    return query.get("is_sufficient", False)

@calibrated_method("saaaaaaa.analysis.derek_beach.DerekBeachProducer.is_necessary")
def is_necessary(self, query: dict[str, Any]) -> bool:
    """Check if mechanism is necessary"""
    return query.get("is_necessary", False)

```

```

@calibrated_method("saaaaaaa.analysis.derek_beach.DerekBeachProducer.is_effect_stable")
def is_effect_stable(self, query: dict[str, Any]) -> bool:
    """Check if effect is stable"""
    return query.get("effect_stable", False)

# =====
# RISK AGGREGATION API
# =====

def aggregate_risk(
    self,
    auditor: BayesianCounterfactualAuditor,
    omission_score: float,
    insufficiency_score: float,
    unnecessity_score: float,
    causal_effect: float,
    feasibility: float = get_parameter_loader().get("saaaaaaa.analysis.derek_beach.Dere
kBeachProducer.is_effect_stable").get("auto_param_L6592_29", 0.8),
    cost: float = get_parameter_loader().get("saaaaaaa.analysis.derek_beach.DerekBeachP
roducer.is_effect_stable").get("auto_param_L6593_22", 1.0)
) -> dict[str, Any]:
    """Aggregate risk and calculate priority"""
    return auditor.aggregate_risk_and_prioritize(
        omission_score,
        insufficiency_score,
        unnecessity_score,
        causal_effect,
        feasibility,
        cost
    )

@calibrated_method("saaaaaaa.analysis.derek_beach.DerekBeachProducer.get_risk_score")
def get_risk_score(self, aggregation: dict[str, Any]) -> float:
    """Extract risk score from aggregation"""
    return aggregation.get("risk_score", get_parameter_loader().get("saaaaaaa.analysis.
derek_beach.DerekBeachProducer.get_risk_score").get("auto_param_L6608_45", 0.0))

    @calibrated_method("saaaaaaa.analysis.derek_beach.DerekBeachProducer.get_success_probab
ility")
def get_success_probability(self, aggregation: dict[str, Any]) -> dict[str, float]:
    """Extract success probability from aggregation"""
    return aggregation.get("success_probability", {})

@calibrated_method("saaaaaaa.analysis.derek_beach.DerekBeachProducer.get_priority")
def get_priority(self, aggregation: dict[str, Any]) -> float:
    """Extract priority from aggregation"""
    return aggregation.get("priority", get_parameter_loader().get("saaaaaaa.analysis.de
rek_beach.DerekBeachProducer.get_priority").get("auto_param_L6618_43", 0.0))

@calibrated_method("saaaaaaa.analysis.derek_beach.DerekBeachProducer.get_recommendations")
def get_recommendations(self, aggregation: dict[str, Any]) -> list[str]:
    """Extract recommendations from aggregation"""
    return aggregation.get("recommendations", [])

# =====
# REFUTATION API
# =====

def refutation_checks(
    self,
    auditor: BayesianCounterfactualAuditor,
    dag: nx.DiGraph,
    target: str,
    treatment: str,
    confounders: list[str] | None = None
) -> dict[str, Any]:
    """Execute refutation and sanity checks"""

```

```

        return auditor.refutation_and_sanity_checks(
            dag, target, treatment, confounders
        )

    @calibrated_method("saaaaaaa.analysis.derek_beach.DerekBeachProducer.get_negative_controls")
    def get_negative_controls(self, refutation: dict[str, Any]) -> dict[str, Any]:
        """Extract negative controls from refutation"""
        return refutation.get("negative_controls", {})

    @calibrated_method("saaaaaaa.analysis.derek_beach.DerekBeachProducer.get_placebo_effect")
    def get_placebo_effect(self, refutation: dict[str, Any]) -> dict[str, Any]:
        """Extract placebo effect from refutation"""
        return refutation.get("placebo_effect", {})

    @calibrated_method("saaaaaaa.analysis.derek_beach.DerekBeachProducer.get_sanity_violations")
    def get_sanity_violations(self, refutation: dict[str, Any]) -> list[dict[str, Any]]:
        """Extract sanity violations from refutation"""
        return refutation.get("sanity_violations", [])

```

```

    @calibrated_method("saaaaaaa.analysis.derek_beach.DerekBeachProducer.all_checks_passed")
    def all_checks_passed(self, refutation: dict[str, Any]) -> bool:
        """Check if all refutation checks passed"""
        return refutation.get("all_checks_passed", False)

    @calibrated_method("saaaaaaa.analysis.derek_beach.DerekBeachProducer.get_refutation_recommendation")
    def get_refutation_recommendation(self, refutation: dict[str, Any]) -> str:
        """Extract recommendation from refutation"""
        return refutation.get("recommendation", "")

```

===== FILE: src/saaaaaaa/analysis/enhance\_recommendation\_rules.py =====

```
#!/usr/bin/env python3
```

```
"""
```

Script to enhance recommendation rules with 7 advanced features:

1. Template parameterization
2. Rule execution logic
3. Measurable indicators
4. Unambiguous time horizons
5. Testable verification
6. Cost tracking
7. Authority mapping

```
"""
```

```
import copy
import json
from pathlib import Path
from typing import Any
from saaaaaaa.core.calibration.decorators import calibrated_method
```

```
def enhance_template(rule: dict[str, Any]) -> dict[str, Any]:
    """
```

```
    Feature 1: Eliminate Hardcoded Template Strings
    Replace with template_id and template_params
    """
```

```
    template = rule.get('template', {})
    rule_id = rule.get('rule_id', "")
    level = rule.get('level', "")
```

```
    # Extract parameters from template strings
    template_params = {}
```

```
    if level == 'MICRO':
        when = rule.get('when', {})
```

```

template_params = {
    'pa_id': when.get('pa_id', 'PA01'),
    'dim_id': when.get('dim_id', 'DIM01'),
    'question_id': 'Q001' # Would be derived from context
}
if level == 'MESO':
    when = rule.get('when', {})
    template_params = {
        'cluster_id': when.get('cluster_id', 'CL01')
    }

# Create enhanced template with ID
enhanced_template = copy.deepcopy(template)
enhanced_template['template_id'] = f"TPL-{rule_id}"
enhanced_template['template_params'] = template_params

return enhanced_template

def add_execution_logic(rule: dict[str, Any]) -> dict[str, Any]:
    """
    Feature 2: Add Rule Execution Logic
    """
    when = rule.get('when', {})
    level = rule.get('level', "")

    # Build trigger condition string
    trigger_parts = []
    if level == 'MICRO':
        pa_id = when.get('pa_id', "")
        dim_id = when.get('dim_id', "")
        score_lt = when.get('score_lt', 1.65)
        trigger_parts.append(f"score < {score_lt}")
        trigger_parts.append(f"pa_id = '{pa_id}'")
        trigger_parts.append(f"dim_id = '{dim_id}'")
    elif level == 'MESO':
        cluster_id = when.get('cluster_id', "")
        score_band = when.get('score_band', "")
        if score_band:
            trigger_parts.append(f"score_band = '{score_band}'")
        if cluster_id:
            trigger_parts.append(f"cluster_id = '{cluster_id}'")
    elif level == 'MACRO':
        macro_band = when.get('macro_band', "")
        if macro_band:
            trigger_parts.append(f"macro_band = '{macro_band}'")

    trigger_condition = " AND ".join(trigger_parts) if trigger_parts else "true"

    return {
        "trigger_condition": trigger_condition,
        "blocking": False,
        "auto_apply": False,
        "requires_approval": True,
        "approval_roles": ["Secretaría de Planeación", "Secretaría de Hacienda"]
    }

def enhance_indicator(indicator: dict[str, Any], rule_id: str, level: str) -> dict[str, Any]:
    """
    Feature 3: Make Indicators Measurable
    """
    enhanced = copy.deepcopy(indicator)

    # Add formula based on indicator type
    indicator.get('name', "")
    if 'proporción' in indicator.get('unit', ""):
        enhanced['formula'] = 'COUNT(compliant_items) / COUNT(total_items)'
        enhanced['acceptable_range'] = [0.6, 1.0]


```

```

        elif 'porcentaje' in indicator.get('unit', ""):
            enhanced['formula'] = '(achieved / target) * 100'
            enhanced['acceptable_range'] = [60.0, 100.0]
        else:
            enhanced['formula'] = 'SUM(verified_artifacts)'
            enhanced['acceptable_range'] = [indicator.get('target', 1) * 0.7,
                indicator.get('target', 1)]
    }

    # Add measurement metadata
    enhanced['baseline_measurement_date'] = '2024-01-01'
    enhanced['measurement_frequency'] = 'mensual'
    enhanced['data_source'] = 'Sistema de Seguimiento de Planes (SSP)'
    enhanced['data_source_query'] = f"SELECT COUNT(*) FROM indicators WHERE indicator_id = {rule_id}-IND"
    enhanced['responsible_measurement'] = 'Oficina de Planeación Municipal'
    enhanced['escalation_if_below'] = enhanced['acceptable_range'][0]

    return enhanced

def enhance_horizon(horizon: dict[str, str], rule_id: str) -> dict[str, Any]:
    """
    Feature 4: Define Unambiguous Time Horizons
    """

    # Map T0, T1, T2, T3 to actual durations
    duration_map = {
        'T0': 0,
        'T1': 6, # 6 months
        'T2': 12, # 12 months
        'T3': 24 # 24 months
    }

    start = horizon.get('start', 'T0')
    end = horizon.get('end', 'T1')

    duration_months = duration_map.get(end, 6) - duration_map.get(start, 0)

    # Create milestones
    milestones = []
    if duration_months >= 6:
        milestones.append({
            "name": "Inicio de implementación",
            "offset_months": 1,
            "deliverables": ["Plan de trabajo aprobado"],
            "verification_required": True
        })
    if duration_months >= 12:
        milestones.append({
            "name": "Revisión intermedia",
            "offset_months": duration_months // 2,
            "deliverables": ["Informe de avance"],
            "verification_required": True
        })
    milestones.append({
        "name": "Entrega final",
        "offset_months": duration_months,
        "deliverables": ["Todos los productos esperados"],
        "verification_required": True
    })

    return {
        "start": start,
        "end": end,
        "start_type": "plan_approval_date",
        "duration_months": duration_months,
        "milestones": milestones,
        "dependencies": [],
        "critical_path": duration_months <= 6
    }
}

```

```

def enhance_verification(verification: list[str], rule_id: str) -> list[dict[str, Any]]:
    """
    Feature 5: Make Verification Testable
    """
    enhanced_verifications = []

    for idx, artifact_text in enumerate(verification, 1):
        # Determine artifact type
        artifact_type = "DOCUMENT"
        if any(word in artifact_text.lower() for word in ['sistema', 'repositorio', 'registro', 'base de datos']):
            artifact_type = "SYSTEM_STATE"

        # Create structured verification
        ver_obj = {
            "id": f"VER-{rule_id}-{idx:03d}",
            "type": artifact_type,
            "artifact": artifact_text,
            "format": "PDF" if artifact_type == "DOCUMENT" else "DATABASE_QUERY",
            "required_sections": ["Objetivo", "Alcance", "Resultados"] if artifact_type ==
            "DOCUMENT" else [],
            "approval_required": True,
            "approver": "Secretaría de Planeación",
            "due_date": "T1",
            "automated_check": artifact_type == "SYSTEM_STATE"
        }
        }

        # Add validation for system states
        if artifact_type == "SYSTEM_STATE":
            ver_obj["validation_query"] = f"SELECT COUNT(*) FROM artifacts WHERE
artifact_id = '{ver_obj['id']}'"
            ver_obj["pass_condition"] = "COUNT(*) >= 1"

        enhanced_verifications.append(ver_obj)

    return enhanced_verifications

def add_budget(rule: dict[str, Any]) -> dict[str, Any]:
    """
    Feature 6: Integrate Cost Tracking
    """
    level = rule.get('level', '')

    # Estimate costs based on level and complexity
    if level == 'MICRO':
        base_cost = 45_000_000 # COP
    elif level == 'MESO':
        base_cost = 150_000_000
    elif level == 'MACRO':
        base_cost = 500_000_000
    else:
        base_cost = 50_000_000

    # Cost breakdown
    personal_cost = int(base_cost * 0.55)
    consultancy_cost = int(base_cost * 0.30)
    technology_cost = int(base_cost * 0.15)

    return {
        "estimated_cost_cop": base_cost,
        "cost_breakdown": {
            "personal": personal_cost,
            "consultancy": consultancy_cost,
            "technology": technology_cost
        },
        "funding_sources": [
            {

```

```

        "source": "SGP - Sistema General de Participaciones",
        "amount": int(base_cost * 0.60),
        "confirmed": False
    },
    {
        "source": "Recursos Propios",
        "amount": int(base_cost * 0.40),
        "confirmed": False
    }
],
"fiscal_year": 2025
}

def enhance_responsible(responsible: dict[str, Any]) -> dict[str, Any]:
"""
Feature 7: Map Authority for Accountability
"""

enhanced = copy.deepcopy(responsible)

# Add legal mandate
entity = responsible.get('entity', '')
if 'Mujer' in entity:
    legal_mandate = "Ley 1257 de 2008 - Normas para la prevención de violencias contra
la mujer"
elif 'Planeación' in entity:
    legal_mandate = "Ley 152 de 1994 - Ley Orgánica del Plan de Desarrollo"
elif 'Hacienda' in entity:
    legal_mandate = "Ley 819 de 2003 - Responsabilidad Fiscal"
else:
    legal_mandate = "Estatuto Orgánico Municipal"

enhanced['legal_mandate'] = legal_mandate

# Add approval chain
enhanced['approval_chain'] = [
{
    "level": 1,
    "role": "Director/Coordinador de Programa",
    "decision": "Aprueba plan de trabajo"
},
{
    "level": 2,
    "role": "Secretario/a de la entidad responsable",
    "decision": "Aprueba presupuesto y recursos"
},
{
    "level": 3,
    "role": "Secretaría de Planeación",
    "decision": "Valida coherencia con PDM"
},
{
    "level": 4,
    "role": "Alcalde Municipal",
    "decision": "Aprobación final (si aplica)"
}
]

# Add escalation path
enhanced['escalation_path'] = {
    "threshold_days_delay": 15,
    "escalate_to": "Secretaría de Planeación",
    "final_escalation": "Despacho del Alcalde",
    "consequences": ["Revisión presupuestal", "Reasignación de responsables"]
}

return enhanced

def enhance_rule(rule: dict[str, Any]) -> dict[str, Any]:

```

```

"""Enhance a single rule with all 7 features"""
enhanced_rule = copy.deepcopy(rule)

# 1. Template parameterization
enhanced_rule['template'] = enhance_template(rule)

# 2. Execution logic
enhanced_rule['execution'] = add_execution_logic(rule)

# 3. Measurable indicators
if 'indicator' in enhanced_rule['template']:
    enhanced_rule['template']['indicator'] = enhance_indicator(
        enhanced_rule['template']['indicator'],
        rule.get('rule_id', ''),
        rule.get('level', '')
    )

# 4. Unambiguous time horizons
if 'horizon' in enhanced_rule['template']:
    enhanced_rule['template']['horizon'] = enhance_horizon(
        enhanced_rule['template']['horizon'],
        rule.get('rule_id', '')
    )

# 5. Testable verification
if 'verification' in enhanced_rule['template']:
    enhanced_rule['template']['verification'] = enhance_verification(
        enhanced_rule['template']['verification'],
        rule.get('rule_id', '')
    )

# 6. Budget tracking
enhanced_rule['budget'] = add_budget(rule)

# 7. Authority mapping
if 'responsible' in enhanced_rule['template']:
    enhanced_rule['template']['responsible'] = enhance_responsible(
        enhanced_rule['template']['responsible']
    )

return enhanced_rule

def main() -> None:
    """Main enhancement process"""
    # Delegate to factory for I/O operations
    from .factory import load_json, save_json

    # Load existing rules
    rules_path = Path('config/recommendation_rules.json')
    rules_data = load_json(rules_path)

    print(f"Loaded {len(rules_data['rules'])} rules from {rules_path}")

    # Enhance all rules
    enhanced_rules = []
    for i, rule in enumerate(rules_data['rules'], 1):
        try:
            enhanced = enhance_rule(rule)
            enhanced_rules.append(enhanced)
            if i % 10 == 0:
                print(f"Enhanced {i}/{len(rules_data['rules'])} rules...")
        except Exception as e:
            print(f"Error enhancing rule {rule.get('rule_id', 'UNKNOWN')}: {e}")
            enhanced_rules.append(rule) # Keep original if enhancement fails

    # Create enhanced data structure
    enhanced_data = {
        'version': '2.0', # Increment version

```

```

'enhanced_features': [
    'template_parameterization',
    'execution_logic',
    'measurable_indicators',
    'unambiguous_time_horizons',
    'testable_verification',
    'cost_tracking',
    'authority_mapping'
],
'rules': enhanced_rules
}

# Save enhanced rules
output_path = Path('config/recommendation_rules_enhanced.json')
save_json(enhanced_data, output_path)

print(f"\nEnhanced {len(enhanced_rules)} rules saved to {output_path}")
print(f"Original file preserved at {rules_path}")

# Show sample enhanced rule
if enhanced_rules:
    print("\n==== Sample Enhanced Rule ===")
    print(json.dumps(enhanced_rules[0], indent=2, ensure_ascii=False)[:2000])
    print("...")

```

# Note: Main entry point removed to maintain I/O boundary separation.

# For usage examples, see examples/ directory.

===== FILE: src/saaaaaa/analysis/factory.py =====

"""

### Factory Layer for Analysis Module I/O Operations

This module provides centralized I/O operations for the analysis package, implementing a clean separation between I/O and business logic following the Ports and Adapters (Hexagonal Architecture) pattern.

All file I/O for the analysis package should be handled through this factory.

"""

```

import csv
import json
import logging
from functools import lru_cache
from pathlib import Path
from typing import Any
from saaaaaa.core.calibration.decorators import calibrated_method

try:
    import yaml
except ImportError:
    yaml = None

try:
    import fitz # PyMuPDF
except ImportError:
    fitz = None

try:
    import pdfplumber
except ImportError:
    pdfplumber = None

try:
    import spacy
except ImportError:
    spacy = None

logger = logging.getLogger(__name__)

```

```

_PROJECT_ROOT = Path(__file__).resolve().parents[3]
_CALIBRATION_SEARCH_PATHS: tuple[Path, ...] = (
    _PROJECT_ROOT / "config" / "calibraciones",
    _PROJECT_ROOT / "config",
    _PROJECT_ROOT,
)
# =====
# JSON I/O OPERATIONS
# =====

def load_json(file_path: str | Path) -> dict[str, Any]:
    """
    Load JSON data from file.

    Args:
        file_path: Path to JSON file

    Returns:
        Dict containing the loaded JSON data

    Raises:
        FileNotFoundError: If file doesn't exist
        json.JSONDecodeError: If file contains invalid JSON
    """
    file_path = Path(file_path)

    if not file_path.exists():
        raise FileNotFoundError(f"File not found: {file_path}")

    with open(file_path, encoding="utf-8") as f:
        data = json.load(f)

    logger.info(f"Loaded JSON from {file_path}")
    return data

def save_json(data: dict[str, Any], file_path: str | Path, indent: int = 2) -> None:
    """
    Save data to JSON file with formatted output.

    Args:
        data: Dictionary to save
        file_path: Path to output JSON file
        indent: Indentation level for formatting
    """
    file_path = Path(file_path)
    file_path.parent.mkdir(parents=True, exist_ok=True)

    with open(file_path, "w", encoding="utf-8") as f:
        json.dump(data, f, ensure_ascii=False, indent=indent)

    logger.info(f"Saved JSON to {file_path}")

# =====
# YAML I/O OPERATIONS
# =====

def load_yaml(file_path: str | Path) -> dict[str, Any]:
    """
    Load YAML data from file.

    Args:
        file_path: Path to YAML file

    Returns:
        Dict containing the loaded YAML data

```

Raises:

ImportError: If PyYAML is not installed